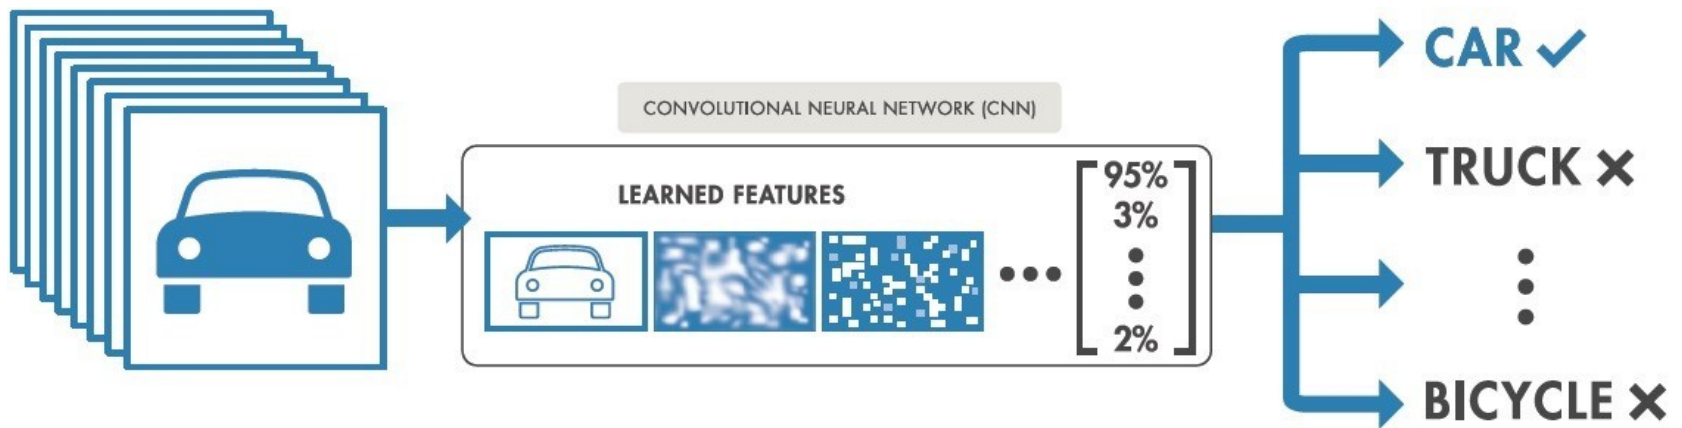



Convolutional Neural Network (CNN) and --- Recursive Neural Networks (RNN)


Image Classification



Hierarchical Features of Images



Low Level
(Edges, dark spots)

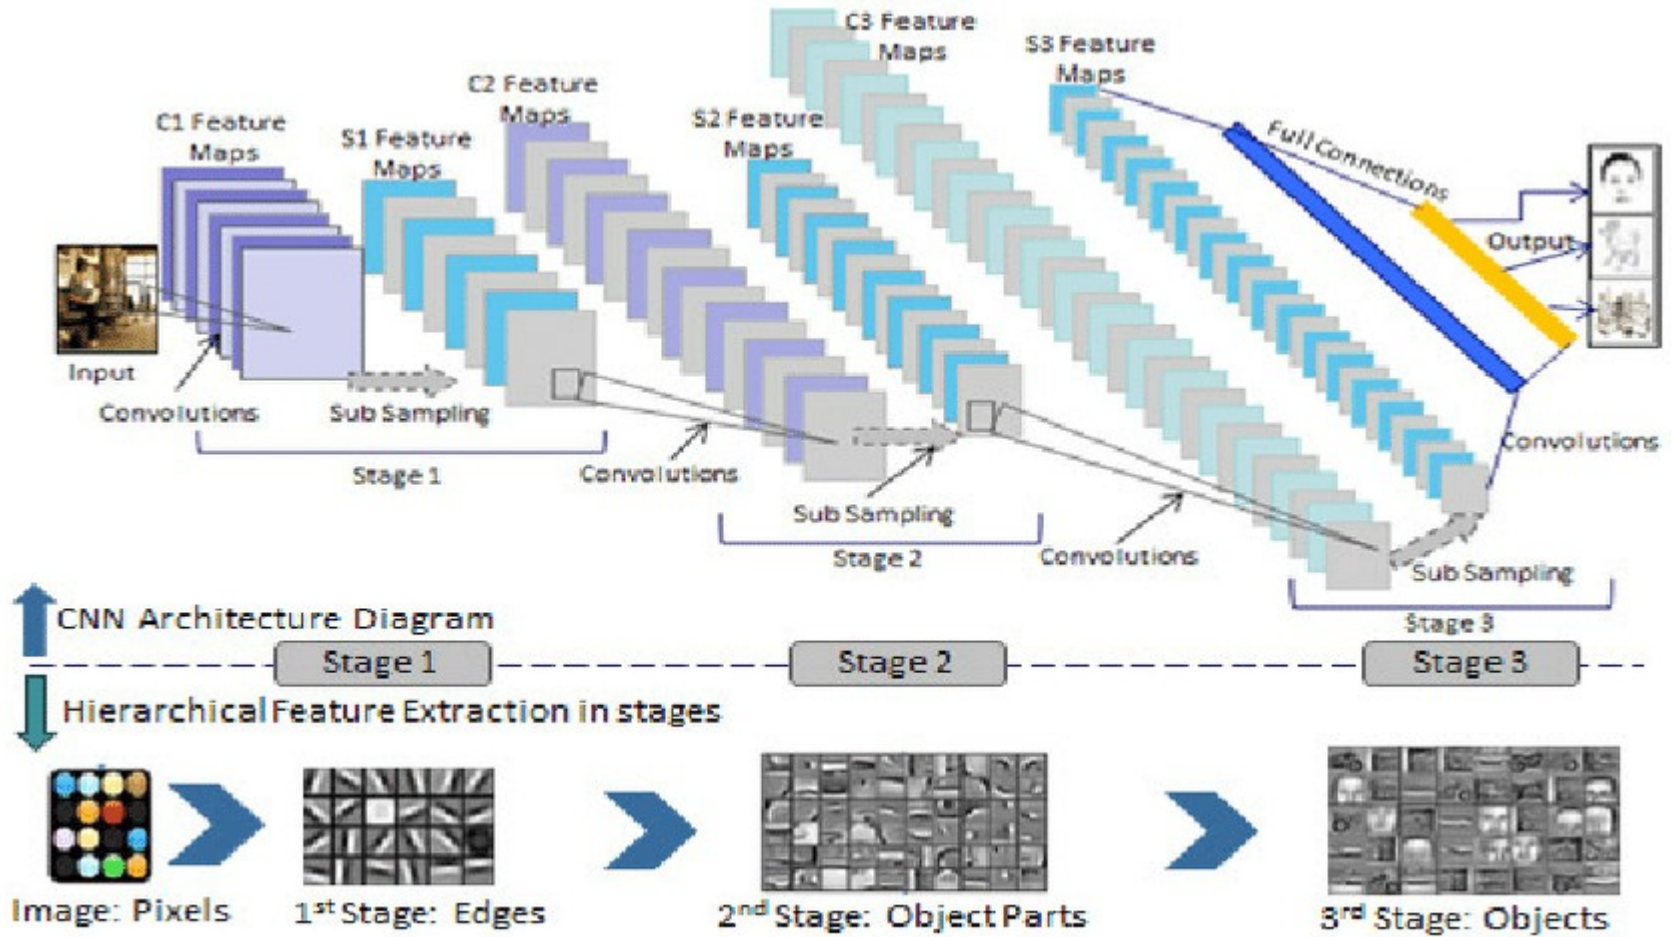


Mid Level
(Eyes, ears, etc.)



High Level
(Facial structure)

Feature Detection for Classification



CNN: Keywords

- **Convolutional layer**

- Feature matching
- Convolutional block/filter

- **Pooling layer:**

- Dimension reduction
- Max-pooling

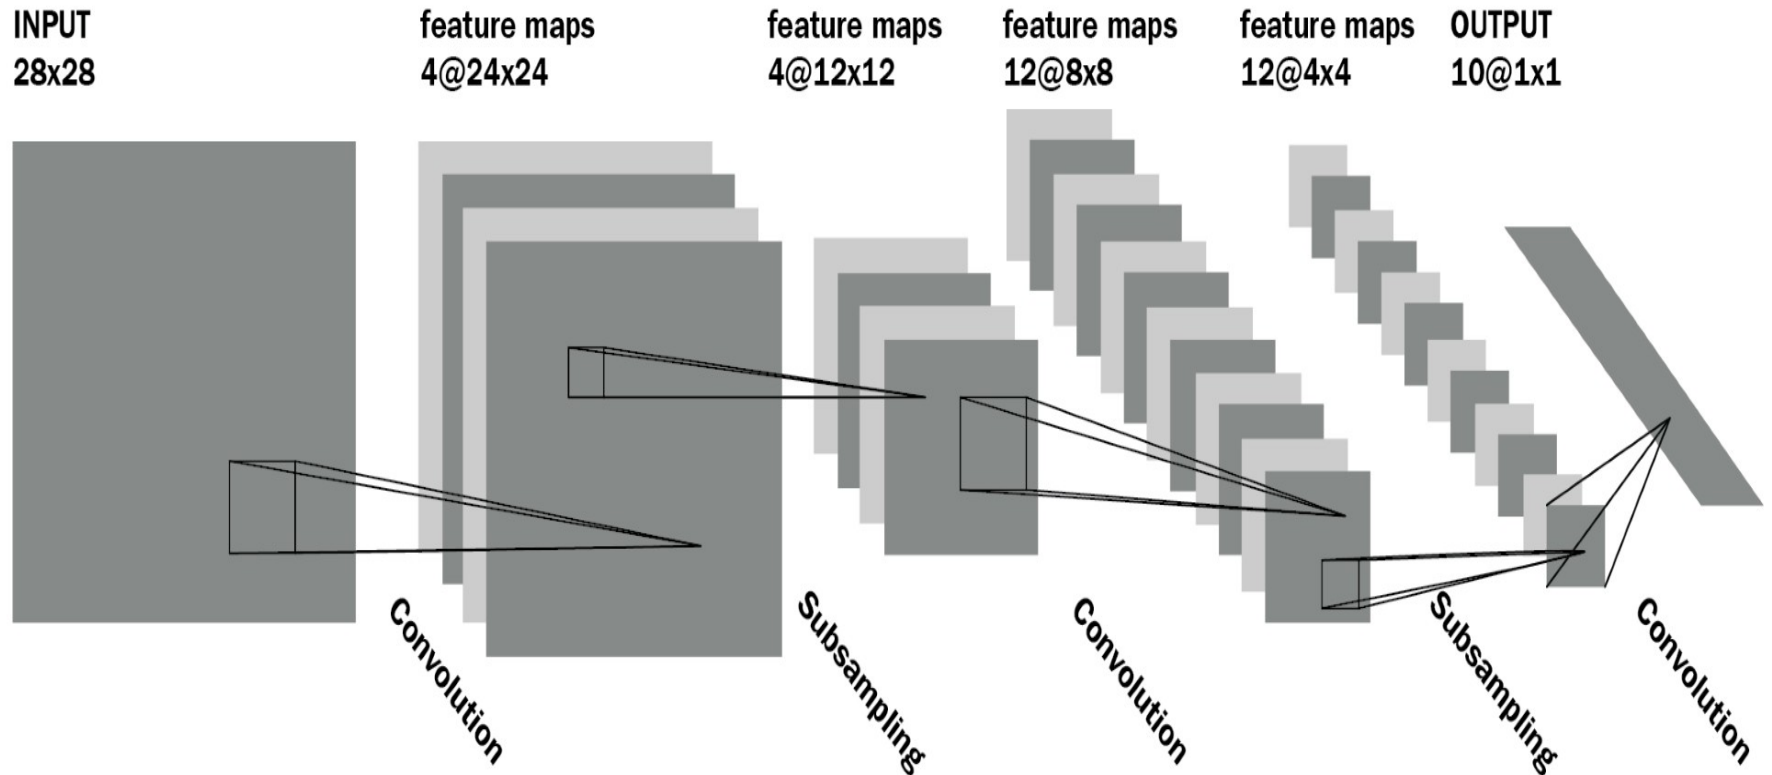
- **Dropout:**

- Address the overfitting issue
- Removing nodes randomly from a network temporarily during training to prevent one of a small number of nodes from getting largest weights and dominating the outputs

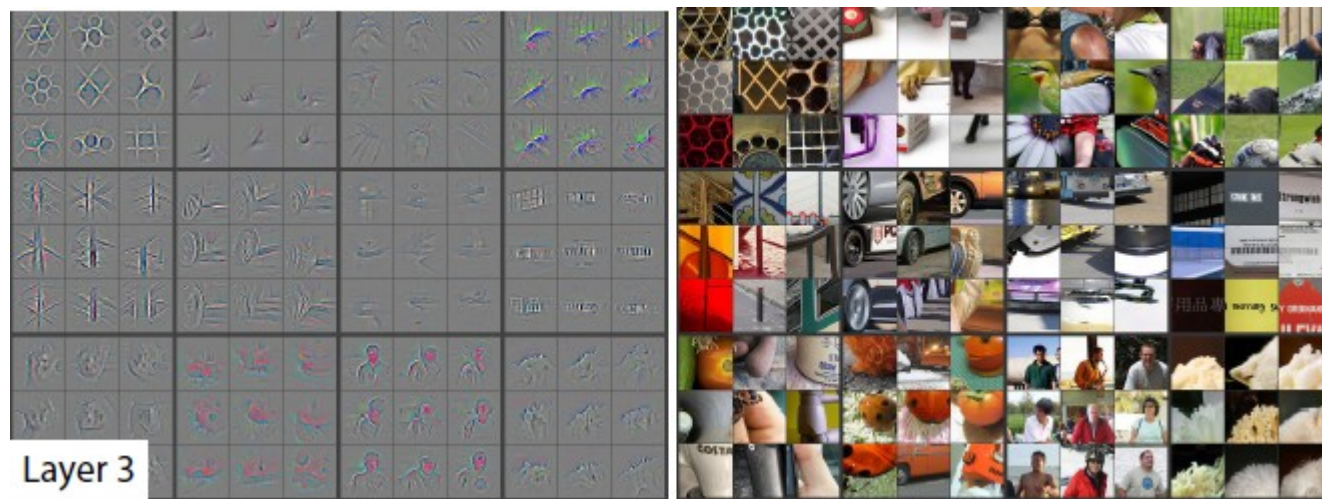

- **Flatten layer**

- Reshape the data to the shape of the required output
- Sigmoid function and softmax function

CNN: The **LeNet** architecture



CNN: Convolution



CNN: Convolution

	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						

	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						

	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						

	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						

	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						
	A	B	C	D	E	F
1						
2						
3						
4						
5						
6						

CNN: Convolution

Input Layer							Conv. Block			Output Layer			
	A	B	C	D	E	F							
1	0.5	0.3	0.1				3	1	1	6.3	XX	XX	XX
2	0.2	0.6	0.1				1	3	1	XX	XX	XX	XX
3	0.1	0.1	0.7				1	1	3	XX	XX	XX	XX
4				0.5	0.6	0.7				XX	XX	XX	3.6
5				0.2	0.1	0.1							
6				0.1	0.1	0.0							

CNN: Example

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

?

=

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	1	-1	-1
-1	-1	-1	1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1


?

=

X


CNN: Example

Features of X




CNN: Example

Filters to Detect X Features

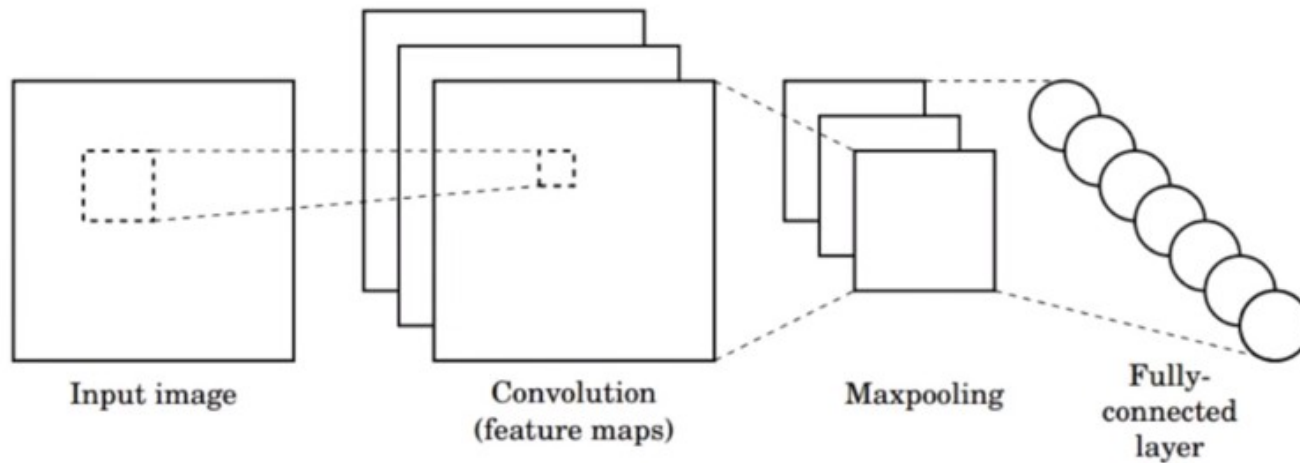


CNN: Example

The Convolution Operation



CNN: Architecture of Operation




1. **Convolution:** Apply filters with learned weights to generate feature maps.
2. **Non-linearity:** Often ReLU.
3. **Pooling:** Downsampling operation on each feature map.

Train model with image data.

Learn weights of filters in convolutional layers.

CNN: Architecture of Operation



4x4 filter: matrix
of weights w_{ij}

$$\sum_{i=1}^4 \sum_{j=1}^4 w_{ij} x_{i+p,j+q} + b$$

for neuron (p,q) in hidden layer


For a neuron in hidden layer:

- Take inputs from patch
- Compute weighted sum
- Apply bias


- 1) applying a window of weights
- 2) computing linear combinations
- 3) activating with non-linear function

CNN: Non-Linearity

- Apply after every convolution operation (i.e., after convolutional layers)
- ReLU: pixel-by-pixel operation that replaces all negative values by zero. **Non-linear operation!**



Rectified Linear Unit (ReLU)

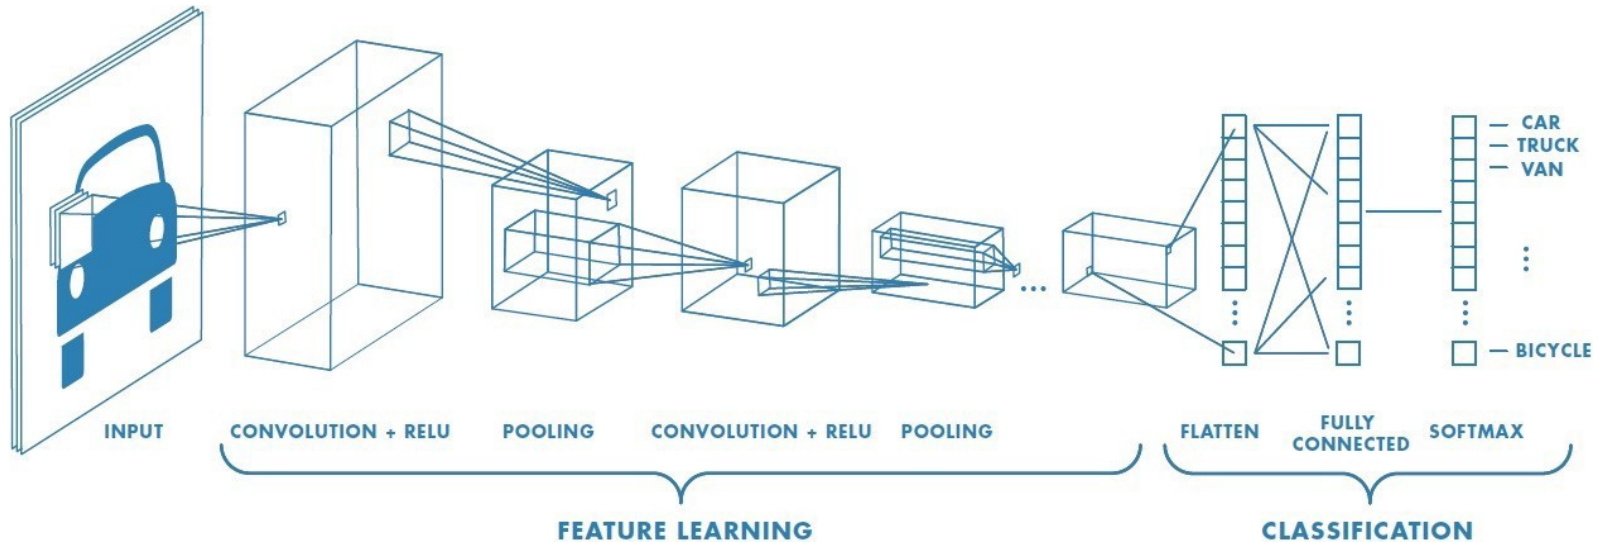


CNN: Max-Pooling

	A	B	C	D	E	F
1	7	0	8	6	2	1
2	6	6	0	8	4	2
3	2	2	1	2	3	6
4	5	0	2	2	2	5
5	4	2	2	1	9	7
6	7	2	0	5	2	4

7	8	4
5	2	6
7	5	9

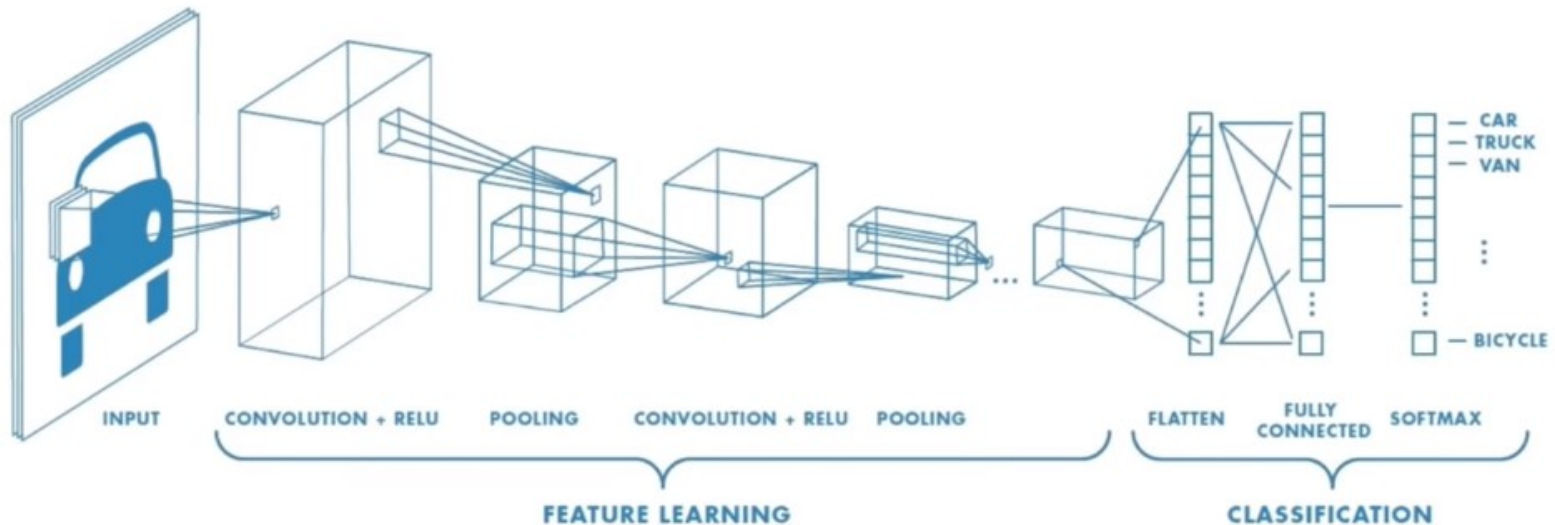
CNN: Architecture of Operation



- CONV and POOL layers output high-level features of input
- Fully connected layer uses these features for classifying input image
- Express output as **probability** of image belonging to a particular class

$$\text{softmax}(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

CNN: Backward Propagation for Training




Learn weights for convolutional filters and fully connected layers

Backpropagation: cross-entropy loss

$$J(\theta) = \sum_i y^{(i)} \log(\hat{y}^{(i)})$$

ImageNet Challenge: Classification



2012: AlexNet. First CNN to win.

- 8 layers, 61 million parameters

2013: ZFNet

- 8 layers, more filters

2014: VGG

- 19 layers




2014: GoogLeNet

- "Inception" modules
- 22 layers, 5 million parameters

2015: ResNet

- 152 layers

CNN: Dropout




CNN: Beyond Classification

Semantic Segmentation



CAT

Object Detection



CAT, DOG, DUCK

Image Captioning



The cat is in the grass.


Recurrent Neural Network

- **Memory**

- Your understanding of something is based on your understanding of previous things

- **RNN:** can be thought of as multiple copies of the same network, each passing a message to a successor

Recurrent Neural Network



Apply a **recurrence relation** at every time step to process a sequence:


$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, x_t)$$

cell state function parameterized by W old state

Update Hidden State

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$


Recurrent Neural Network



LSTM: Long Short Term Memory

- A specific case of RNN that remember information for long periods of time.


“I grew up in France, ... and I speak fluent ____.”



LSTM: Keywords


- **Cell State**

- The information passing through



- **Gate**

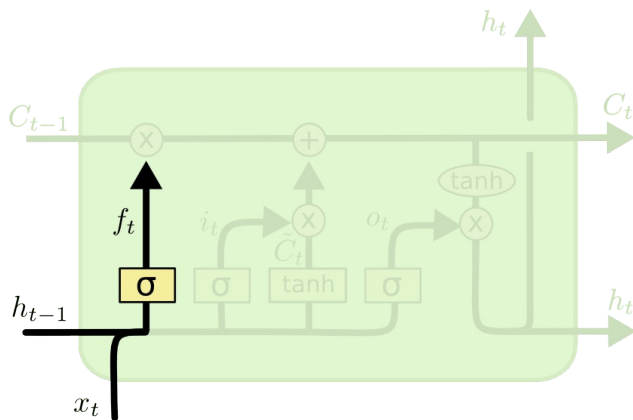
- To protect or filter information, control the information to pass through (sigmoid layer and pointwise multiplication)



LSTM: Keywords

■ Step 1: Forget

- A **forget gate layer** decides what information we're going to throw away from the cell state, usually a *sigmoid* layer that outputs a value between 0 and 1
- Sigmoid: outputs between 0 – completely forget vs. 1- completely keep




$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

LSTM: Keywords

■ Step 2: Update

- Decide what information we're going to store in the cell state
- First, a sigmoid layer called the “input gate layer” decides which values we'll update.
- Next, a *tanh* layer creates a vector of new candidate values that could be added to the state.

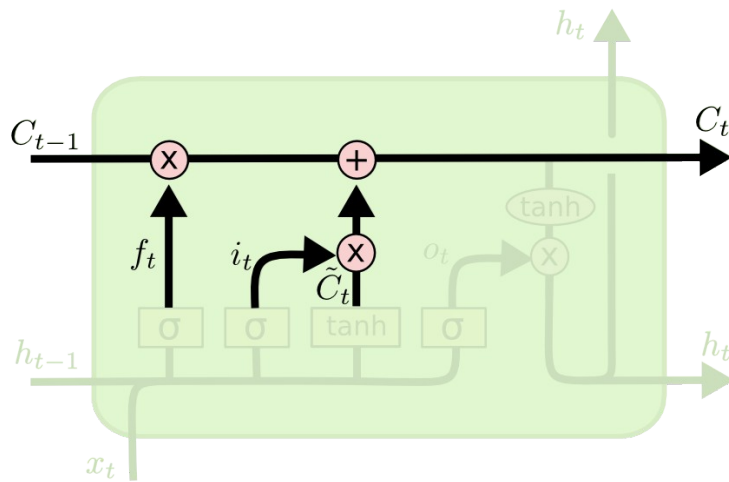


$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

LSTM: Keywords

■ Step 2: Update



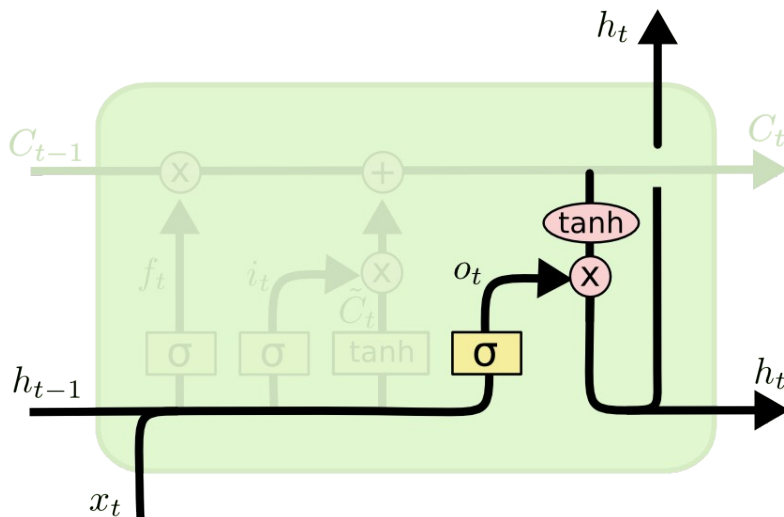
$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

- Apply forget operation to previous internal cell state: $f_t * C_{t-1}$
- Add new candidate values, scaled by how much we decided to update: $i_t * \tilde{C}_t$

LSTM: Keywords

■ Step 3: Output

- Decide what we're going to output
- Tanh: to push the values to be between -1 and 1




- Sigmoid layer: decide what parts of state to output

$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$


$$h_t = o_t * \tanh(C_t)$$

- Tanh layer: squash values between -1 and 1
- $o_t * \tanh(C_t)$: output filtered version of cell state

LSTM: Training




RNN: Sentiment Analysis




Input: sequence of words

Output: probability of having positive sentiment


```
 loss = tf.nn.softmax_cross_entropy_with_logits(  
    labels=model.y, logits=model.pred  
)
```

RNN: Music Generation



Input: sheet music

Output: next character in sheet music



RNN: Machine Translation

