

# PySpark Regression Example

For Python programmers

# Data process steps for regression

1. Read the data
2. Review the data
3. Split data into training & testing sets
4. Specify the regression (model / formula)
5. Fit the data to the model
6. Check the model fit

# Read data

```
1 %sh
2 wget https://raw.githubusercontent.com/bcbarsness/machine-learning/master/USA_Housing.csv
```

```
1 df = spark.read.csv('file:/databricks/driver/USA_Housing.csv', inferSchema=True, header=True, mode='DROPMALFORMED')
```

# Review data

1display(df)

▶ (1) Spark Jobs

	Avg_Area_Income ▲	Avg_Area_House_Age ▲	Avg_Area_Number_of_Rooms ▲	Avg_Area_Number_of_Bedrooms ▲	Area_Population ▲
1	79545.45857431678	5.682861321615587	7.009188142792237	4.09	23086.800502686456
2	79248.64245482568	6.0028998082752425	6.730821019094919	3.09	40173.07217364482

1

display(df.summary())

▶ (2) Spark Jobs

	summary ▲	Avg_Area_Income ▲	Avg_Area_House_Age ▲	Avg_Area_Number_of_Rooms▲	Avg_Area_Number_of_Bedrooms▲
1	count	5000	5000	5000	5000
2	mean	68583.10898395971	5.97722203528029	6.987791850907942	3.9813299999999967
3	stddev	10657.991213830363	0.9914561798281722	1.0058332312773866	1.2341372654846832
4	min	17796.631189543397	2.644304186036705	3.2361940234262048	2.0
5	25%	61478.633929567324	5.322269839263871	6.298966338516728	3.14
6	50%	68803.55207659505	5.969905376273397	7.002864274301249	4.05
7	75%	75782.33514026614	6.650746733224263	7.665643100697559	4.49

# Split data into training & testing sets

```
1 train_data,test_data = df.randomSplit([0.6, 0.4], 24) # proportions [], seed for random
2
3 print("Number of training records: " + str(train_data.count()))
4 print("Number of testing records : " + str(test_data.count()))
```

# Specify the regression formula

```
1  from pyspark.ml.feature import RFormula
2  columns = df.columns # all columns
3  # Not using Price (label) or address in features
4  columns.remove('Price')
5  columns.remove('Address')
6  # Careful! Capitalization does matter Price vs price
7  formula = "{} ~ {}".format("Price", " + ".join(columns))
8  print( "Formula : {}".format(formula))
9  r_formula = RFormula(formula = formula)
10 r_formula
```

```
Formula : Price ~ Avg_Area_Income + Avg_Area_House_Age + Avg_Area_Number_of_Rooms + Avg_Area_Number_of_Bedrooms + Area_Population
Out[35]: RFormula_2ca730639436
```

# Create RFormula for data

PySpark: it's common to fit & transform, which customizes the model to the data and then fits the model to the data

`fit()`

Create RFormula model given the training data.  
RFormula looks at the data (column types) to create the model

`transform()`

Create a DataFrame containing the fitted RFormula model

```
1 # RFormula must review the data (fit) to handle categorical (string) variables before it can run its transformation
2 trained_RF_model = r_formula.fit(train_data) # create model based on data
3 # Using the RFormula model, create a DataFrame of the transformed model
4 trained_model_DF = trained_RF_model.transform(train_data) # model of RFormula for data
5 display(trained_model_DF)
```

▶ (1) Spark Jobs

▶ `trained_model_DF`: `pyspark.sql.dataframe.DataFrame` = [Avg\_Area\_Income: double, Avg\_Area\_House\_Age: double ... 7 more fields]

		Address	features	label
1	33597895555	9932 Eric Circles	<code>{"vectorType": "dense", "length": 5, "values": [17796.631189543397, 4.9495570055571125, 6.713905444702088, 2.5, 47162.183643191434]}</code>	302355.83597895555
2	.577726322	Unit 4700 Box 1880	<code>{"vectorType": "dense", "length": 5, "values": [35454.714659475445, 6.855708363901107, 6.018646502679608, 4.5, 59636.40255302499]}</code>	1077805.577726322

Two important columns:  
features (the Xs from Python), as DenseVector of numbers, same order as data columns  
label (the Y from Python)

# Apply regression with RFormula & test the model

fit()

Create a model given the data and RFormula model  
LinearRegression will use the RFormula for the regression

transform()

Create a DataFrame applying the fitted model to the data

```
1 from pyspark.ml.regression import LinearRegression
2 lr = LinearRegression(labelCol="label", featuresCol="features")
3 train_fittedLR = lr.fit(trained_model_DF)
4 test_transformedLR = train_fittedLR.transform(test_preparedDF)
5 display(test_transformedLR)
```

► (3) Spark Jobs

► test\_transformedLR: pyspark.sql.dataframe.DataFrame = [Avg\_Area\_Income: double, Avg\_Area\_House\_Age: double ... 8 more fields]

	ess	features	label	prediction
1	8 Terrance Pines	▶ {"vectorType": "dense", "length": 5, "values": [37971.20756623529, 4.291223903128535, 5.807509527238798, 3.24, 33267.767727560946]}	31140.517620186045	99409.91783725284

label	prediction
31140.517620186045	99409.91783725284
723750.0652577134	572142.5458968622
401148.5687913792	483203.61772650667
759044.6879907805	828305.733864381
1042814.0978200927	951102.6562018218

Prediction column added by LinearRegression



# Common data processing for modeling aka ML pipeline

- Begin with data
- Prepare a transformation, fit()
  - Model (estimator) customized to data
- Apply a transformation, transform()
  - Model applied to data to create DataFrame
- Prepare the next transformation, fit()
- Apply the next transformation, transform()
- ...

```
# RFormula must review the data (fit) to handle categorical
trained_RF_model = r_formula.fit(train_data) # create model
# Using the RFormula model, create a DataFrame of the transformed data
trained_model_DF = trained_RF_model.transform(train_data)
```

```
1 test_transformedLR = train_fittedLR.transform(test_preparedDF)
2 display(test_transformedLR)
```


# Check the model fit

- Recall, we applied the trained model to the test data

```
1 test_transformedLR = train_fittedLR.transform(test_preparedDF)
2 display(test_transformedLR)
```

```
1 labeledPredictions = test_transformedLR.select("label", "prediction")
2 labeledPredictions.show(5)
```


► (1) Spark Jobs

►  labeledPredictions: pyspark.sql.dataframe.DataFrame = [label: double, prediction: double]

label	prediction
31140.517620186045	99409.91783725284
723750.0652577134	572142.5458968622
401148.5687913792	483203.61772650667
759044.6879907805	828305.733864381
1042814.0978200927	951102.6562018218

```
5 y_test, predictions = zip(*labeledPredictions.collect())
6 fig, ax = plt.subplots()
7 plt.scatter(y_test, predictions)
8 display(fig)
```

► (1) Spark Jobs



# Important to remember

- PySpark regression is like Python regression
- RFormula must be fitted & transformed before using it in Regression
- Model parameters and data are placed into a single Features DataFrame column, using a DenseVector
- PySpark modeling assumes a pipeline
  - An estimator prepares model parameters, like RFormula
  - A transformation updates a DataFrame using a prepared model
  - Like Python's ML pipeline