

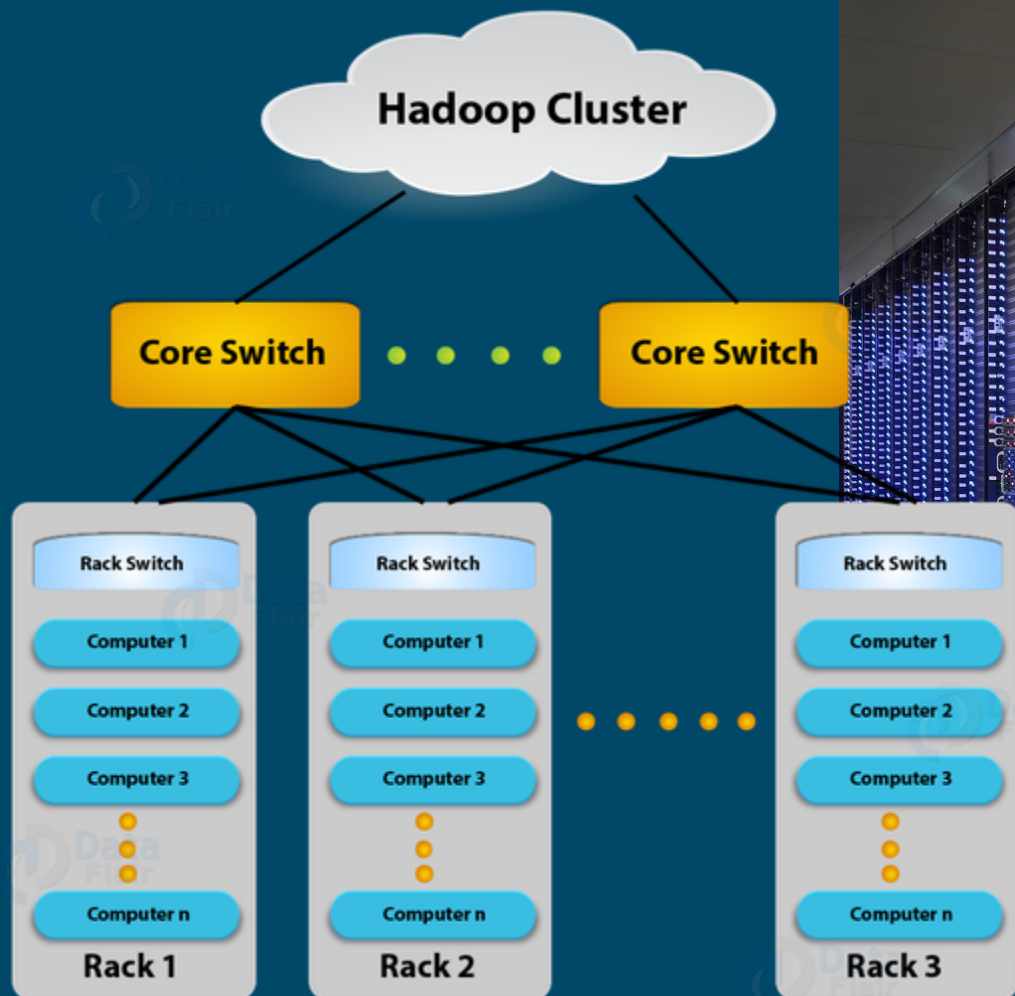


Cluster Networking Concepts

Bandwidth, latency, etc.




Computing cluster network




Bandwidth & throughput

- **Bandwidth** is the maximum amount of data that can be transmitted in a period. Measured as bits per second (bps)
- **Throughput** is the actual number of bits that flows through a network connection in a period
 - Can also measure job throughput—number of jobs processed per period
- Why a difference?
 - **Contention** is when two more processors attempt to access the same resource (e.g., network, disk drive, memory); called network **congestion** when due to heavy network usage
 - Resource processing limitation (CPU, RAM) of network devices


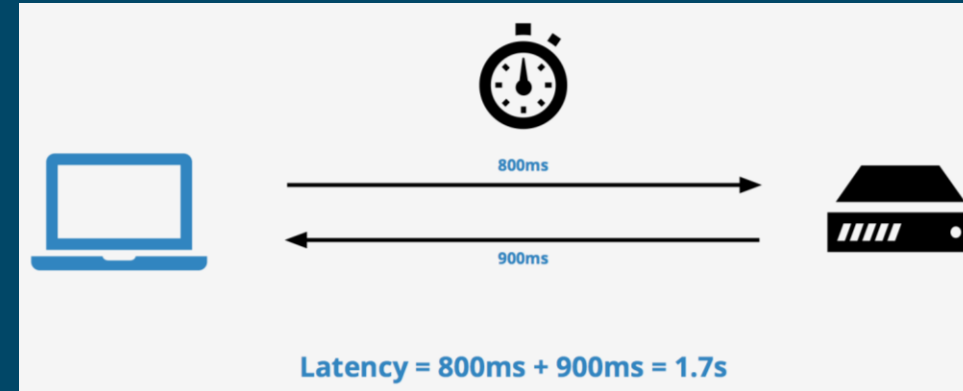


Cluster congestion often because of data transfer among nodes; for example, a large job




Latency & response time

- **Network latency** is the time it takes for a packet of data to get from one designated point to another
- **Response time** is the amount of time it takes to complete a job from the time it is submitted
 - Response time = 1 / throughput
- Both increase with cluster contention (network & processing)



Increasing the cluster speed

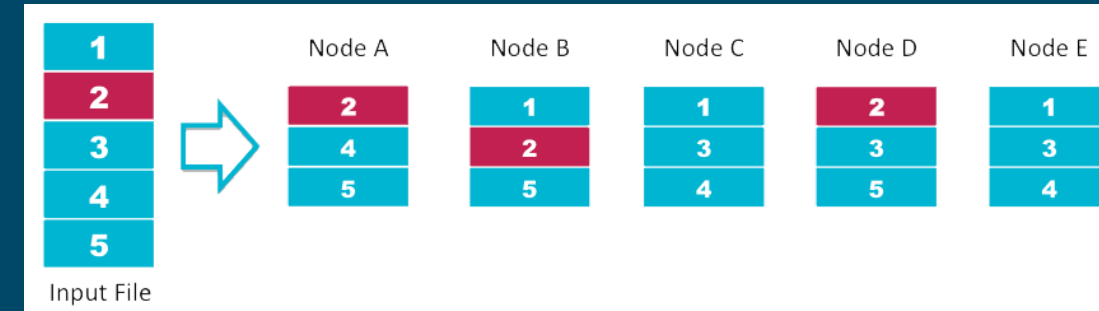
- Vertical Scaling (**scale-up**)
 - adding more processors and RAM, buying a more expensive and robust server
- Horizontal scaling (**scale out**)
 - adding more nodes to a system



Increasing the data processing speed

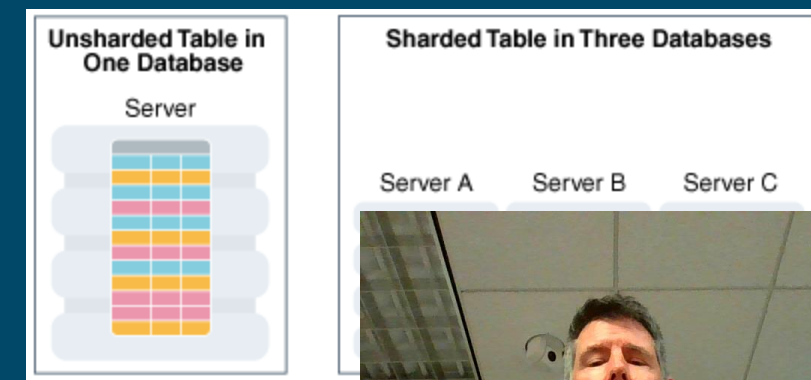
- **Data Replication**

- storing data in multiple nodes, improving availability, dependability,




- **Database shard**

- a partition of data (e.g., SQL rows), which may be stored on a node, improving availability, dependability, and processing time (through parallel processing)



Batch and Transaction Scaleup

- **Batch scaleup**
 - A single large job; typical of big data analysis and database
 - Use an N -times larger computer on N -times larger problem.
- **Transaction scaleup**
 - Numerous small queries submitted by independent users to a shared database; typical transaction processing and timesharing systems.
 - N -times as many users submitting requests (hence, N -times as many requests) to an N -times larger database, on an N -times larger computer.
- **Big Data processing (Spark / Hadoop) is better with batch processing**
 - Well-suited to parallel execution, but performs poorly with many clients and data update (and only recently addresses ACID transactions)
 - Note: batch process does allow for iterative algorithms, like K-means, which Spark (while Hadoop does not)



Factors limiting speedup

- Linear speed-up
 - Each added node decreases the processing time by one unit
- Speedup is often **sublinear** due to:
 - **Startup costs**
 - Node startup take time, which delays job computation
 - **Contention**
 - Nodes contend for shared resources (network, CPUs), and thus wait for processes to complete
 - **Skew**
 - Non-uniform data distribution (partitions) causes some tasks to complete before others. The job is complete when the slowest task ends



Hadoop overhead example on small query

- **75% Overhead!**

- 1 second to parse the query
- 9 seconds to submitting the query and launching ApplicationMaster (00:45:43 – 00:45:52)
- 6 seconds to initialize and launch the container for Map task (00:45:52 – 00:45:58)
- 3 seconds to initialize JVM (00:45:58 – 00:46:01)
- **6 seconds for actual Map and cleanup (00:46:01 – 00:46:07)**
 - $6/23.8 = 25\%$

```
00:45:43,041 - Parsing command: select 'A' from dual where 1=1
00:45:44,184 - Starting command: select 'A' from dual where 1=1
00:45:45,232 - Connecting to ResourceManager (by client)
00:45:48,459 - Submitted application
00:45:52,148 - Created MRAppMaster
00:45:55,742 - Connecting to ResourceManager (by AM)
00:45:58,184 - ContainerLauncher - CONTAINER_REMOTE_LAUNCH
00:45:58,246 - Transitioned from ASSIGNED to RUNNING
00:46:01,195 - JVM given task
00:46:04,181 - Progress of TaskAttempt is : 0.0
00:46:04,595 - Progress of TaskAttempt is : 1.0
00:46:04,677 - Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.85 sec
00:46:06,820 - Ended Job
Time taken: 23.8 seconds, Fetched: 1 row(s)
```



Parallel system performance measures

- **Speedup**

- $$\frac{\text{small system elapsed time}}{\text{large system elapsed time}} = \frac{100 \text{ min}}{50 \text{ min}} = 2X$$

linear


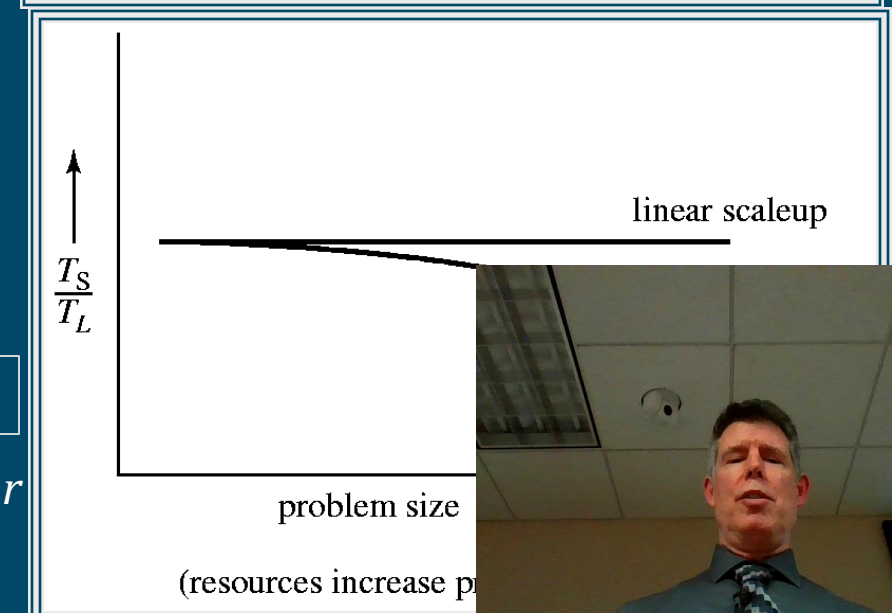

- **Scaleup**

- Do things “at scale”. Linear scaleup = 1
 - N-times larger system to perform n-times larger job

- $$\frac{\text{small system small problem elapsed time}}{\text{big system big problem elapsed time}}$$


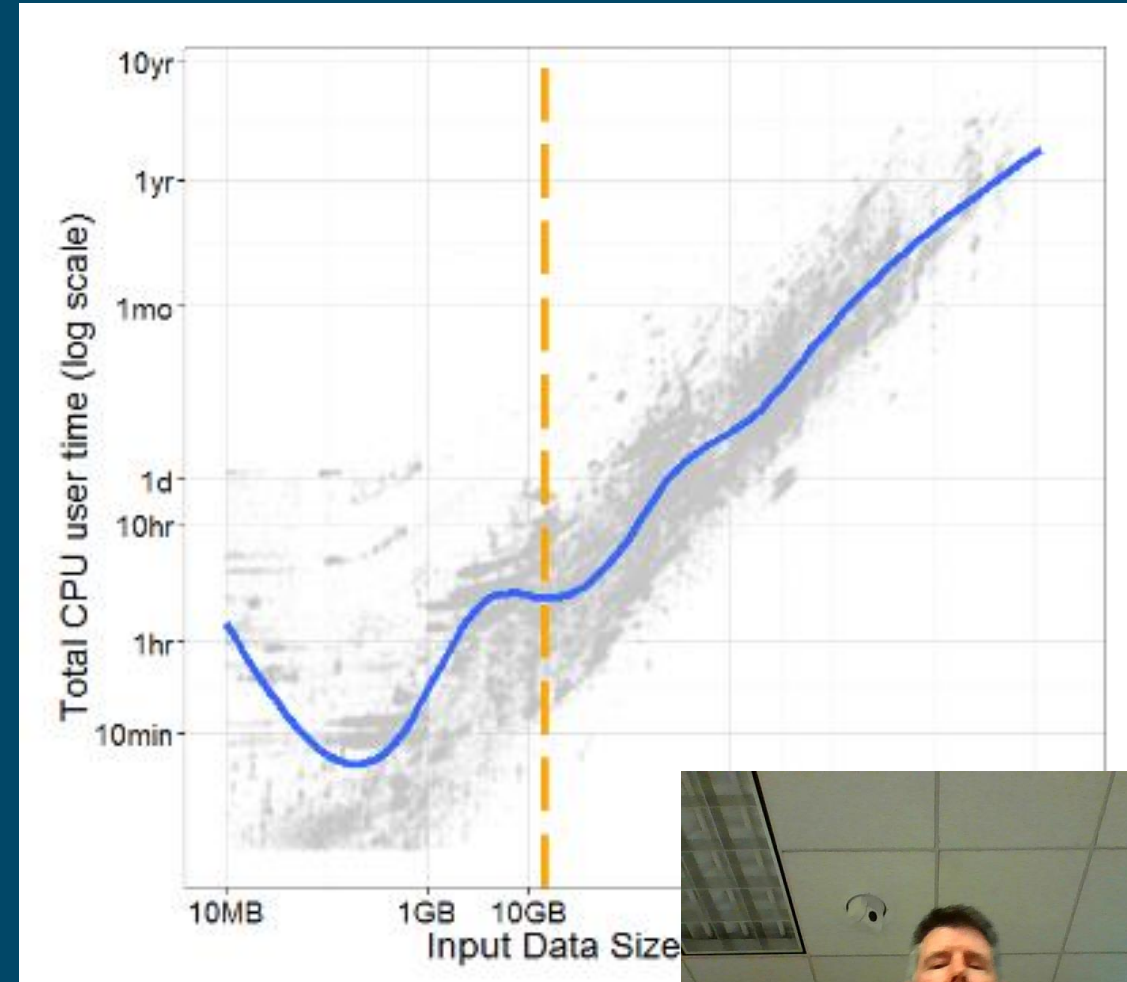
- $$\frac{10 \text{ CPU for 10 G data: } 10 \text{ min}}{100 \text{ CPU for 100G data: } 11 \text{ min}} = 0.91 \text{ scaleup}$$

sub-linear



Clusters don't provide linear improvement

- Cluster study
 - MapReduce 32 cores, 512 GB
 - Initially, as job size increases time decreases
 - Eventually, time increases with job size, with slope sub-linear



Important to remember

- Know networking concepts
 - Bandwidth, throughput, contention, latency, response time, scale up/out, shard, replication
- Factors limiting linear improvement
 - Start up, contention, skew
- Clusters don't provide linear improvement

