



Advance DataFrame Considerations

Operations and optimizations

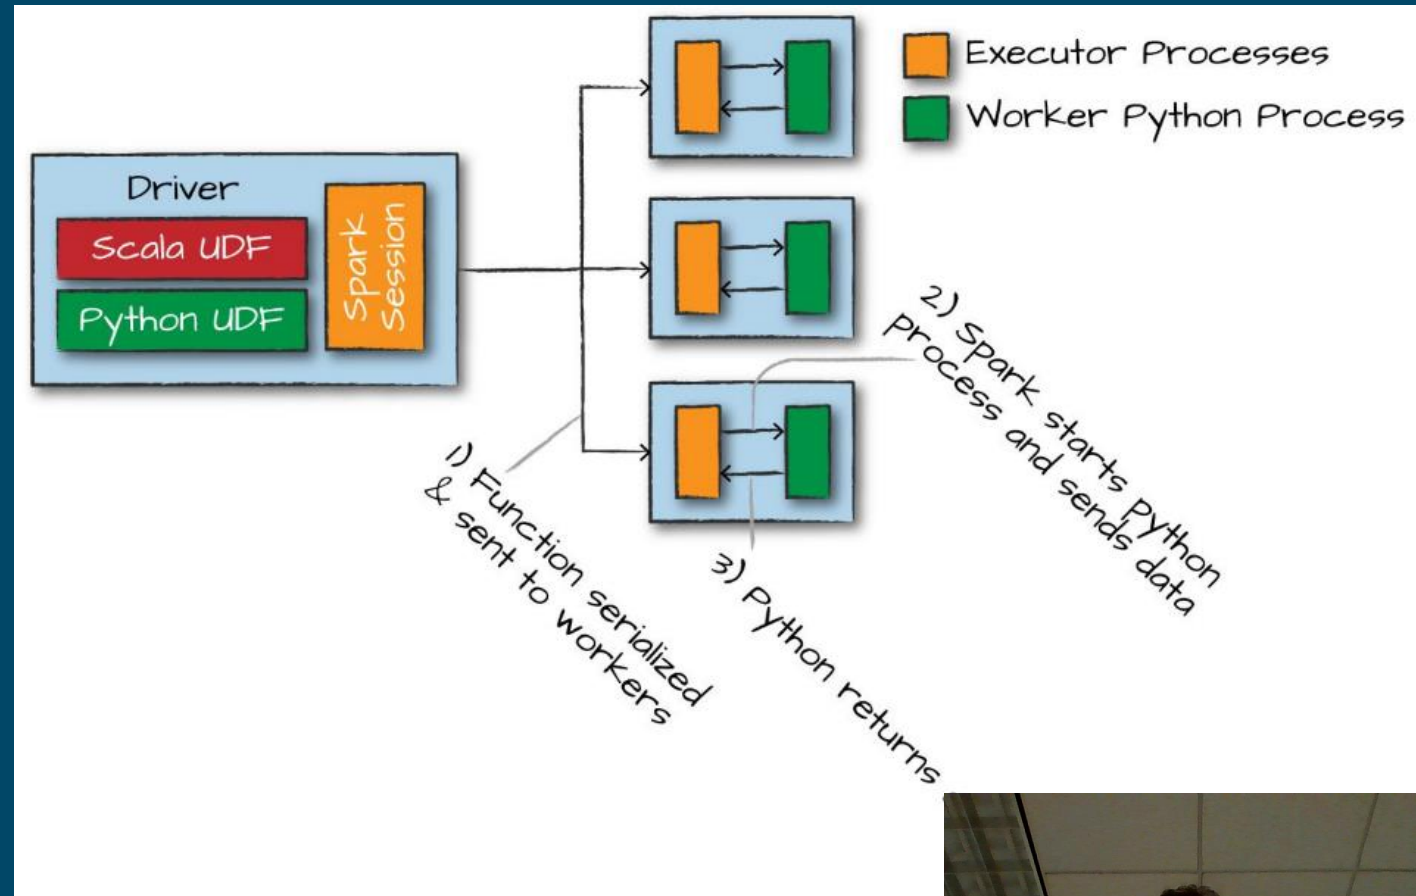


User Defined Function (UDF)

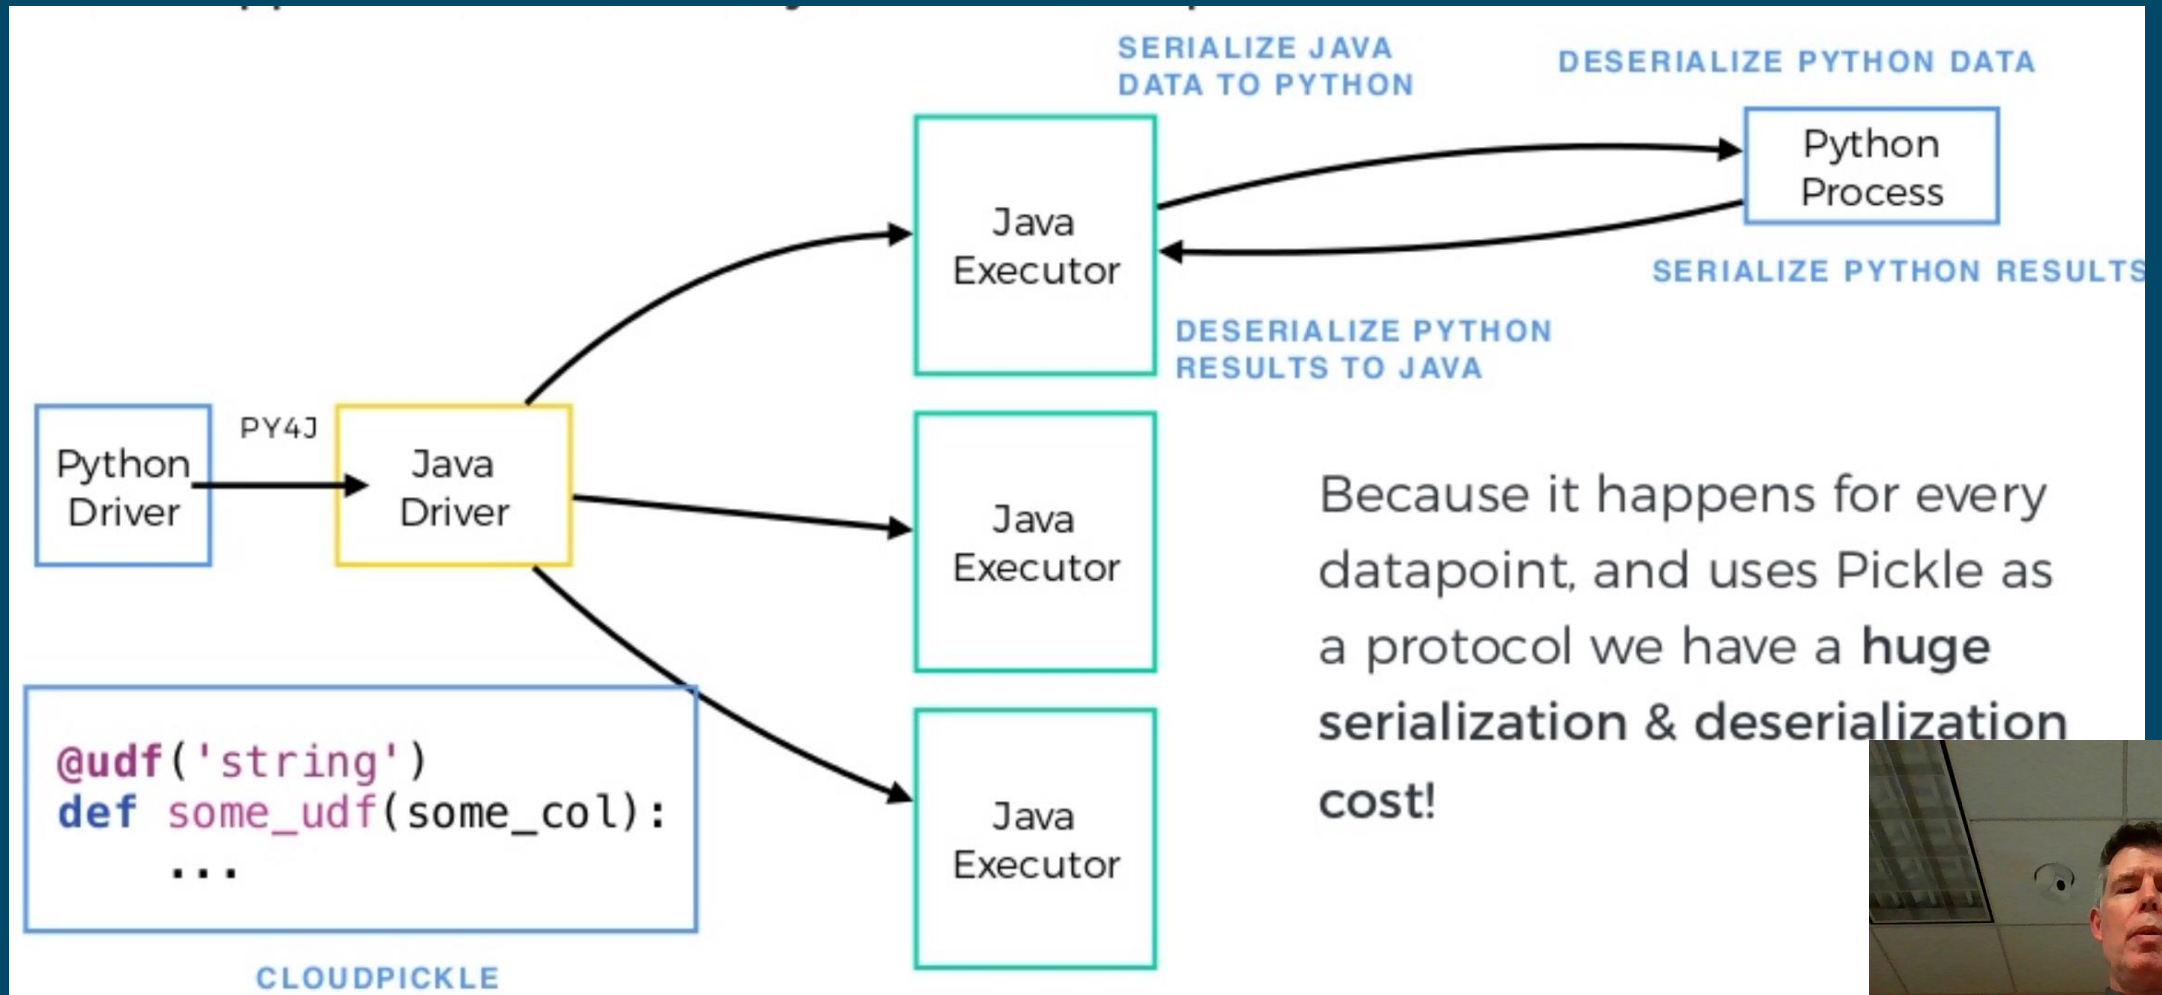


UDF

- UDF runs on each row of a DataFrame/SQL table
- Row is serialized & sent to worker (JVM) process
 - Added overhead of Python process for PySpark UDF
 - Process can fail because of memory, etc.
- For production systems, write UDF in Java or Scala



Python objects serialized and deserialized to Java (over pipe), which is slow (for UDFs)



UDF: User Defined Function is a procedure that runs over a DataFrame or SQL table. Calls to the UDF occur for **each row**. So many calls to the Py



Example PySpark native vs UDF

- Split text, by space, into words and put into vector

Native API split is fast

```
1 from pyspark.sql.functions import split,col
2 df1 = df.select(split(col("Description"), " ").alias("array_col"))
3 df1.show(1)
```

▶ (1) Spark Jobs

▶ df1: pyspark.sql.dataframe.DataFrame = [array_col: array]

```
+-----+
|      array_col|
+-----+
|[WHITE, HANGING, ...|
+-----+
only showing top 1 row
```


UDF split (mySplitUDF) is slow

```
1 from pyspark.sql.functions import udf,col
2
3 def splitUDF(col):
4     if col is not None:
5         return col.split(' ')
6
7 mySplitUDF = udf(splitUDF)
8
9 df2 = df.select(mySplitUDF(col('Description')))
10 display(df2)
```

▶ (1) Spark Jobs

▶ df2: pyspark.sql.dataframe.DataFrame = [splitUDF(Description): array]


	splitUDF(Description)
1	[WHITE, HANGING, HEART, T-LIGHT, HOLDER]



UDF try-catch

- Because PySpark UDF runs in its own process, good practice to catch error

```
3  def splitUDF(col):  
4      try:  
5          if col is not None:  
6              return(col.split(' '))  
7      except:  
8          return(None)
```



UDF for SQL

- UDF is for DataFrame
- Can be registered to work in SQL context

Register UDF in SQL context

```
1 from pyspark.sql.types import StringType
2 sqlContext.udf.register("splitUDF", splitUDF, StringType())
```


```
1 %sql
2 select *,splitUDF(Description)
3 from retail
```

► (1) Spark Jobs

	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	splitUDF(Description)
1	WHITE HANGING HEART T-LIGHT HOLDER	6	2010-12-01 08:26:00	2.55	17850	United Kingdom	[WHITE, HANG
2	WHITE METAL LANTERN	6	2010-12-01 08:26:00	3.39	17850	United Kingdom	[WHITE, META




Complex types in columns



DataFrame supports complex column types

- Array
 - Functions for size, split, and find in arrays
- Structure
 - Defined
 - with Struct
 - Read from JSON
 - Functions to expand into new columns (col.*)
- Maps
 - Functions to create a map from columns, explode a map to



DataFrame supports array column types


- Array
 - Functions to split, find in arrays

```
# in Python  
from pyspark.sql.functions import array_contains  
df.select(array_contains(split(col("Description"), " "), "WHITE")).show(2)  
  
-- in SQL  
SELECT array_contains(split(Description, ' '), 'WHITE') FROM dfTable
```

```
+-----+  
|array_contains(split(Description, ), WHITE)|  
+-----+  
|               true|  
|               true|  
+-----+
```

- Collect into a list

```
# in Python  
from pyspark.sql.functions import collect_set, collect_list  
df.agg(collect_set("Country"), collect_list("Country")).show()
```




DataFrame supports structure column types

- Structure
 - Defined
 - with Struct
 - Read from JSON
 - Functions to expand into new columns (col.*)

```
complexDF.select("complex.*")
```

```
-- in SQL
```

```
SELECT complex.* FROM complexDF
```



Expanding column structure with *

- Id and name attributes within Json in column

	department
1	▶ {"id": "123456", "name": "Computer Science"}
2	▶ {"id": "789012", "name": "Mechanical Engineering"}
3	▶ {"id": "345678", "name": "Theater and Drama"}
4	▶ {"id": "901234", "name": "Indoor Recreation"}

1 display(unionDF.select('department.*').limit(5))

▶ (3) Spark Jobs

	id	name
1	123456	Computer Science
2	789012	Mechanical Engineering
3	345678	Theater and Drama
4	901234	Indoor Recreation



Expanding embedded rows with explode

- Multiple employees in column

employees	
▶	[{"firstName": "michael", "lastName": "armbrust", "email": "no-reply@berkeley.edu", "salary": 100000}, {"firstName": "xiangrui", "lastName": "meng", "email": "no-reply@stanford.edu", "salary": 120000}]
▶	[{"firstName": "matei", "lastName": null, "email": "no-reply@waterloo.edu", "salary": 140000}, {"firstName": null, "lastName": "wendell", "email": "no-reply@berkeley.edu", "salary": 160000}]
▶	[{"firstName": "michael", "lastName": "armbrust", "email": "no-reply@berkeley.edu", "salary": 100000}, {"firstName": null, "lastName": "wendell", "email": "no-reply@berkeley.edu", "salary": 160000}]
▶	[{"firstName": "xiangrui", "lastName": "meng", "email": "no-reply@stanford.edu", "salary": 120000}, {"firstName": null, "lastName": null, "email": "no-reply@waterloo.edu", "salary": 140000}]

```
1 from pyspark.sql.functions import explode
2 exploded_df = unionDF.select(explode("employees").alias("e"))
3 display(exploded_df)
```

▶ (3) Spark Jobs


▶ exploded_df: pyspark.sql.dataframe.DataFrame = [e: struct]

	e
1	▶ {"firstName": "michael", "lastName": "armbrust", "email": "no-reply@berkeley.edu", "salary": 100000}
2	▶ {"firstName": "xiangrui", "lastName": "meng", "email": "no-reply@stanford.edu", "salary": 120000}
3	▶ {"firstName": "matei", "lastName": null, "email": "no-reply@waterloo.edu", "salary": 140000}
4	▶ {"firstName": null, "lastName": "wendell", "email": "no-reply@berkeley.edu", "salary": 160000}
5	▶ {"firstName": "michael", "lastName": "armbrust", "email": "no-reply@berkeley.edu", "salary": 100000}




Joins

Try to reduce the amount of data transmitted over network




Join types

- Inner joins
 - keep rows with keys that exist in the left and right datasets
- Outer joins
 - keep rows with keys in either the left or right datasets
- Left outer joins
 - keep rows with keys in the left dataset
- Right outer joins
 - keep rows with keys in the right dataset
- Left semi joins
 - keep the rows in the left, and only the left, dataset where the key appears in the right
- Left anti joins
 - keep the rows in the left, and only the left, dataset where they do not appear in the right



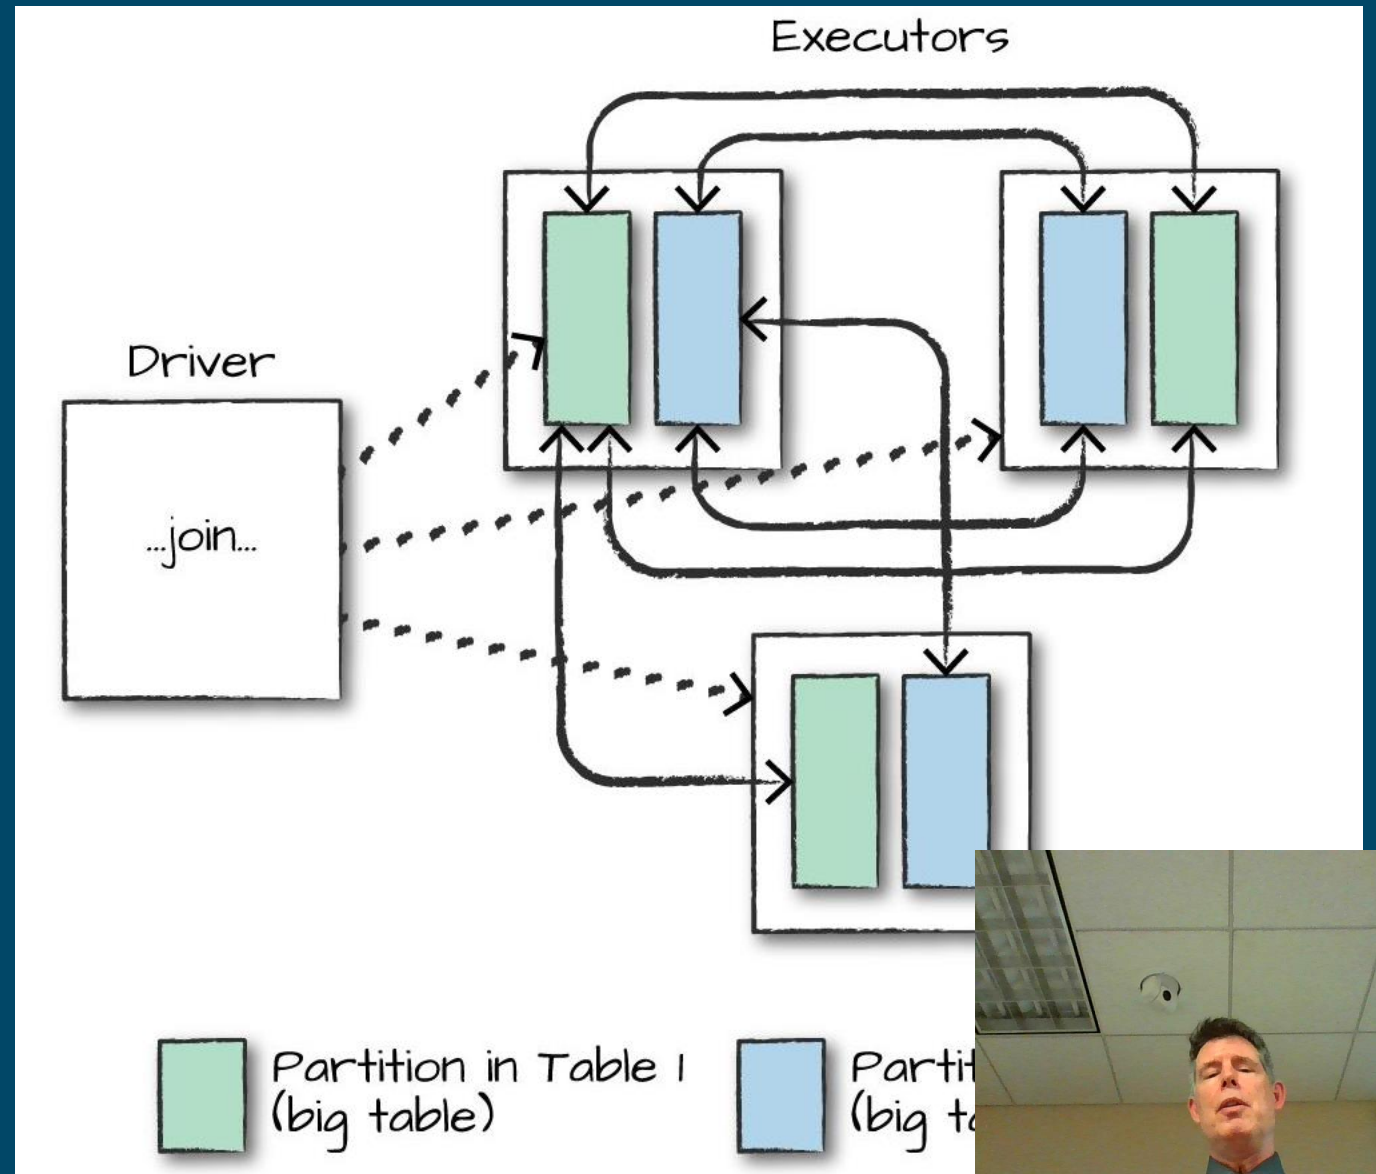
Joins are slow

- Data is distributed over cluster
- To join, each row from each table is compared
- So, each row has to be sent to each node
- Consequently, joins move lots of data over the network
- Network bandwidth becomes a bottleneck
- Improve join by transmitting less data
 - Send small tables to nodes with large tables
 - Filter data (push down predicated) before transmitting




Joining large tables transmits all data

- Each node receives data from all other nodes
- Each partition is sent to all other nodes



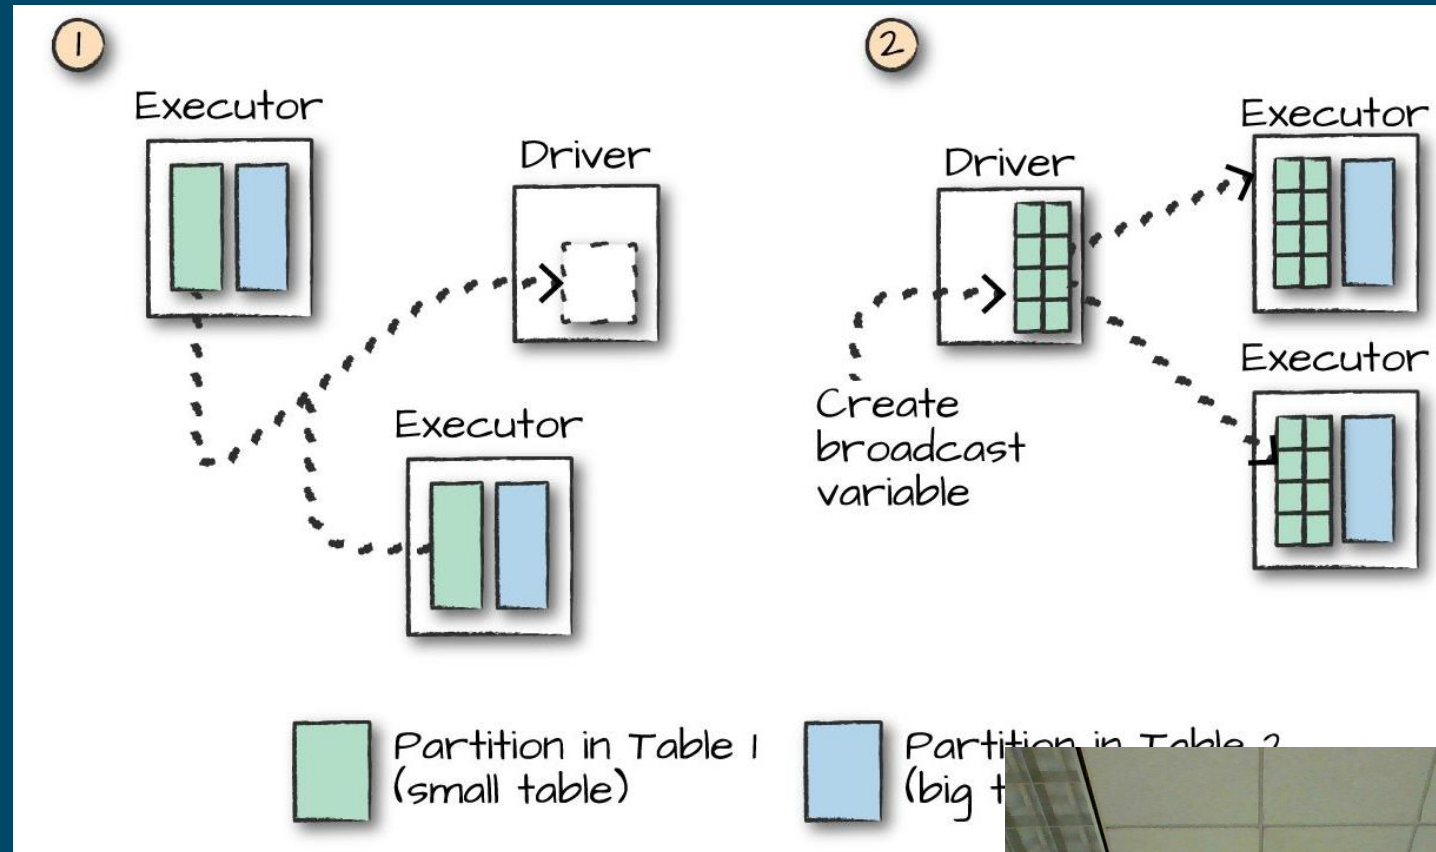
Optimize when there is a smaller table

- Broadcast the small table to improve the join
- When one of the join tables is small enough to fit into memory, then broadcast that table to the nodes
 - Reduces the total data transmitted
 - Can offer a hint to the Catalyst optimizer, but generally it optimizes OK



Broadcasting a table to join


1. Table is sent to Driver
 2. Special *broadcast variable* makes the data in the Driver accessible to all nodes
- Only transmitting the smaller table reduces data transmitted



```
df = blue.join(broadcast(green), blue.id =
```



Reading data and partitions



Partition the data during read for subsequent parallel processing


- Data may already be distributed over the nodes, where each node contains a partition
- Data that is not partitioned, can be partitioned during the read
 - From a file or SQL table

```
sqlContext.setConf("spark.sql.shuffle.partitions", "200")
```

in Python

```
dbDataFrame = spark.read.format("jdbc")\  
    .option("url", url).option("dbtable", tablename).option("driver", driver)\  
    .option("numPartitions", 10).load()
```

- Such partitions distribute the data over the nodes, allows for subsequent parallel processing



Important to remember

- PySpark UDF runs slowly in a separate Python process
 - Catch errors
- DataFrame support complex types
 - Arrays or structures as a column data type
 - Explode pulls multiple rows in a cell into rows
- DataFrame supports a variety of joins
 - Joins are slow because the tables must be copied to all nodes

