

# Resilient Distributed Datasets (RDDs)



# Resilient Distributed Dataset (RDD)

- represents an immutable, partitioned collection of records that can be operated on in parallel
  - list of objects
- the records are just Java, Scala, or Python objects of the programmer's choosing.
  - Unlike DataFrames though, where each record is a structured row containing fields with a known schema
- RDDs provide transformations, which evaluate lazily, and actions, which evaluate eagerly, to manipulate data in distributed fashion.



# Creating an RDD

- From local collection

```
myCollection = "Spark The Definitive Guide : Big Data Processing Made Simple"\  
               .split(" ")  
words = spark.sparkContext.parallelize(myCollection, 2)
```

- From data source

```
spark.sparkContext.wholeTextFiles("/some/path/withTextFiles")
```



# Some examples

```
1 # Map (apply a function to a list)
2 rdd = sc.parallelize(range(10), 5)
3 rdd.map(lambda x: x*2).collect()
```

► (1) Spark Jobs

```
Out[7]: [0, 2, 4, 6, 8, 10, 12, 14, 16, 18]
```

```
1 words = "I like data analytics with PySpark".split(" ") # List of words
2 wordsRDD = sc.parallelize(words)
3 wordsRDD.take(15)
```

► (3) Spark Jobs

```
Out[10]: ['I', 'like', 'data', 'analytics', 'with', 'PySpark']
```

```
1 # Add a suffix to each word
2 wordsRDD.map(lambda x: x + " um").take(15)
```

► (3) Spark Jobs

```
Out[11]: ['I um', 'like um', 'data um', 'analytics um', 'with um', 'PySpark um']
```



# Examples using reduce

```
1 xRDD = sc.parallelize(range(0,10))
2 xRDD.collect()
```

► (1) Spark Jobs

```
Out[15]: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
1 # Sum the list
2 # In the list, replace every pair of numbers with the sum of their numbers and repeat until the list has only one number.
3 xRDD.reduce(lambda x1,x2: x1+x2)
```

► (1) Spark Jobs

```
Out[16]: 45
```



Notice that RDD is the API for partitions

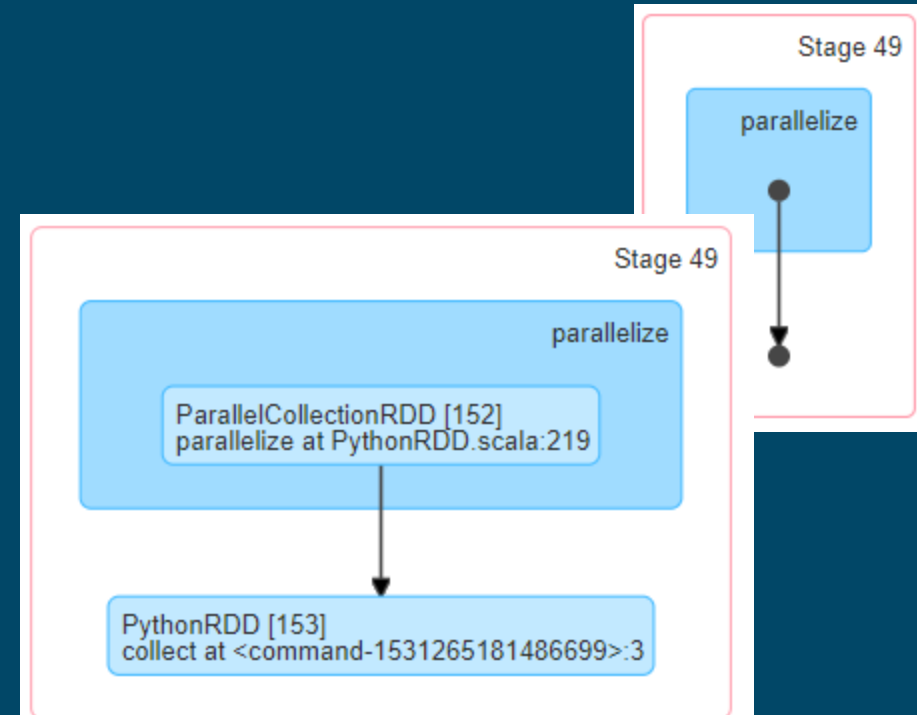


# 2 partitions -> 2 tasks

```
1 data = range(1,100)
2 rdd = sc.parallelize(data,2)
3 rdd.map(lambda x: x * x).collect()
4
```

## ▼ (1) Spark Jobs

► Job 36 [View](#) (Stages: 1/1)



## Tasks (2)

Index	ID	Attempt	Status	Locality Level	Executor ID	Host	Laun
0	206	0	SUCCESS	PROCESS_LOCAL	driver	localhost	2018
1	207	0	SUCCESS	PROCESS_LOCAL	driver	localhost	2018



# 5 partitions -> 5 tasks

```
1 data = range(1,100)
2 rdd = sc.parallelize(data,5)
3 rdd.map(lambda x: x * x).collect()
4
```


## ▼ (1) Spark Jobs

► Job 37 [View](#) (Stages: 1/1)

### Tasks (5)

Index	ID	Attempt	Status	Locality Level	Executor ID	Host	Launch Time	Duration
0	208	0	SUCCESS	PROCESS_LOCAL	driver	localhost	2018/11/02 16:58:40	0.2 s
1	209	0	SUCCESS	PROCESS_LOCAL	driver	localhost	2018/11/02 16:58:40	0.2 s
2	210	0	SUCCESS	PROCESS_LOCAL	driver	localhost	2018/11/02 16:58:40	0.3 s
3	211	0	SUCCESS	PROCESS_LOCAL	driver	localhost	2018/11/02 16:58:40	0.2 s
4	212	0	SUCCESS	PROCESS_LOCAL	driver	localhost	2018/11/02 16:58:40	0.2 s

Stage 50





# RDD API functions



# Some transformations


- Filter

```
def startsWithS(individual):  
    return individual.startswith("S")
```

```
words.filter(lambda word: startsWithS(word)).collect()
```

- Map

```
words2 = words.map(lambda word: (word, word[0], word.startswith(
```



# Some actions

- Count

```
words.count()
```

```
words.cache()
```


- Cache

- Subsequent access to words relies on temporary stored data rather than rerunning the computation



# RDD Transformations and Actions

<b>Transformations</b>	$map(f : T \Rightarrow U) : RDD[T] \Rightarrow RDD[U]$ $filter(f : T \Rightarrow Bool) : RDD[T] \Rightarrow RDD[T]$ $flatMap(f : T \Rightarrow Seq[U]) : RDD[T] \Rightarrow RDD[U]$ $sample(fraction : Float) : RDD[T] \Rightarrow RDD[T]$ (Deterministic sampling) $groupByKey() : RDD[(K, V)] \Rightarrow RDD[(K, Seq[V])]$ $reduceByKey(f : (V, V) \Rightarrow V) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $union() : (RDD[T], RDD[T]) \Rightarrow RDD[T]$ $join() : (RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (V, W))]$ $cogroup() : (RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (Seq[V], Seq[W]))]$ $crossProduct() : (RDD[T], RDD[U]) \Rightarrow RDD[(T, U)]$ $mapValues(f : V \Rightarrow W) : RDD[(K, V)] \Rightarrow RDD[(K, W)]$ (Preserves partitioning) $sort(c : Comparator[K]) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $partitionBy(p : Partitioner[K]) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$
<b>Actions</b>	$count() : RDD[T] \Rightarrow Long$ $collect() : RDD[T] \Rightarrow Seq[T]$ $reduce(f : (T, T) \Rightarrow T) : RDD[T] \Rightarrow T$ $lookup(k : K) : RDD[(K, V)] \Rightarrow Seq[V]$ (On hash/range partitioned RDDs) $save(path : String) : \text{Outputs RDD to a storage system, e.g., HDFS}$



# Key value pairs

- Many RDD are processed as key-value pairs

- Returns pair (word,1)

- Word as key, with 1 as value

```
words.map(lambda word: (word.lower(), 1))
```

- Can be used to count up word occurrences

- Map over pairs

```
keyword.mapValues(lambda word: word.upper()).collect()
```

```
[('s', 'SPARK'),  
 ('t', 'THE'),  
 ('d', 'DEFINITIVE'),  
 ('g', 'GUIDE'),  
 (':', ':'),  
 ('b', 'BIG'),  
 ('d', 'DATA'),  
 ('p', 'PROCESSING'),
```



# Count letter occurrence from words

- Collection of pairs (letter, 1)

```
chars = words.flatMap(lambda word: word.lower())  
KVcharacters = chars.map(lambda letter: (letter, 1))
```

- groupByKey

- First, sort letters into groups by their key (letter)
- Then, using reduce, add up the occurrences

```
KVcharacters.groupByKey().map(lambda row: (row[0], reduce(addFunc, row[1])))\  
.collect()
```


- reduceByKey

- Add occurrences that have the same key

```
KVcharacters.reduceByKey(addFunc).collect()
```



# Class word-count program






# Classic word-count with map-reduce

- Originally used to index web pages (by Google, Yahoo!)
  - Create a map of (word,1) pairs
  - Add up all numbers with for the same word

```
1 words = "I like data analytics with PySpark. It's data analytics are like Python.".split(" ") # List of words
2 wordsRDD = sc.parallelize(words)
3 words_KV_RDD = wordsRDD.map(lambda w: (w.lower(), 1))
4
5 print("Each word (key) with an associated 1: {}".format(words_KV_RDD.collect()))
6
7 word_count_RDD = words_KV_RDD.reduceByKey(lambda left,right: left+right)
8 print("Now, reduce those pairs by summing the 1's when the keys match:\n {}".format(word_count_RDD.collect()))
9 # Notice the '2' values in the output
```

► (2) Spark Jobs

```
Each word (key) with an associated 1: [('i', 1), ('like', 1), ('data', 1), ('analytics', 1), ('with', 1), ('pyspark.', 1), ('it's', 1), ('are', 1), ('like', 1), ('python.', 1)]
Now, reduce those pairs by summing the 1's when the keys match:
[('i', 1), ('like', 2), ('are', 1), ('it's', 1), ('python.', 1), ('data', 2), ('analytics', 2), ('pyspark.', 1), ('with', 1)]
```






# word-count with groupByKey

- Slower than reduceByKey
  - Create a map of (word,1) pairs
  - Group all words
  - Add numbers in group

```
1 # For every pair (key,valueList), return the key and the sum of the valueList
2 words_KV_RDD.groupByKey().map(lambda key_value: (key_value[0],sum(key_value[1]))).collect()
```

► (1) Spark Jobs

```
Out[4]: [('i', 1),
('like', 2),
('are', 1),
("it's", 1),
('python.', 1),
('data', 2),
('analytics', 2),
('pyspark.', 1),
('with', 1)]
```

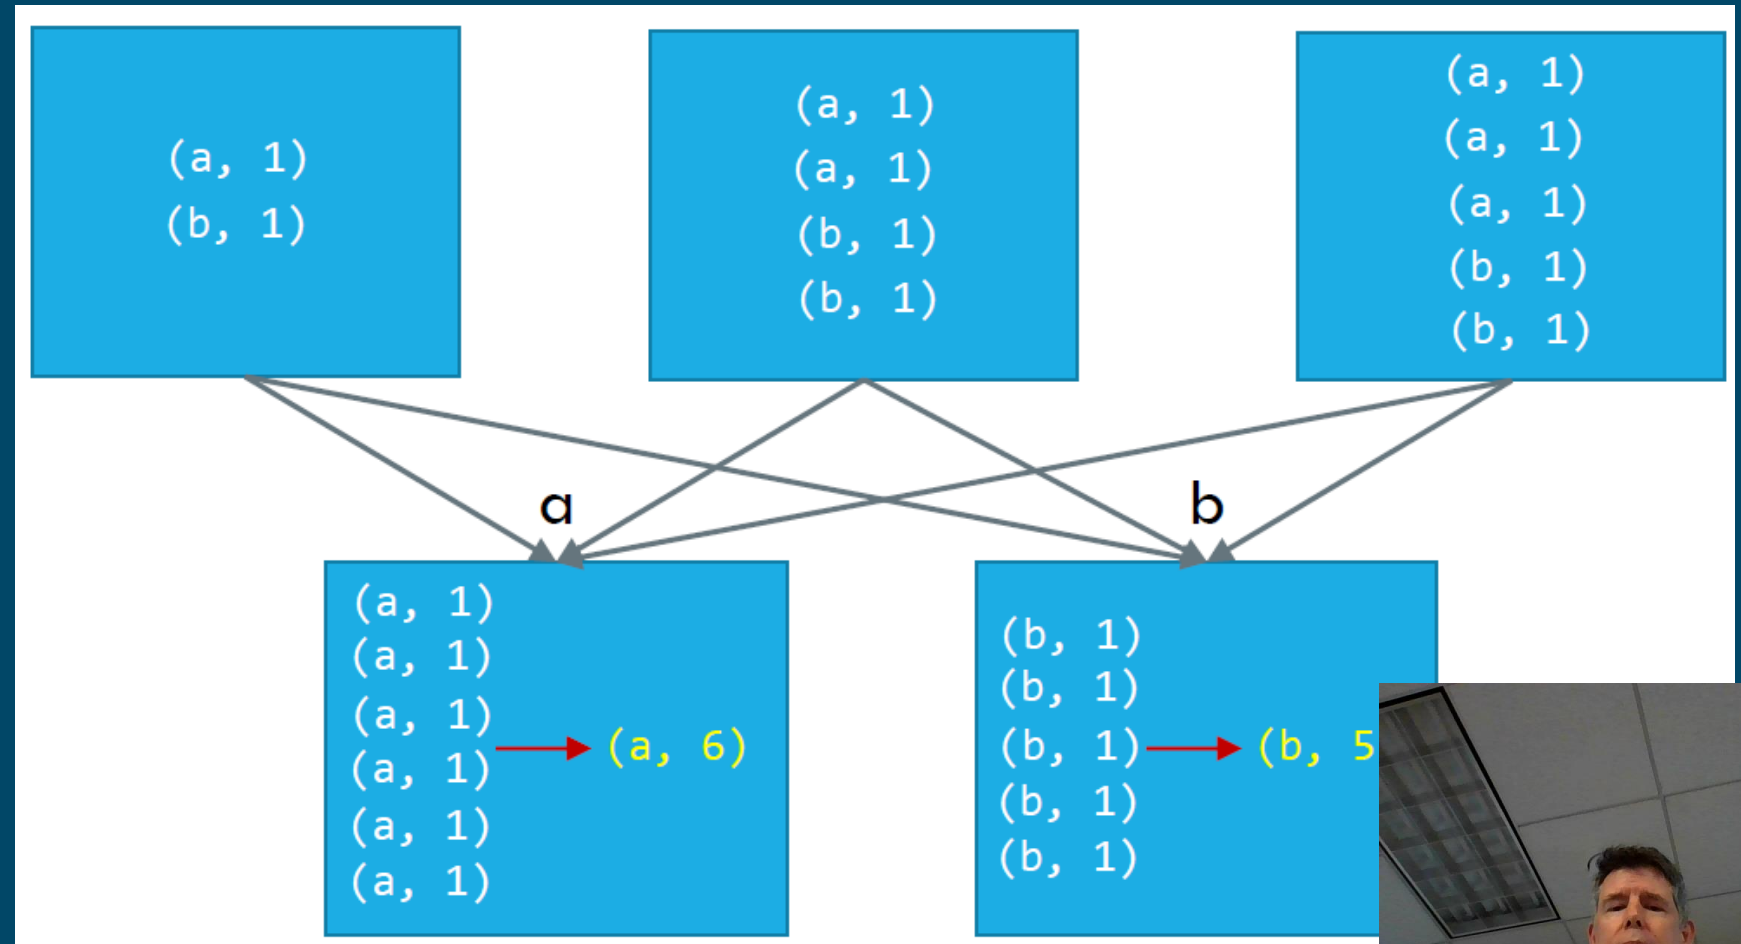


# Group By Key

All the key-value pairs are shuffled around, to obtain the groups  
**More data transfer**

Note data transferred...

- All simple  $\langle K, V \rangle$
- Sum only run in second node

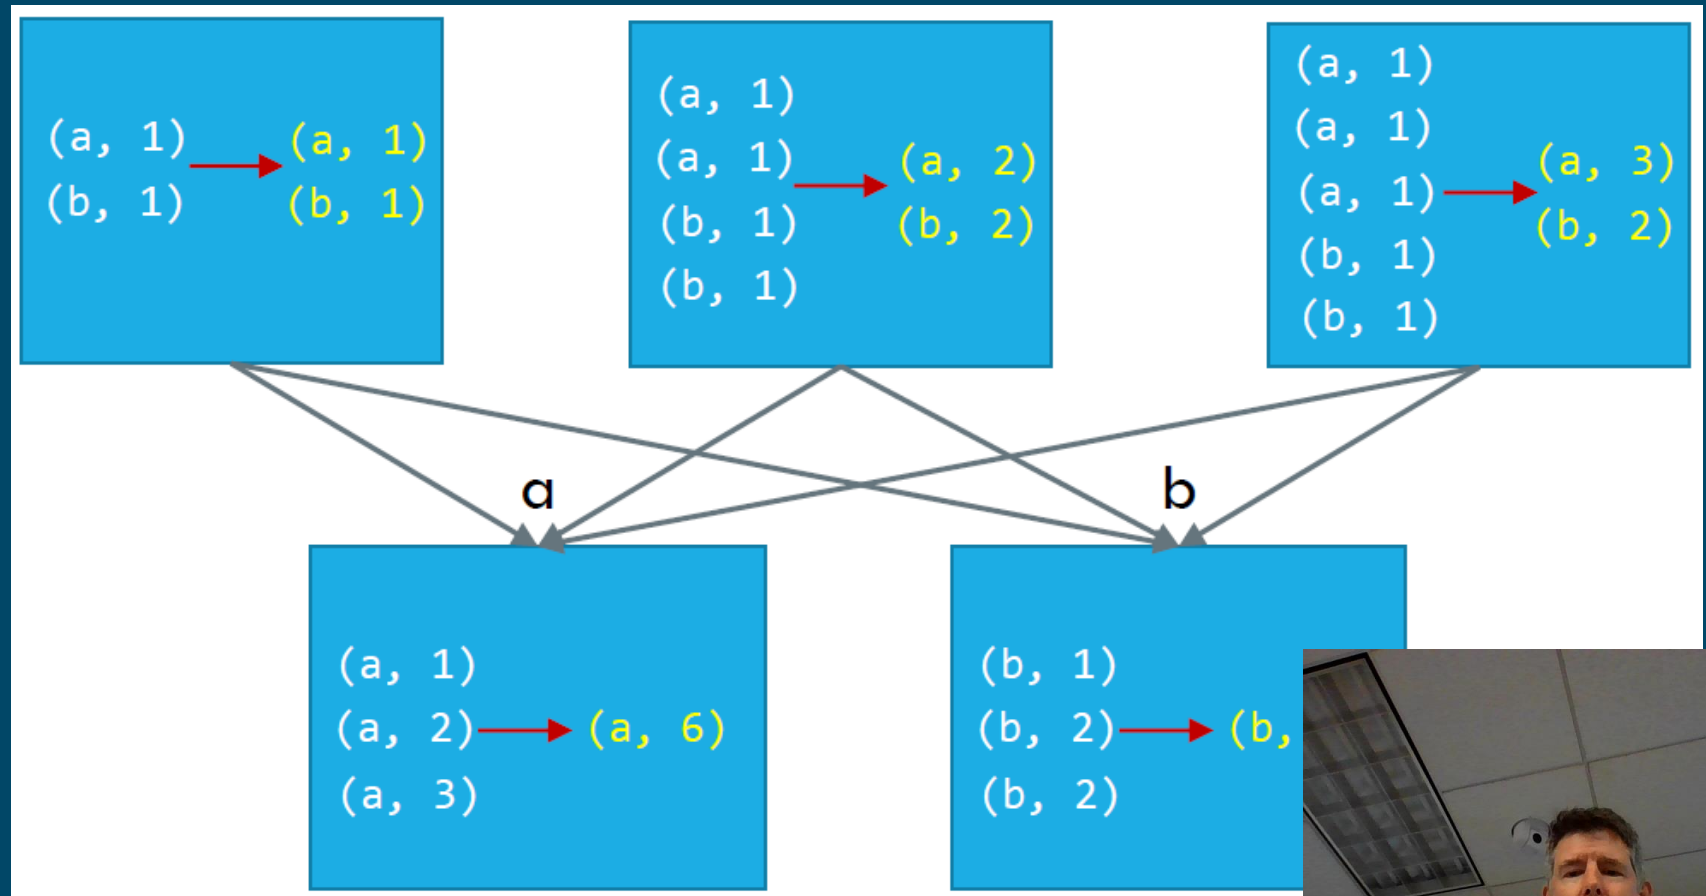


# Reduce By Key

Notice how pairs (e.g.,  $\langle a, 3 \rangle$ ) on the same machine with the **same key** are combined...  
So **less data transfer**

Note data transferred...

- Only  $\langle K, V \rangle$
- Combiner (sum) run in first node
- Sum also run in second node



# Prefer reduceByKey over groupByKey

- Less data transit with reduceByKey
- In general, be aware of how each transformation works
  - Notice how much data is being moved
- Goal is to reduce data movement
  - Do computations local to the node that has the data
  - Don't move the data to perform computations (unless necessary)

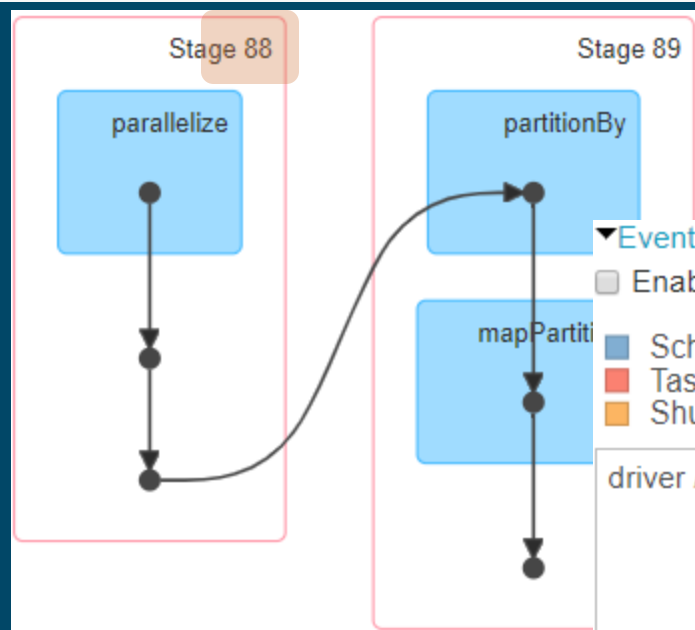


# reduceByKey with 5 tasks

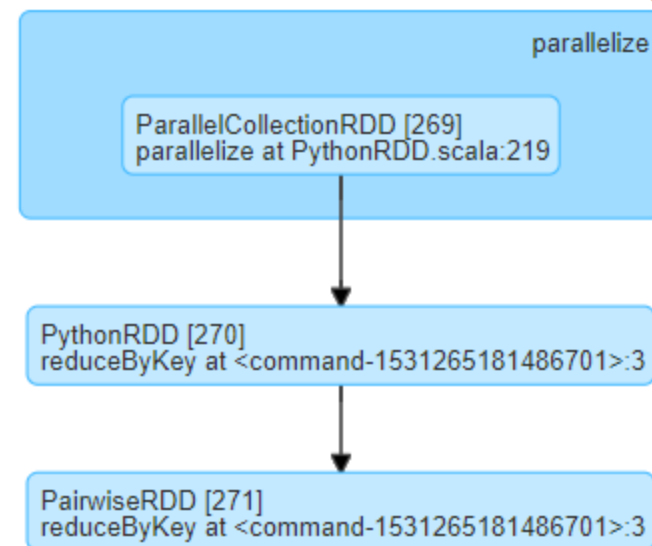
```
1 data = range(1,100)
2 rdd = sc.parallelize(data,5)
3 rdd.map(lambda x: (x%5, x)).reduceByKey(lambda x,y: x + y).collect()
```

## ▼ (1) Spark Jobs

► Job 59 [View](#) (Stages: 2/2)




Shuffle shown @ end




## ▼ Event Timeline

Enable zooming

Scheduler Delay  
Task Deserialization Time  
Shuffle Read Time  
Executor Computing Time  
Shuffle Write Time  
Getting Result Time  
Result Serialization Time



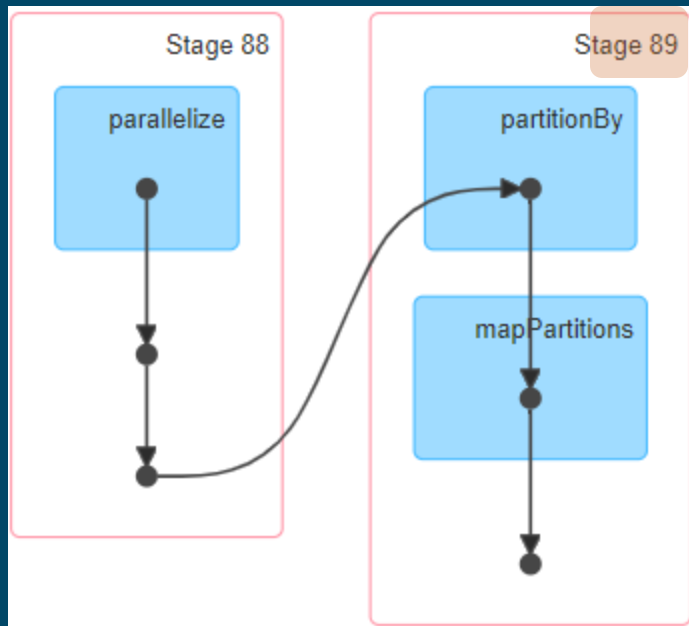
5 Tasks



# reduceByKey with 5 tasks

```
1 data = range(1,100)
2 rdd = sc.parallelize(data,5)
3 rdd.map(lambda x: (x%5, x)).reduceByKey(lambda x,y: x + y).collect()
```

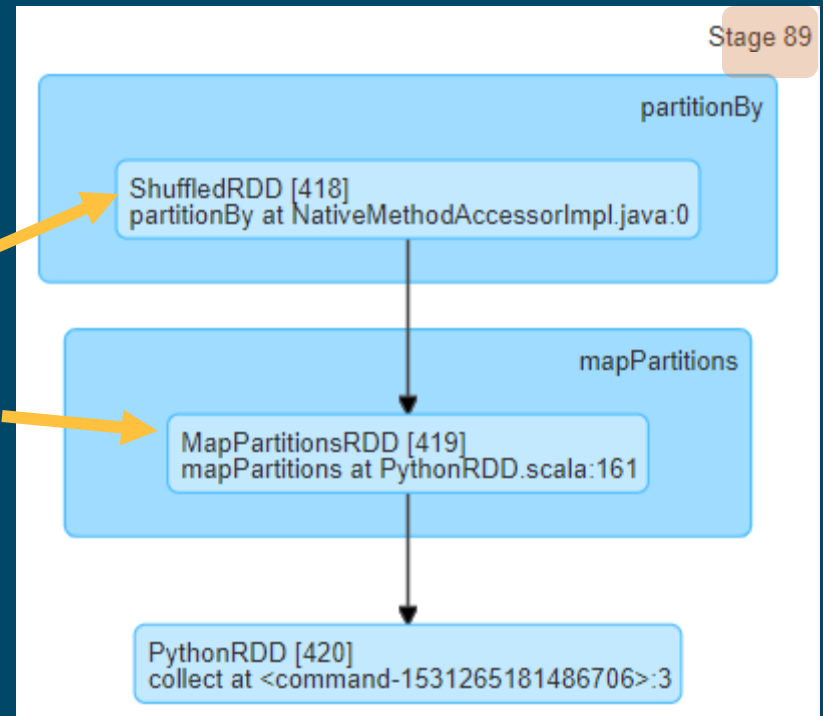
▼ (1) Spark Jobs  
► Job 59 View (Stages: 2/2)



5 Tasks

Shuffled

Reduce




le zooming

eduler Delay  
c Deserialization Time  
ffle Read Time

Executor Computing Time  
Shuffle Write Time  
Result Serialization Time

Getting Result Time



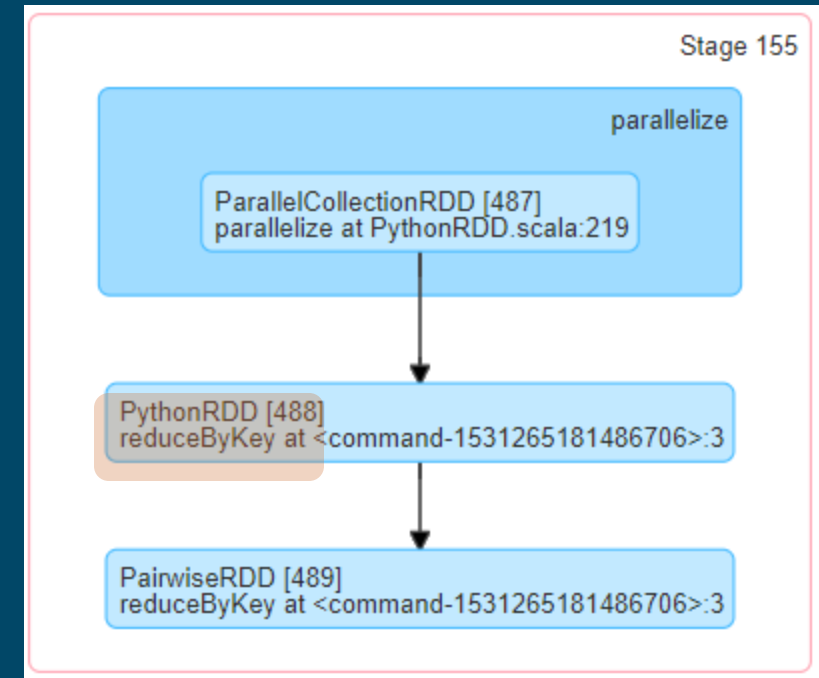
# Map Reduce with 5 tasks

Stage Id ▾	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total
156	8261552548926543378	data = range(1,100) rdd = sc.parallelize(data,5... collect at <command-1531265181486706>:3 +details	2018/11/02 21:09:34	50 ms	5/5
155	8261552548926543378	data = range(1,100) rdd = sc.parallelize(data,5... reduceByKey at <command- 1531265181486706>:3 +details	2018/11/02 21:09:33	0.3 s	






# reduceByKey with 5 tasks



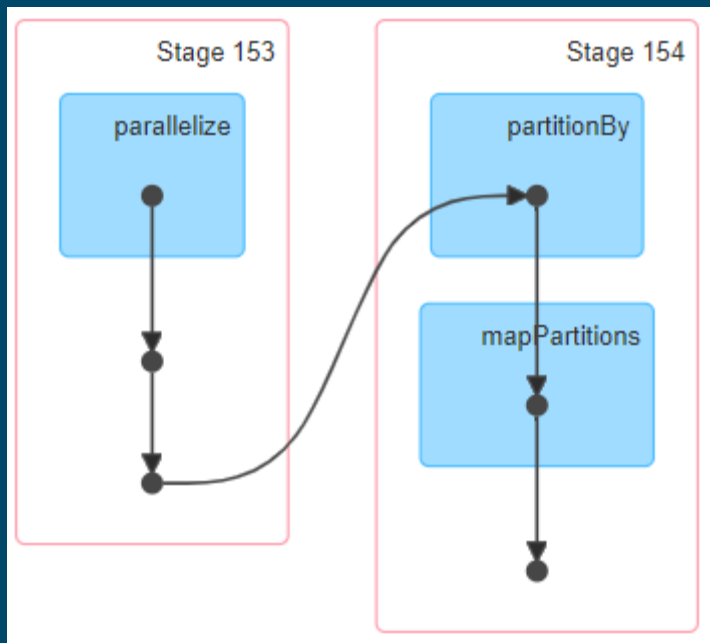
Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0.2 s	0.2 s	0.2 s	0.2 s	0.3 s
GC Time	0 ms	0 ms	0 ms	0 ms	0 ms
Shuffle Write Size / Records	305.0 B / 5	305.0 B / 5	305.0 B / 5	309.0 B / 5	310.0 B / 5





# groupByKey with 5 tasks, using

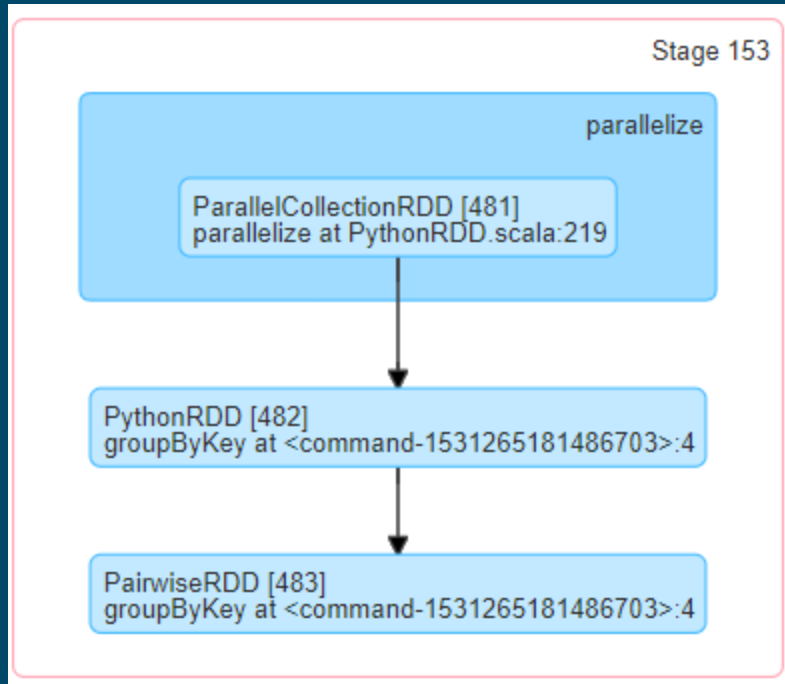
```
1 data = range(1,100)
2 rdd = sc.parallelize(data,5)
3 rdd.map(lambda x: (x%5, x)).groupByKey().map(lambda x: (x[0], sum(list(x[1])))).collect()
```



Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total
154	8261552548926543378	data = range(1,100) rdd = sc.parallelize(data,5... collect at <command-1531265181486703>:4 <a href="#">+details</a>	2018/11/21:07:37		
153	8261552548926543378	data = range(1,100) rdd = sc.parallelize(data,5... groupByKey at <command-1531265181486703>:4 <a href="#">+details</a>	2018/11/21:07:37		



# groupByKey with 5 tasks




Slower than  
reduceByKey

Metric	Min	25th percentile	Median	75th percentile
Duration	0.2 s	0.2 s	0.2 s	0.2 s
GC Time	0 ms	0 ms	0 ms	0 ms
Shuffle Write Size / Records	358.0 B / 5	358.0 B / 5	358.0 B / 5	358.0 B / 5



# reduceByKey shuffles smaller data set


- reduceByKey
  - Duration 50 ms
  - Shuffle min: 305 ms
- groupByKey
  - Duration 61 ms
  - Shuffle min: 358 ms



To reduce processing time,

**Minimize the number of stages**  
(wide transformation, shuffle & sort)

a key to good Spark programming



# Important to remember

- RDDs are the primitive structure that supports DataFrames
  - RDD data functions are simpler
  - RDD allows manipulation of the partitions
- When possible, minimize data transmission
  - Prefer `reduceByKey` over `groupByKey`
    - `reduce` aggregates local keys, and then send to combine results
    - `group by` sends all data to create groups, then aggregates

