

Labs illustrating PySpark Development

Common Spark development techniques

Spark demonstrations

- PySpark on your computer
 - Run from shell (terminal)
 - Run within IDE (PyCharm)
- PySpark on managed cluster
 - Databricks
 - Google Dataproc
- PySpark in Docker on Kubernetes cluster
 - On your computer
 - On Google Kubernetes engine (GKE)

Why on your local computer?

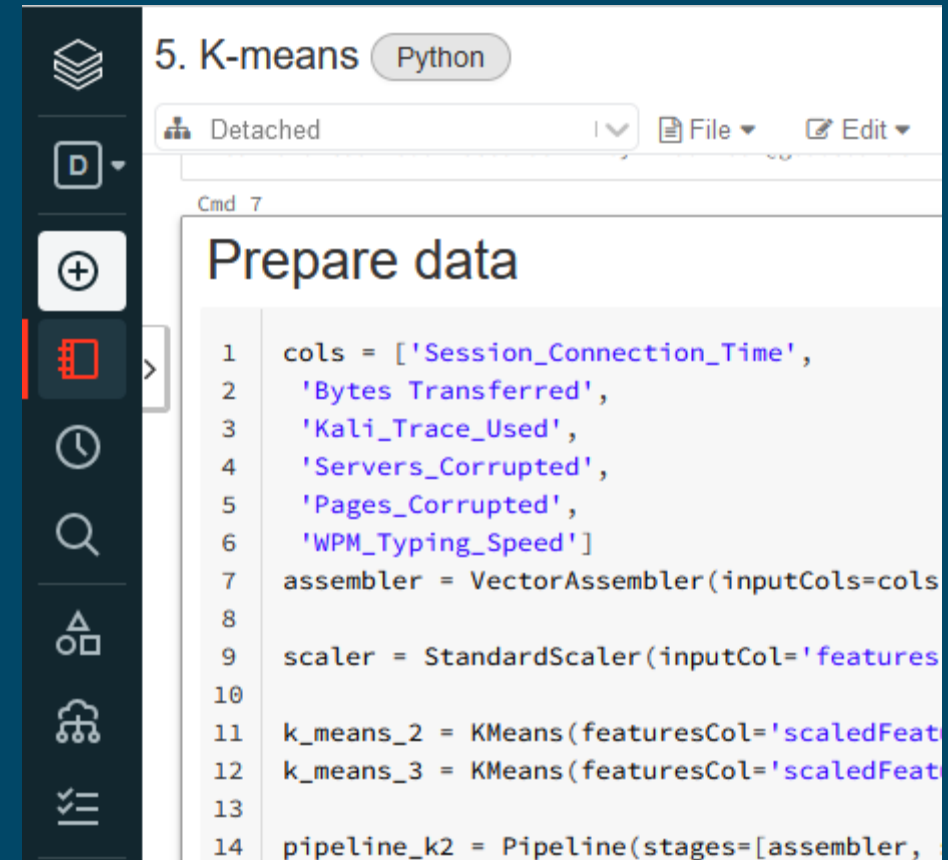
- Typical development life-cycle
 - Create & test with small data on local computer before moving to cluster
- No need for internet resources
 - Test on airplane
 - No expenses
- No need for remote (complex) debugging tools
 - Can use debugger to step through code, including Spark internals
 - Easier for debugging and understanding how it works

Management of Spark cluster

- Managed Spark cluster
 - User requests access to a cluster, perhaps specifying the general size information (e.g., number of nodes)
 - Examples: Databricks, Google Dataproc
- Unmanaged Spark cluster
 - User specifies details of cluster manager and worker nodes, including CPU, memory, OS version, package installations, network configuration, etc.
 - Example: Google Kubernetes Engine or Compute Engine
- Level (higher-level Spark application or lower-level Kubernetes containers)
 - GKE manages Kubernetes, which can run Spark (which is unmanaged)
 - So, GKE automatically manages Kubernetes (managed at the Kubernetes level), but a user must specify (manage) the images and their architecture
- Mostly unmanaged example
 - Use Compute Engine to install Kubernetes (which user must manage)
 - On the installed Kubernetes, run Spark images (which user must manage)
 - Note that Google still manages the virtual computers, which is yet a lower level

Why on managed cluster?

- Running Spark on **managed cluster is much simpler to use**, costs more, and leaves some architectural considerations to the managing software
 - Databricks
 - Google Dataproc



The screenshot shows a Databricks notebook interface. The title bar indicates the notebook is named '5. K-means' and is written in Python. The status is 'Detached'. The left sidebar contains icons for workspace, recent clusters, search, and other functions. The main area shows a code editor with the following Python code:

```
1 cols = ['Session_Connection_Time',  
2         'Bytes_Transferred',  
3         'Kali_Trace_Used',  
4         'Servers_Corrupted',  
5         'Pages_Corrupted',  
6         'WPM_Typing_Speed']  
7 assembler = VectorAssembler(inputCols=cols,  
8                               outputCol='features')  
9 scaler = StandardScaler(inputCol='features',  
10                          outputCol='scaledFeatures')  
11 k_means_2 = KMeans(featuresCol='scaledFeatures',  
12                    numClusters=2, seed=1)  
13 k_means_3 = KMeans(featuresCol='scaledFeatures',  
14                    numClusters=3, seed=1)  
15 pipeline_k2 = Pipeline(stages=[assembler, scaler, k_means_2])  
16 pipeline_k3 = Pipeline(stages=[assembler, scaler, k_means_3])
```

Why on unmanaged cluster?

- Running Spark on **unmanaged cluster is more complex**, costs less, and leaves architectural considerations to the user
 - Google Kubernetes Engine
 - Compute Engine

```
SPARK_CMD="$SPARK_HOME/bin/spark-submit \  
--master ${MASTER} \  
--deploy-mode cluster \  
--driver-memory ${DRIVER_MEMORY} \  
--executor-memory ${EXECUTOR_MEMORY} \  
--conf spark.executor.instances=${EXECUTORS} \  
--conf spark.dynamicAllocation.enabled=false \  
--conf spark.executor.heartbeatInterval=20s \  
--conf spark.shuffle.io.retryWait=60s \  
--conf spark.sql.shuffle.partitions=1000 \  
--conf spark.executor.cores=${EXECUTOR_CORES} \  
--conf spark.kubernetes.container.image=${IMAGE} \  
--conf spark.kubernetes.container.image.pullPolicy=Always \  
--name ${PYSPARK_APP_MODULE} \  
--conf spark.kubernetes.authenticate.driver.serviceAccountName=spark \  
--conf spark.kubernetes.executor.label.app=${PYSPARK_APP_MODULE} \  
--conf spark.kubernetes.driver.label.app=${PYSPARK_APP_MODULE} \  
--conf spark.kubernetes.driverEnv.PYSPARK_APP_MODULE=${PYSPARK_APP_MODULE} \  
--conf spark.kubernetes.driverEnv.PYTHONPATH=${WORKING_DIR}/${PYSPARK_APP_MODULE}.zip \  
--conf spark.kubernetes.executorEnv.PYTHONPATH=${WORKING_DIR}/${PYSPARK_APP_MODULE}.zip \  
${SCRIPT} \  
--master=${MASTER} --py_files=${WORKING_DIR}/${PYSPARK_APP_MODULE}.zip"
```

Illustrative problem for Spark developer

- Program requires other software to run
 - Python modules, Java jars, etc.
- Program & modules must be distributed over cluster
 - PySpark .py file or Java jar sent
 - Dependencies are a problem, because many (large) files may be required
- Managed cluster
 - First, install dependencies on cluster nodes (large files copied)
 - Then, distribute the program (--py_files)
- Unmanaged cluster
 - Create Docker image of program & dependencies
 - **Distribute image** to cluster nodes (very efficient, only updates sent)

1	pyspark~=2.4.0
2	pandas~=1.2.1
3	wget~=3.2
4	setuptools~=51.3.3
5	<u>findspark~=1.4.2</u>

Important to remember

- Labs illustrate common Spark development techniques
 - Development on local computer
 - Development on managed cluster
 - Modules installed on nodes before program runs
 - Development on unmanaged cluster
 - Docker image created & distributed to cluster