

AWS SageMaker

<https://aws.amazon.com/getting-started/tutorials/build-train-deploy-machine-learning-model-sagemaker/>

Create & deploy a model

- Create s3 file folder for data
- Create an XGBoost container
 - Which contains the XGBoost code, in a format known by the AWS API
- Train the model
- Deploy the model, in a new container
- Send the deployed model data
- → All model and containers managed by AWS API



Services ▾

Resource Groups ▾



Amazon SageMaker



Oregon ▾

S

AWS Management Console

AWS services

Find Services

You can enter names, keywords or acronyms.

🔍 SageMaker ✕

Amazon SageMaker

Build, Train, and Deploy Machine Learning Models

▼ Recently visited services



Amazon SageMaker



Rekognition



Billing



EC2



S3

▶ All services

Access resources on the go



Access the Management Console using the AWS Console Mobile App. [Learn more](#) ↗

Explore AWS

Visit AWS around the world at a Summit


AWS Global Summits bring the cloud computing community together to connect, collaborate, and learn about AWS. [Learn more](#) ↗

Data Lake Storage



Services ▾

Resource Groups ▾

 Amazon SageMaker



Oregon ▾

Support ▾

Amazon SageMaker



Dashboard

Search

▼ Ground Truth

Labeling jobs

Labeling datasets

Labeling workforces

▼ Notebook

Notebook instances

Lifecycle configurations

Git repositories

▼ Training

Algorithms

Training jobs

Hyperparameter tuning jobs

▼ Inference

Amazon SageMaker > Dashboard

Overview

Hide



Ground Truth

Set up and manage labeling jobs for highly accurate training datasets using active learning and human labeling.

Labeling jobs



Notebook

Availability of AWS and SageMaker SDKs and sample notebooks to create training Jobs and deploy models.

Notebook instances



Training

Train and tune models at any scale. Leverage high performance AWS algorithms or bring your own.

Training jobs



Inference

Create models from training jobs or import external models for hosting to run inferences on new data.


Models

Recent activity

Recent activity within the

Last 7 days ▾

Create notebook instance

Amazon SageMaker provides pre-built fully managed notebook instances that run Jupyter notebooks. The notebook instances include example code for common model training and hosting exercises. [Learn more](#) 

Notebook instance settings

Notebook instance name

MySageMakerInstance

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Notebook instance type

ml.t2.medium

Elastic Inference [Learn more](#) 

none

► Additional configuration

Permissions and encryption

IAM role

Notebook instances require permissions to call other services including SageMaker and S3. Choose a role or let us create a role with the [AmazonSageMakerFullAccess](#) IAM policy attached.

Choose an IAM role

Create a new role

Enter a custom IAM role ARN

Prepare file space

Create an IAM role

Passing an IAM role gives Amazon SageMaker permission to perform actions in other AWS services on your behalf. Creating a role here will grant permissions described by the [AmazonSageMakerFullAccess](#) IAM policy to the role you create.

The IAM role you create will provide access to:

- ☒ S3 buckets you specify - *optional*
 - ☐ Specific S3 buckets
 -
 - Comma delimited. ARNs, "*" and "/" are not supported.
 - ☒ Any S3 bucket
 - Allow users that have access to your notebook instance access to any bucket and its contents in your account.
 - ☐ None
- ☒ Any S3 bucket with "sagemaker" in the name
- ☒ Any S3 object with "sagemaker" in the name
- ☒ Any S3 object with the tag "sagemaker" and value "true" [See Object tagging](#)
- ☒ S3 bucket with a Bucket Policy allowing access to SageMaker [See S3 bucket policies](#)

CancelCreate role

Wait for notebook instance to be ready

Amazon SageMaker

Dashboard

Search

▼ Ground Truth

Labeling jobs

Labeling datasets

Labeling workforces

▼ Notebook

Notebook instances

Success! Your notebook instance is being created.

Open the notebook instance when status is InService and open a template notebook to get started.

View details

Amazon SageMaker > Notebook instances

Notebook instances

Search notebook instances

Actions

Create notebook instance

< 1 > ⌕

	Name	Instance	Creation time	Status	Actions
○	MySageMakerInstance	ml.t2.medium	Aug 13, 2019 20:51 UTC	⌚ Pending	-

Amazon SageMaker



Dashboard

Search

▼ Ground Truth

Labeling jobs

Labeling datasets

Labeling workforces

▼ Notebook

Notebook instances

Amazon SageMaker > Notebook instances

Notebook instances

Search notebook instances

Name	Instance	Creation time	Status	
MySageMakerInstance	ml.t2.medium	Aug 13, 2019 20:51 UTC	InService	Open JupyterLab

Actions

Open Jupyter

Open JupyterLab

Stop

Start

Update settings

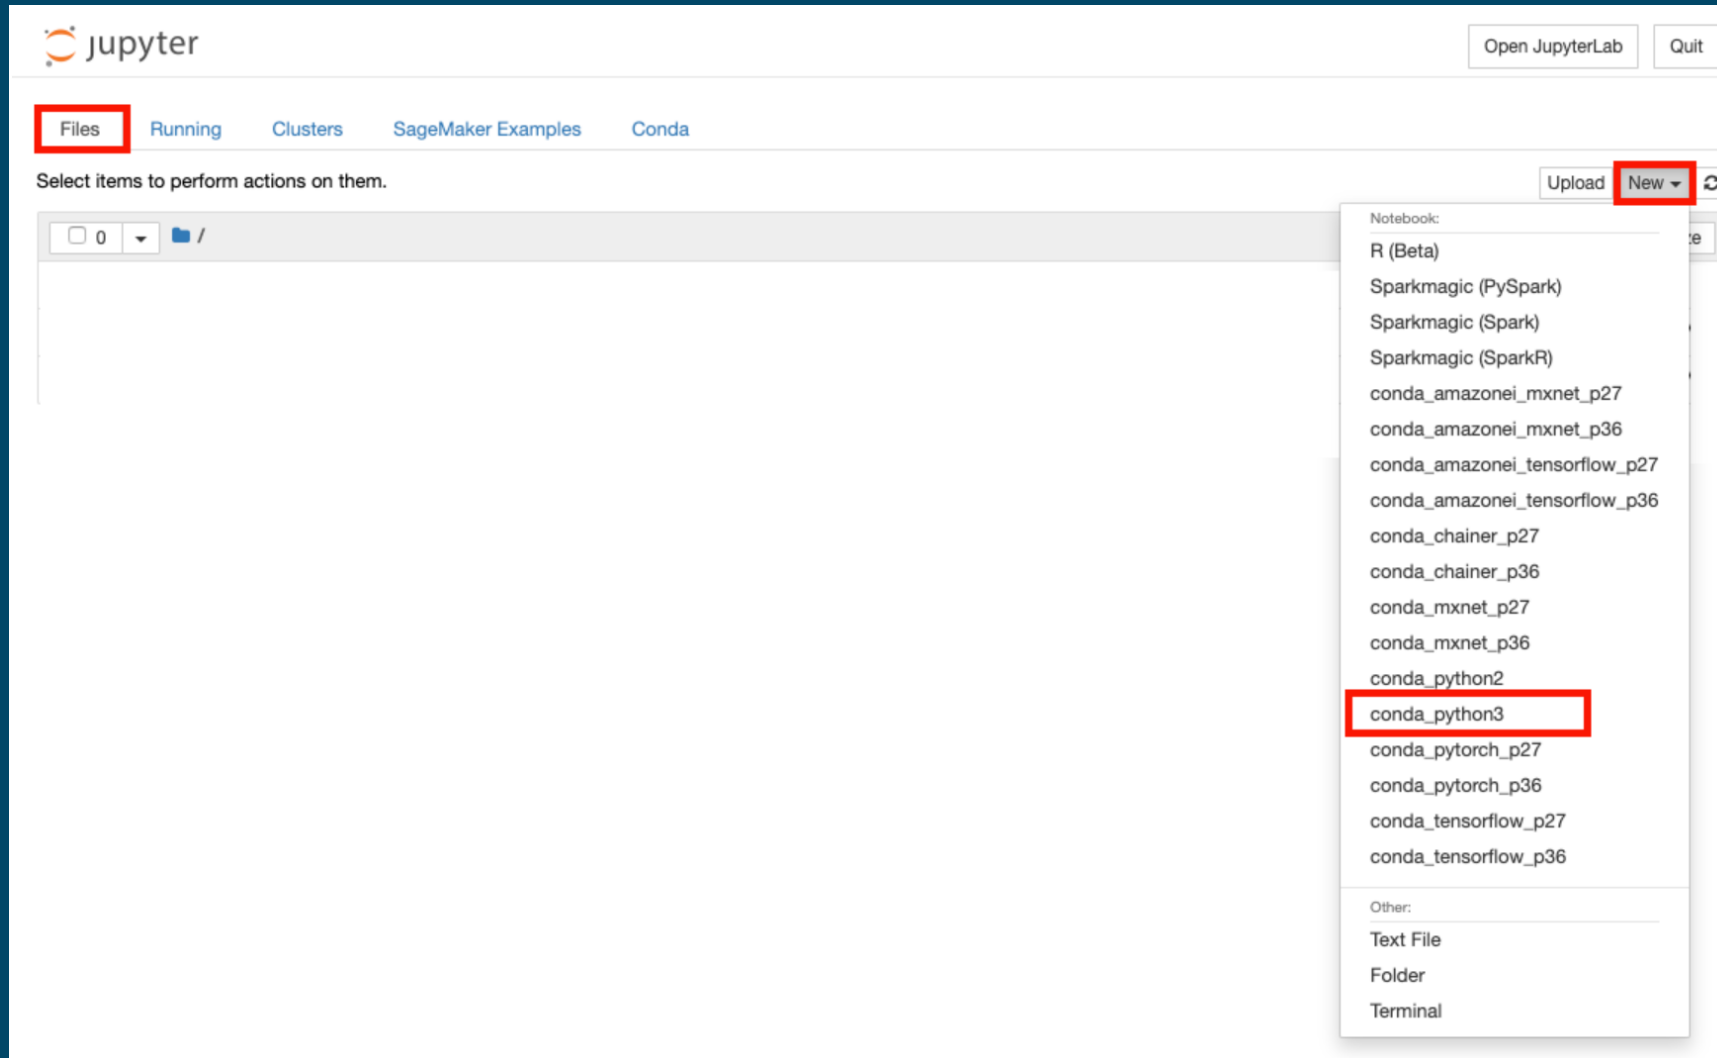
Add/Edit tags

Delete

Create notebook instance

< 1 > ⚙

Simple Python notebook



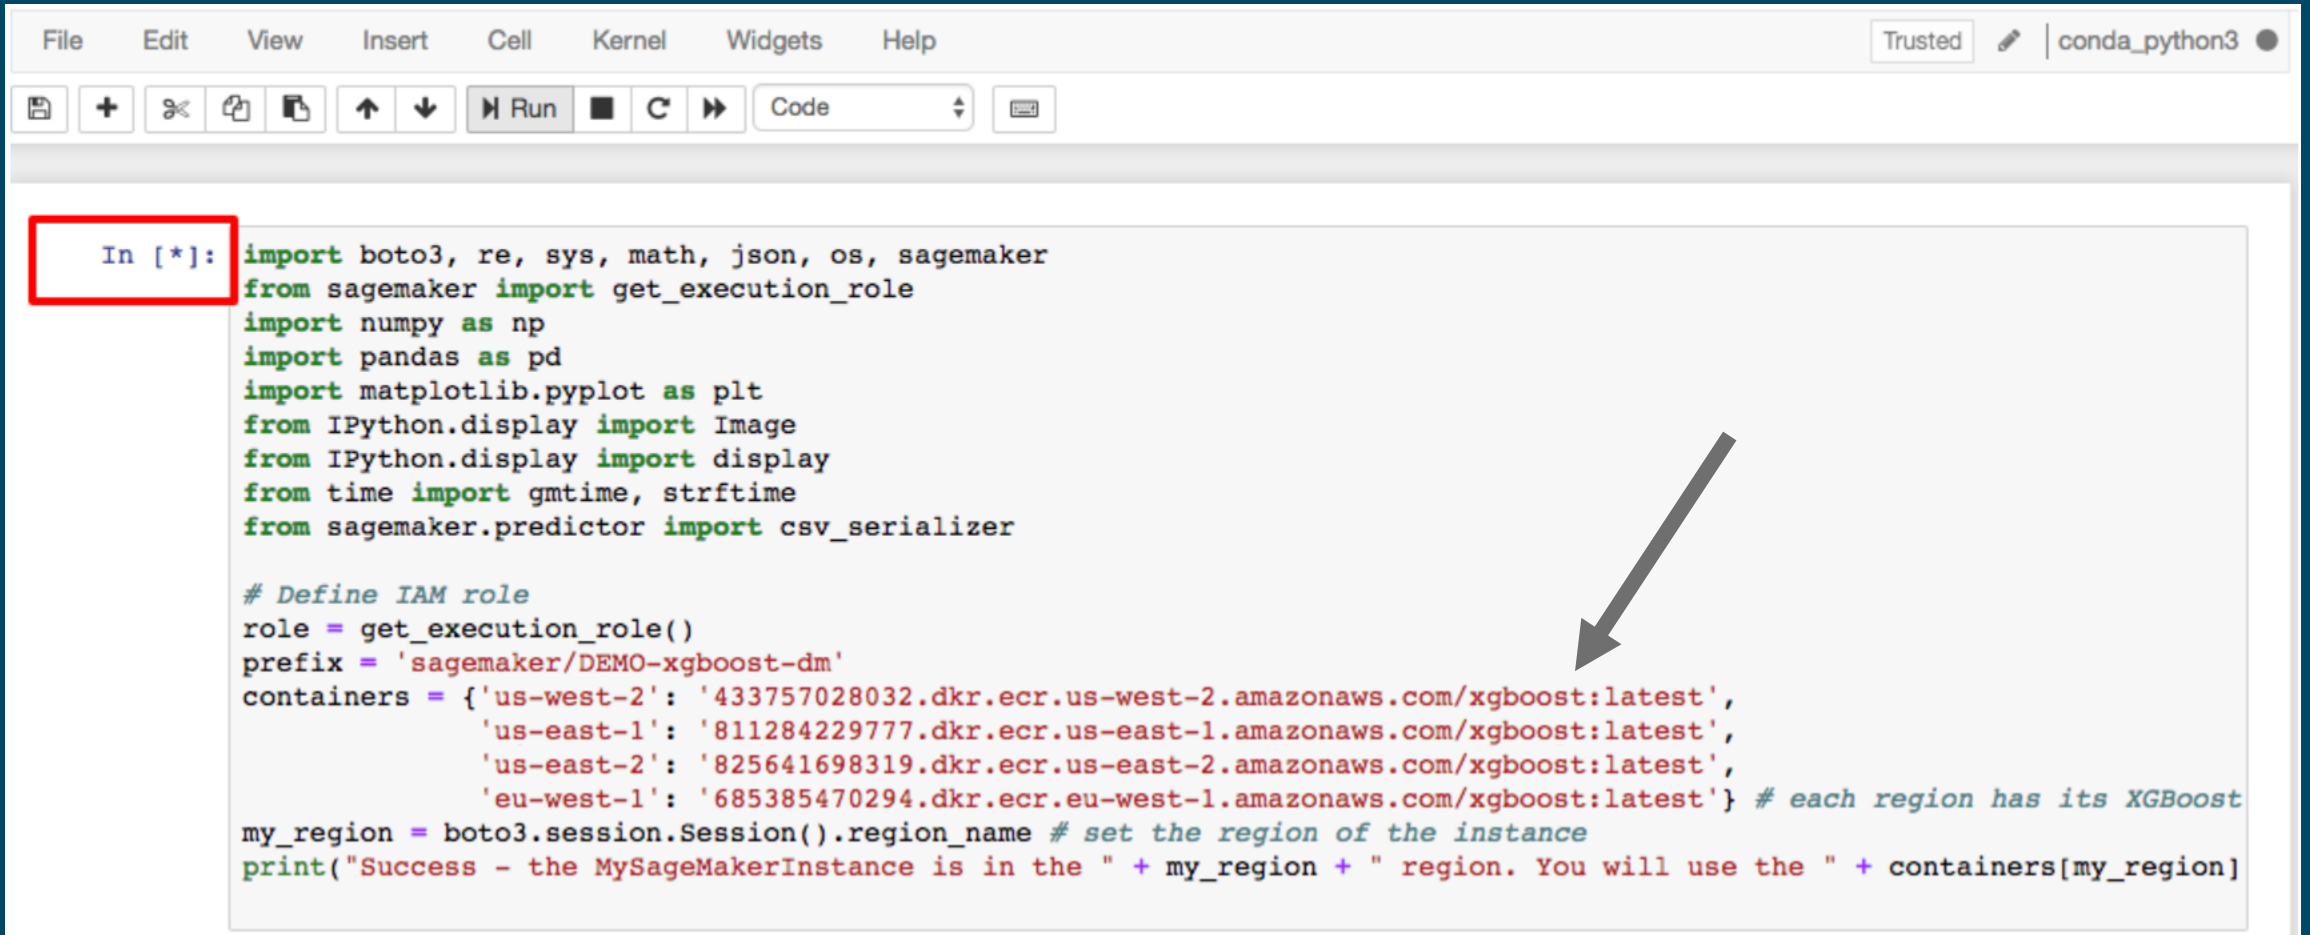
Prepare Docker XGBoost container to run

(the estimator inside the container is XGBoost)

```
File Edit View Insert Cell Kernel Widgets Help Trusted | conda_python3 ●
```

```
In [*]: import boto3, re, sys, math, json, os, sagemaker
        from sagemaker import get_execution_role
        import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        from IPython.display import Image
        from IPython.display import display
        from time import gmtime, strftime
        from sagemaker.predictor import csv_serializer

        # Define IAM role
        role = get_execution_role()
        prefix = 'sagemaker/DEMO-xgboost-dm'
        containers = {'us-west-2': '433757028032.dkr.ecr.us-west-2.amazonaws.com/xgboost:latest',
                     'us-east-1': '811284229777.dkr.ecr.us-east-1.amazonaws.com/xgboost:latest',
                     'us-east-2': '825641698319.dkr.ecr.us-east-2.amazonaws.com/xgboost:latest',
                     'eu-west-1': '685385470294.dkr.ecr.eu-west-1.amazonaws.com/xgboost:latest'} # each region has its XGBoost
        my_region = boto3.session.Session().region_name # set the region of the instance
        print("Success - the MySageMakerInstance is in the " + my_region + " region. You will use the " + containers[my_region])
```



File Edit View Insert Cell Kernel Widgets Help Trusted conda_python3

Save Add Split Copy Paste Up Down Run Stop Refresh Run All Code

```
In [1]: import boto3, re, sys, math, json, os, sagemaker
from sagemaker import get_execution_role
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from IPython.display import Image
from IPython.display import display
from time import gmtime, strftime
from sagemaker.predictor import csv_serializer

# Define IAM role
role = get_execution_role()
prefix = 'sagemaker/DEMO-xgboost-dm'
containers = {'us-west-2': '433757028032.dkr.ecr.us-west-2.amazonaws.com/xgboost:latest',
              'us-east-1': '811284229777.dkr.ecr.us-east-1.amazonaws.com/xgboost:latest',
              'us-east-2': '825641698319.dkr.ecr.us-east-2.amazonaws.com/xgboost:latest',
              'eu-west-1': '685385470294.dkr.ecr.eu-west-1.amazonaws.com/xgboost:latest'} # each region has its XGBoost
my_region = boto3.session.Session().region_name # set the region of the instance
print("Success - the MySageMakerInstance is in the " + my_region + " region. You will use the " + containers[my_region])
```

Success - the MySageMakerInstance is in the us-east-1 region. You will use the 811284229777.dkr.ecr.us-east-1.amazonaws.com/xgboost:latest container for your SageMaker endpoint.

In []: |

A file folder (bucket) for processing

```
In [2]: bucket_name = 'testyourname' # <--- change this variable to a unique name for your bucket
s3 = boto3.resource('s3')
try:
    if my_region == 'us-east-1':
        s3.create_bucket(Bucket=bucket_name)
    else:
        s3.create_bucket(Bucket=bucket_name, CreateBucketConfiguration={ 'LocationConstraint': my_region })
    print('S3 bucket created successfully')
except Exception as e:
    print('S3 error: ',e)
```

S3 bucket created successfully

Data available for processing

```
In [3]: try:
        urllib.request.urlretrieve ("https://d1.awsstatic.com/tmt/build-train-deploy-machine-learning-model-sagemaker/bank_clean.csv")
        print('Success: downloaded bank_clean.csv.')
    except Exception as e:
        print('Data load error: ',e)

    try:
        model_data = pd.read_csv('./bank_clean.csv',index_col=0)
        print('Success: Data loaded into dataframe.')
    except Exception as e:
        print('Data load error: ',e)
```

```
Success: downloaded bank_clean.csv.
Success: Data loaded into dataframe.
```

Split data


```
In [4]: train_data, test_data = np.split(model_data.sample(frac=1, random_state=1729), [int(0.7 * len(model_data))])  
        print(train_data.shape, test_data.shape)  
(28831, 61) (12357, 61)
```

Format the input data

Prep the input data file

```
pd.concat([train_data['y_yes'], train_data.drop(['y_no', 'y_yes'], axis=1)],  
axis=1).to_csv('train.csv', index=False, header=False)  
boto3.Session().resource('s3').Bucket(bucket_name).Object(os.path.join(prefix,  
'train/train.csv')).upload_file('train.csv')  
s3_input_train = sagemaker.s3_input(s3_data='s3://{}/{}/train'.format(bucket_name, prefix),  
content_type='csv')
```

Create a sized instance of the XGBoost container



```
# Create the instance (Docker) with the XGBost estimator
sess = sagemaker.Session()
xgb = sagemaker.estimator.Estimator(containers[my_region],role, train_instance_count=1,
train_instance_type='ml.m4.xlarge',output_path='s3://{}/{}/output'.format(bucket_name,
prefix),sagemaker_session=sess)
xgb.set_hyperparameters(max_depth=5,eta=0.2,gamma=4,min_child_weight=6,subsample=0
.8,silent=0,objective='binary:logistic',num_round=100)
```


Run the container


```
In [7]: xgb.fit({'train': s3_input_train})
```

```
[92]#011train-error:0.095314  
[17:36:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 10 extra nodes, 14 pruned nodes, max_depth=5  
[93]#011train-error:0.095314  
[17:36:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 24 extra nodes, 30 pruned nodes, max_depth=5  
[94]#011train-error:0.095314  
[17:36:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 6 extra nodes, 24 pruned nodes, max_depth=3  
[95]#011train-error:0.095314  
[17:36:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 12 extra nodes, 30 pruned nodes, max_depth=5  
[96]#011train-error:0.095279  
[17:36:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 18 extra nodes, 12 pruned nodes, max_depth=5  
[97]#011train-error:0.094828  
[17:36:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 4 extra nodes, 22 pruned nodes, max_depth=2  
[98]#011train-error:0.094863  
[17:36:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 30 extra nodes, 12 pruned nodes, max_depth=5  
[99]#011train-error:0.094759
```

```
2019-08-15 17:36:34 Uploading - Uploading generated training model  
2019-08-15 17:36:34 Completed - Training job completed  
Billable seconds: 56
```

Deploy the generated model

- Note that the tools know that this is a XGBoost container
 - Thus it knows where the model is stored
 - It uses that information to create a new container with



```
In [9]: xgb_predictor = xgb.deploy(initial_instance_count=1,instance_type='ml.m4.xlarge')
```

```
INFO:sagemaker:Creating model with name: xgboost-2018-07-13-14-29-39-425
```

```
INFO:sagemaker:Creating endpoint with name xgboost-2018-07-13-14-25-03-272
```

```
-----!
```

Send data to the deployed model

```
In [20]: test_data_array = test_data.drop(['y_no', 'y_yes'], axis=1).values #load the data into an array
xgb_predictor.content_type = 'text/csv' # set the data type for an inference
xgb_predictor.serializer = csv_serializer # set the serializer type
predictions = xgb_predictor.predict(test_data_array).decode('utf-8') # predict!
predictions_array = np.fromstring(predictions[1:], sep=',') # and turn the prediction into an array
print(predictions_array.shape)
```

```
(12357,)
```

Create & deploy a model

- Create s3 file folder for data
- Create an XGBoost container
 - Which contains the XGBoost code, in a format known by the AWS API
- Train the model
- Deploy the model, in a new container
- Send the deployed model data
- → All model and containers managed by AWS API