GCP Dataproc

Google GCP

Data Proc Cluster setup on Google Cloud Platform

https://cloud.google.com/dataproc/docs/tutorials/jupyter-notebook

Signing up

- GCP offers $300 credit for new users.
- Navigate to this link to get started.
- Login with your Gmail ID to begin.
- Select “United States” as country and “Continue”.
- Select Account Type as “Individual” and fill in your personal details.
- You will be asked for credit card details to make sure you’re not a robot. **Google does not charge you even if you exhaust $300 credit.**
- You’re set if you reach “Getting Started” page.

Setting up the project

- Open the link given in the title slide and follow the instructions.
- For all the works you do in GCP, it is mandatory to select a project.
- Click the “Go to the project selector page” button to set up a project.
- Once you’re in the dashboard, select “Create Project” on top right.
- Give a name to the project.
- Click “Create”.

New Project

You have 20 projects remaining in your quota. Request an increase or delete projects. Learn more

MANAGE QUOTAS

Project name *

CIS8795-demo

Project ID: cis8795-demo. It cannot be changed later. EDIT

Billing account *

CIS 8795: IT Infrastructure for Big Data

Any charges for this project will be billed to the account you select here.

Location *

No organization BROWSE

Parent organization or folder

CREATE CANCEL

Enable APIs

- Since we are working on Data Proc which uses the Compute Engine for underlying infrastructure, we need to enable the corresponding APIs before spinning up the cluster.
- Click on "Enable the APIs" button.
  - See the button on the https://cloud.google.com/dataproc/docs/tutorials/jupyter-notebook page
- Select the project you just created and hit "Continue".

Search for API

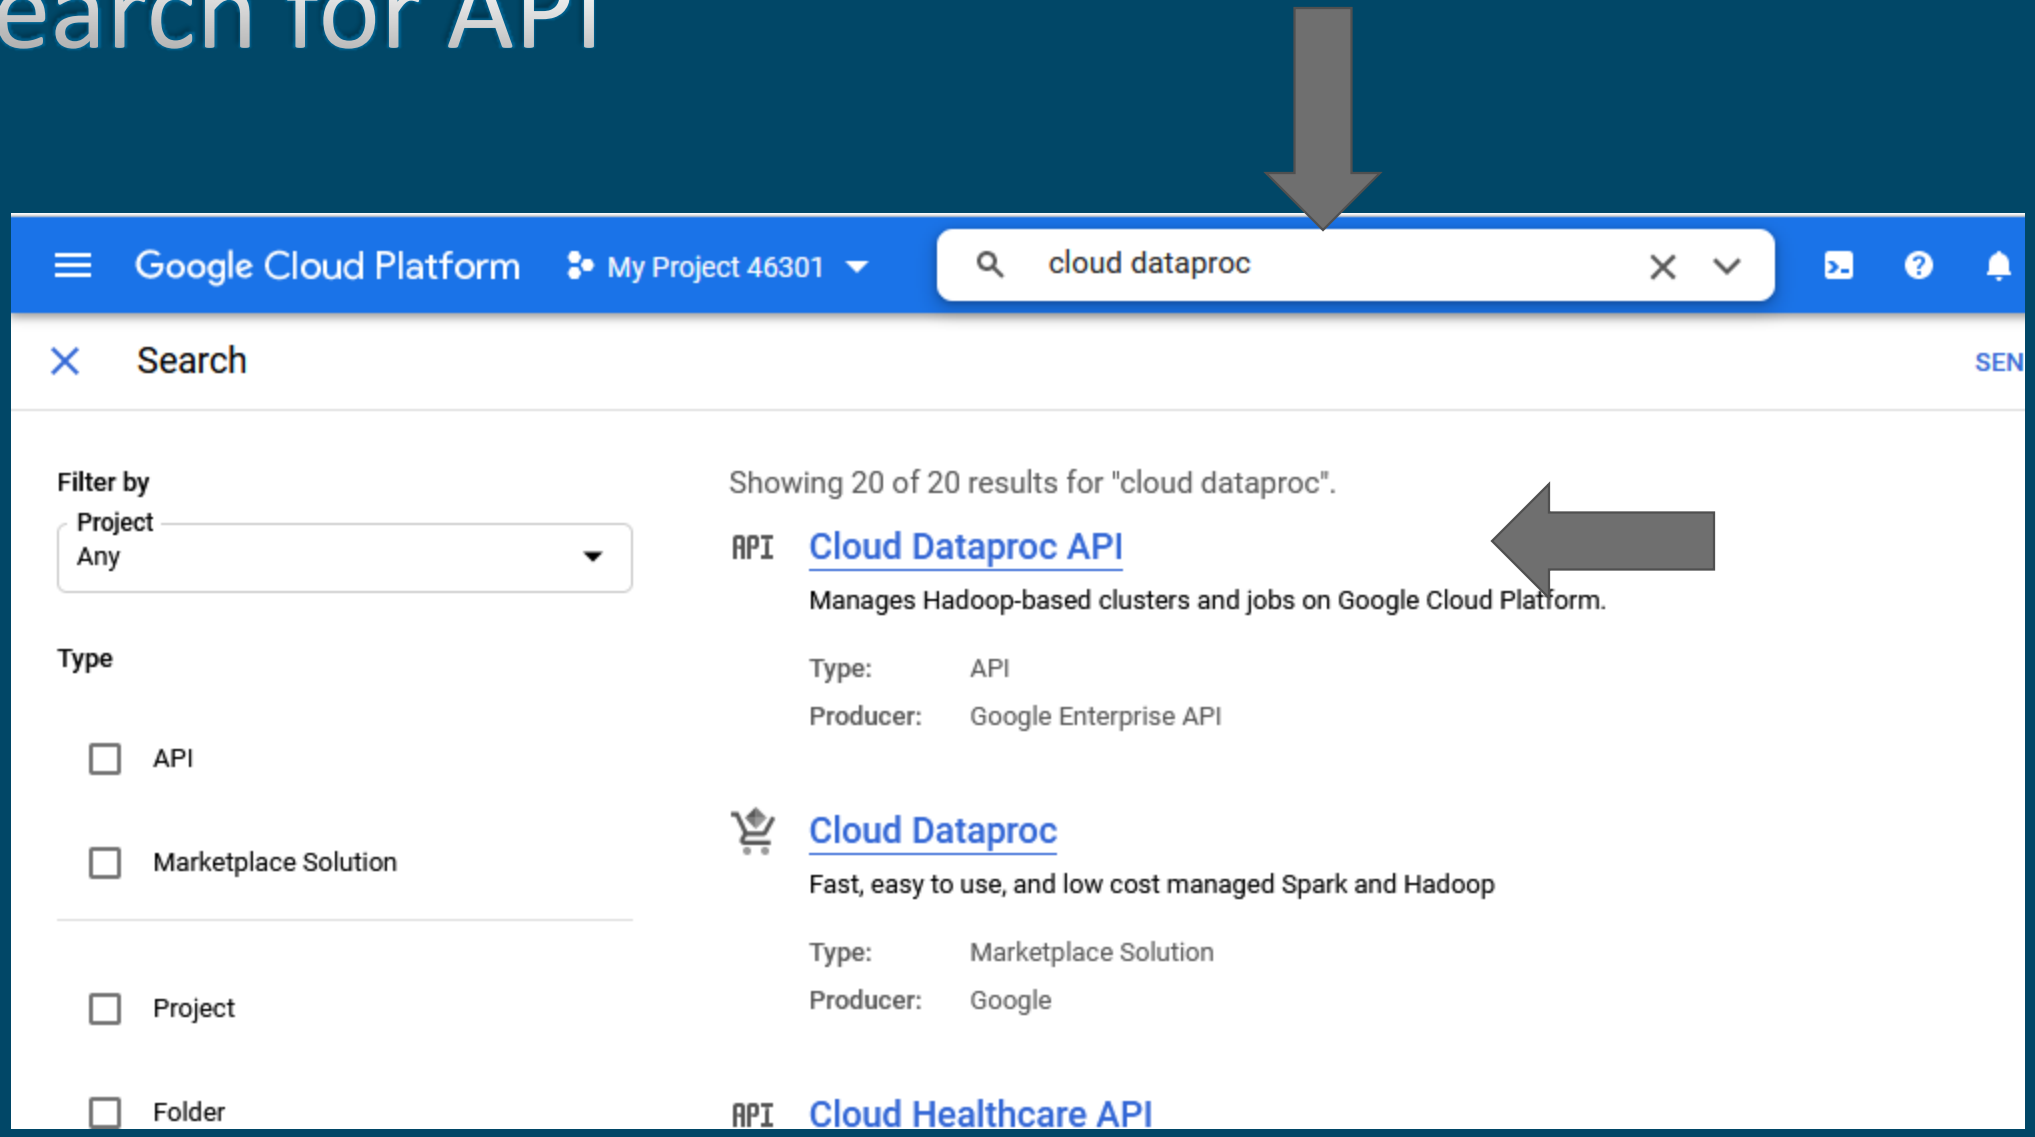

Search for: Cloud Dataproc API
select enable (ensure your project is selected)

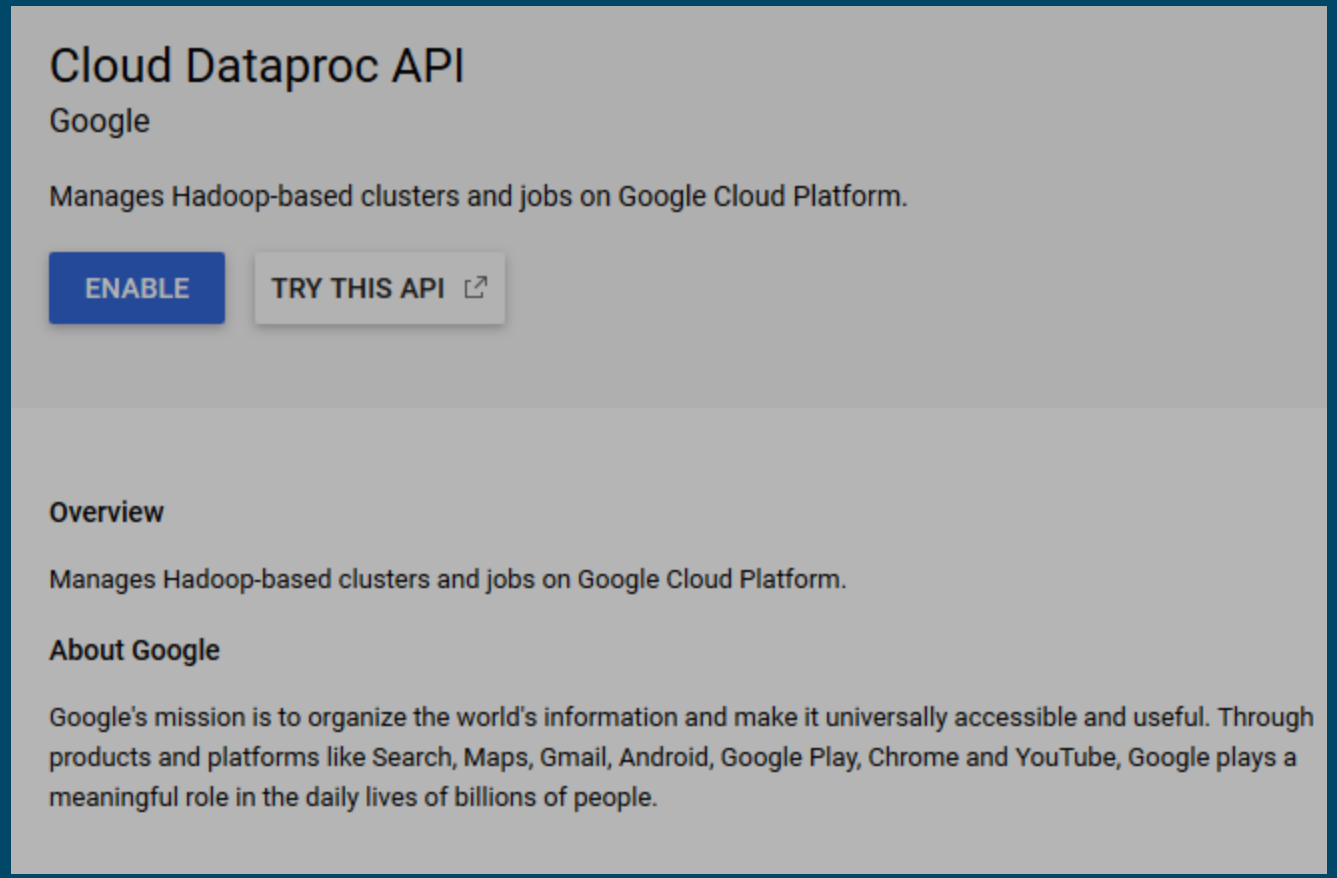

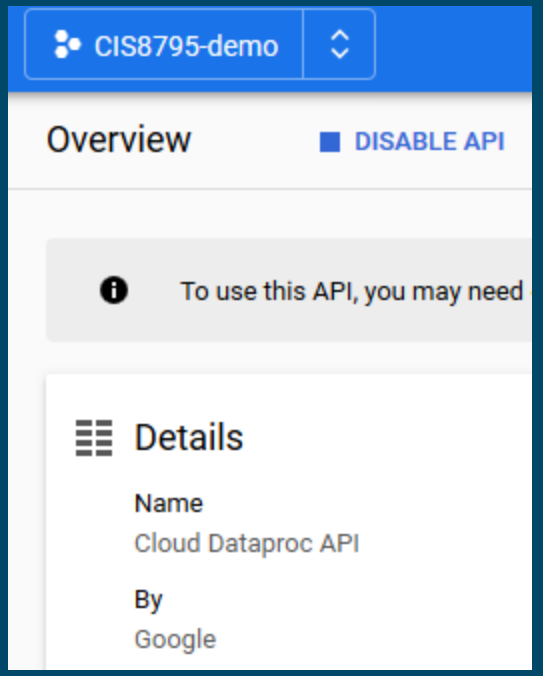

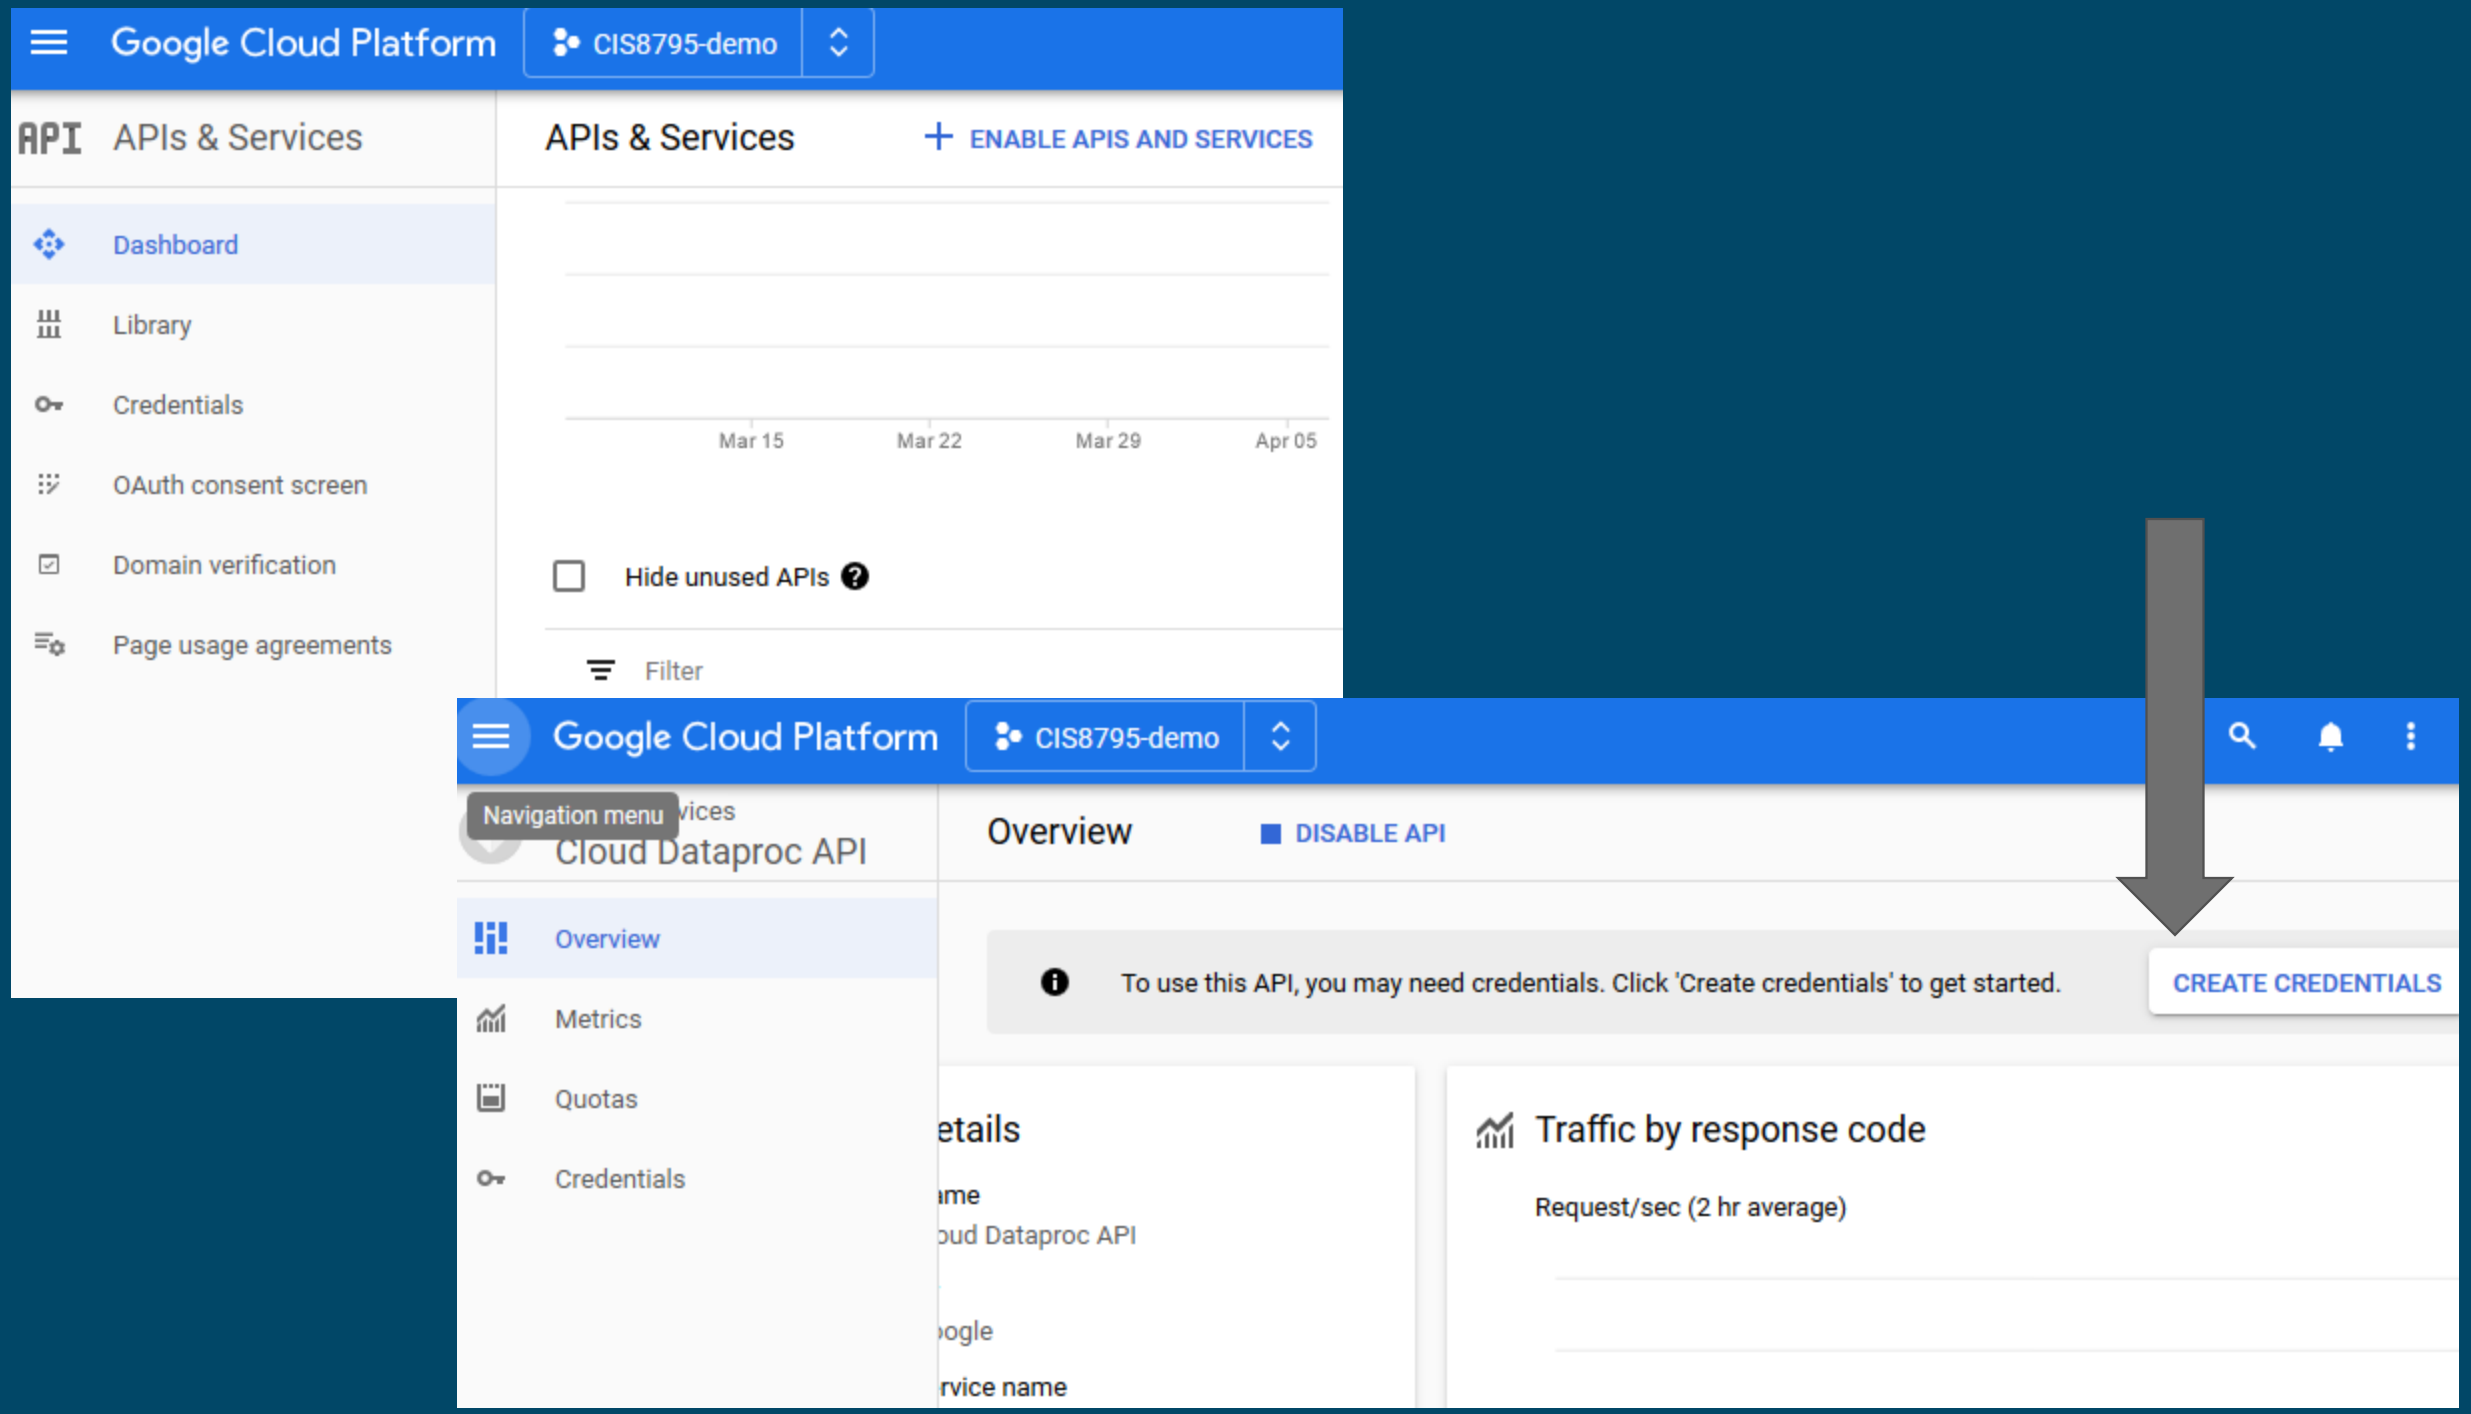
Google Cloud Platform
CIS8795-demo
APIs & Services
APIs & Services
ENABLE APIS AND SERVICES
Dashboard
Library
Credentials
OAuth consent screen
Domain verification
Page usage agreements
Mar 15
Mar 22
Mar 29
Apr 05
Hide unused APIs
Filter
Google Cloud Platform
CIS8795-demo
Navigation menu
Cloud Dataproc API
Overview
Metrics
Quotas
Credentials
Overview
DISABLE API
To use this API, you may need credentials. Click 'Create credentials' to get started.
CREATE CREDENTIALS
Traffic by response code
Request/sec (2 hr average)

Establishing the Credentials

- Once the APIs are enabled, you need to set your credentials by clicking “Create Credentials”.
  - You’ll need to wait for the API’s to be enabled
  - API & Services -> Credentials
- Select “Application Data” and “Yes, I’m using one or more”
- And click “Done”
- Once credentials are established, Click on “CONFIGURE CONSENT SCREEN”
  - External
- Give a name for the application, and click on “Save”

Create credentials

1 Credential Type

Which API are you using?

Different APIs use different auth platforms and some credentials can be restricted to only call certain APIs.

Select an API *
Cloud Dataproc API

What data will you be accessing? *

Different credentials are required to authorize access depending on the type of data that you request. Learn more

This Google Cloud Platform API is usually accessed from a server using a service account. To create a service account, select "Application data".

- ○ User data
  Data belonging to a Google user, like their email address or age. User consent required. This will create an OAuth client.
- ◉ Application data
  Data belonging to your own application, such as your app's Cloud Firestore backend. This will create a service account.

Are you planning to use this API with Compute Engine, Kubernetes Engine, App Engine, or Cloud Functions?

Applications running on GCE, GKE, GAE, and GCF can use Application Default Credentials and don't require that you create a credential.

- ◉ Yes, I'm using one or more
- ○ No, I'm not using them

NEXT

2 Your Credentials

DONE CANCEL

2021 style credentials

Credentials + CREATE CREDENTIALS DELETE

Create credentials to access your enabled APIs. Learn more

Remember to configure the OAuth consent screen with information about your application. CONFIGURE CONSENT SCREEN

API Keys

Name Key Actions

No API keys to display

OAuth consent screen

Choose how you want to configure and register your app, including your target users. You can only associate one app with your project.

User Type

Internal

Only available to users within your organization. You will not need to submit your app for verification. Learn more about user type

External

Available to any test user with a Google Account. Your app will start in testing mode and will only be available to users you add to the list of test users. Once your app is ready to push to production, you may need to verify your app. Learn more about user type

CREATE

Let us know what you think about our OAuth experience

Application registration

- Give app name
- Email for
  - User support
  - Developer contact
- Save

- Go to Home Dashboard

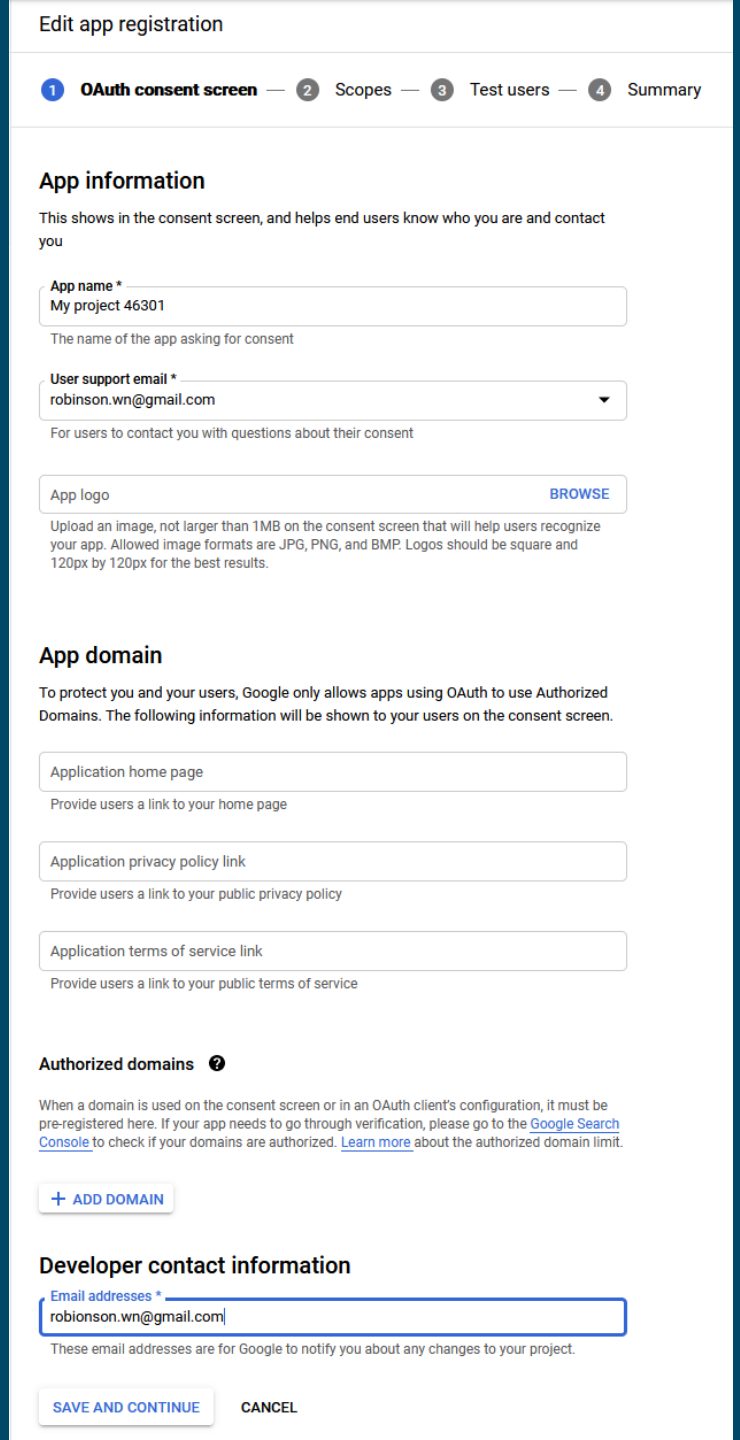

Home – BQ Project – Google C | j5mm5zaq5vco7b74ozfvx3ikpy | j5mm5zaq5vco7b74ozfvx3ikpy

https://console.cloud.google.com/home/dashboard?project=bq-project-1

Getting Started | GSU | writing | docker | DA | my | stats

Google Cloud Platform | BQ Project

DASHBOARD | ACTIVITY | CUSTOMIZE

Project info

Project name
BQ Project

Project ID
bq-project-1

Project number
915265256313

ADD PEOPLE TO THIS PROJECT

Go to project settings

Resources

This project has no resources

Trace

No trace data from the past 7 days

Get started with Stackdriver Trace

Getting Started

Explore and enable APIs

Deploy a prebuilt solution

APIs

Requests (requests/sec)

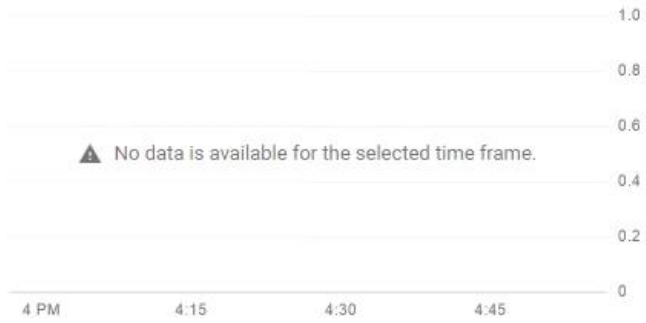

Go to APIs overview

Google Cloud Platform status

All services normal

Go to Cloud status dashboard

Error Reporting

No sign of any errors. Have you set up Error Reporting?

Learn how to set up Error Reporting

News

Announcing the winners of our Google Cloud 2019 Partner Awards
7 hours ago

Machine learning with XGBoost gets faster with Dataproc on GPUs
7 hours ago

How Google Cloud is helping U.S public sector agencies during the COVID-19 pandemic and beyond
3 days ago

Read all news

Documentation

Learn about Compute Engine

Create a Storage Bucket

- Click on Go to the Cloud Storage Browser page
- Click on CREATE BUCKET
- Give your bucket a name.
- Select Region and hit Continue
- Select Standard as default storage class and hit Continue.
- Select Fine-grained access control and hit Continue.
- Select Google-managed key and hit Continue.
- Navigate to the bucket and click on “Upload files” and upload the .csv and .ipynb files.

Create a bucket for your project

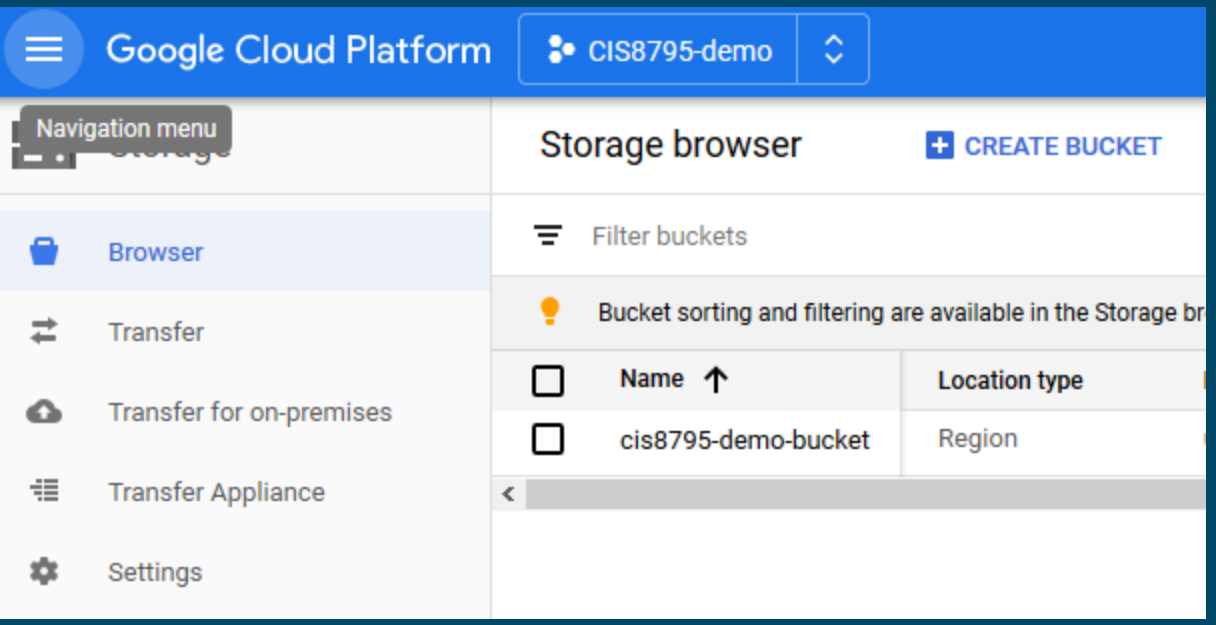

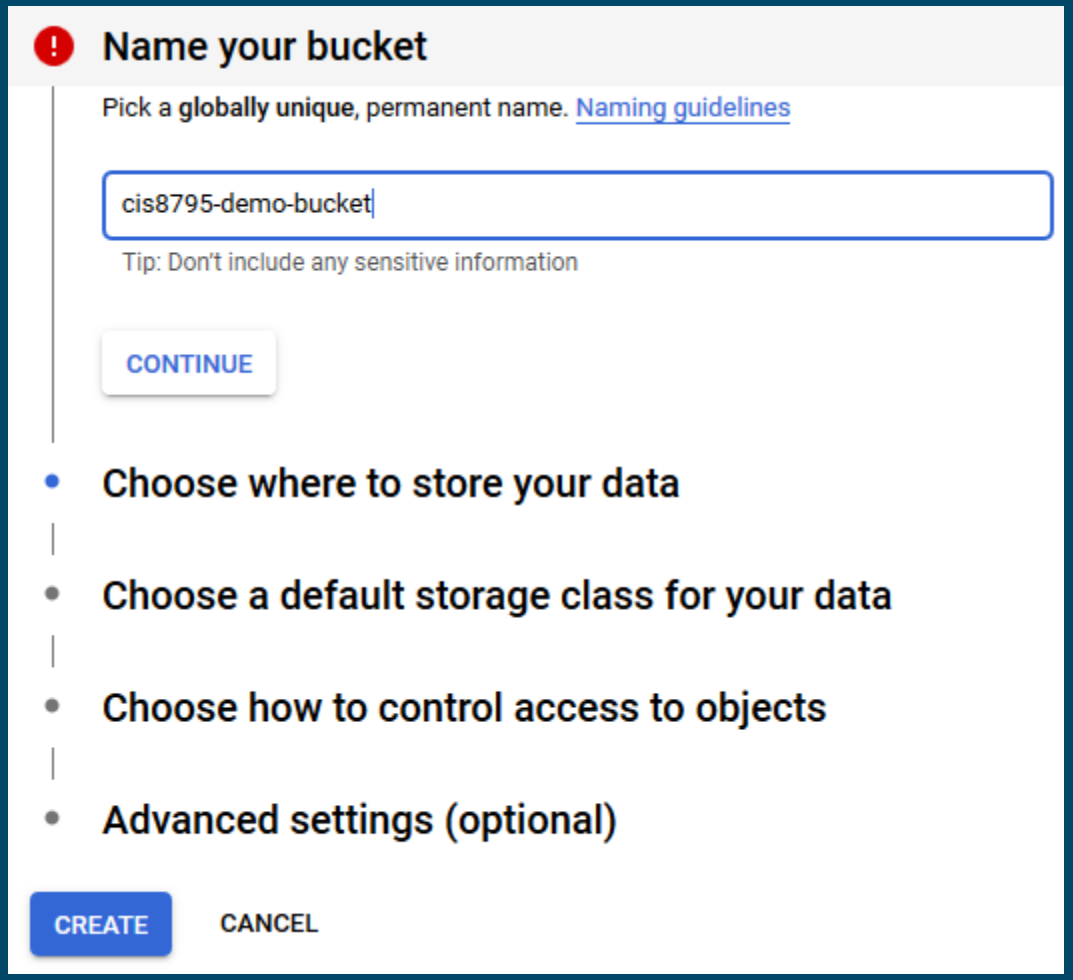

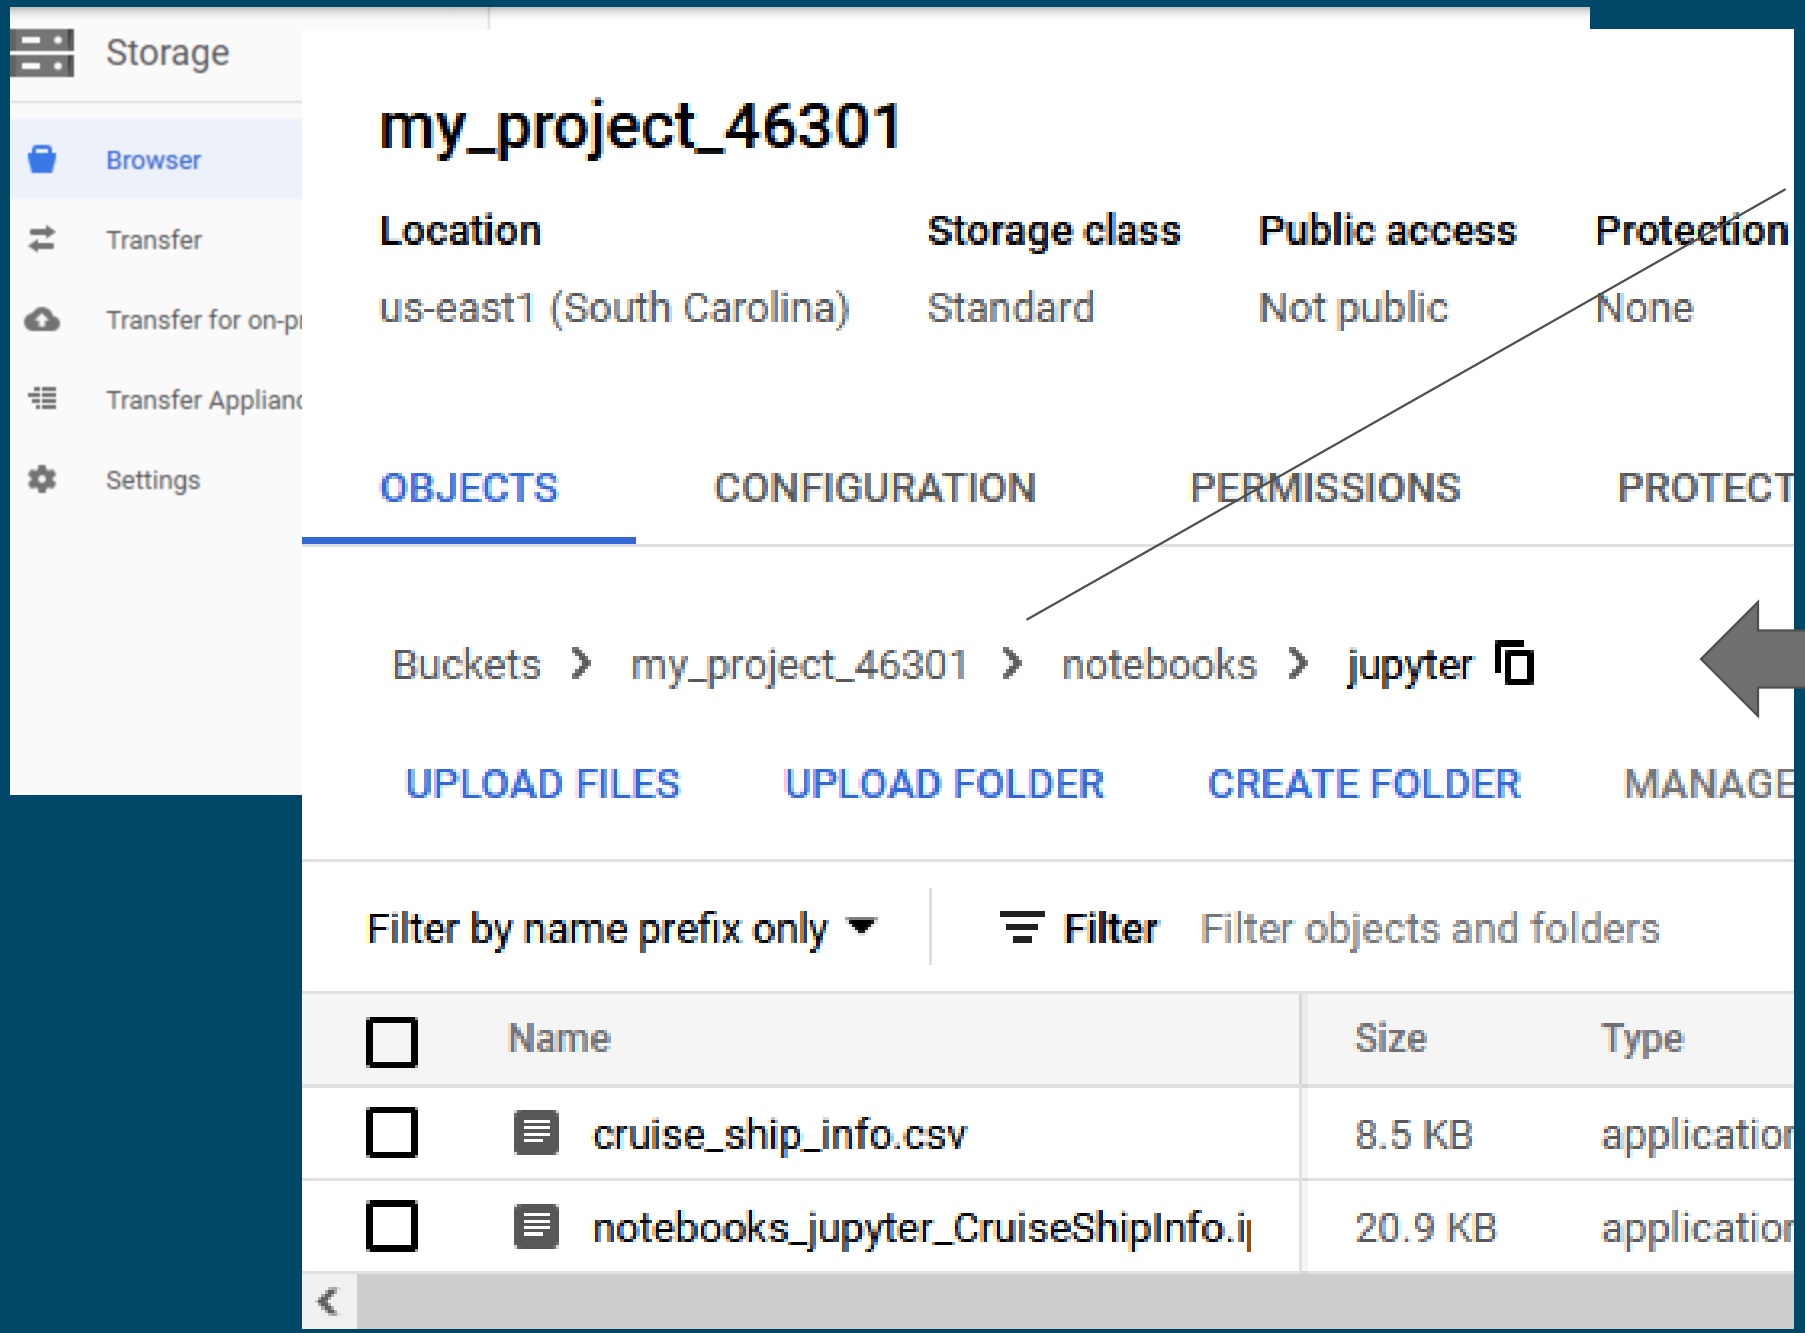

Upload notebooks & data after cluster creation

Path
project/notebooks/jupyter

Optional:
Install and Initialize the Cloud SDK

- Click on “Install and initialize the Cloud SDK”.
- Choose your operating system and make sure Python is installed.
- Download the SDK file.
- And run the “gcloud init” command on your terminal.
- Select option 1 - Re-initialize this configuration [default] with new settings
- Choose your email ID.
- Choose your project.
- Choose No for default region.

Creating Cluster

- Click on “Cloud Console” and choose “Dataproc”
- Select “Clusters” and “Create cluster”
- Set up Cluster
  - Under “Component gateway”, Enable (for web access)
  - Under Optional components, Select “Jupyter notebook”
- Configure nodes
  - Master Node to have 2 CPUs
  - Worker Node to have 2 CPUs
- Customize cluster
  - Under “Cloud Storage staging bucket” configure the bucket as your default storage option.
- Click “Create”

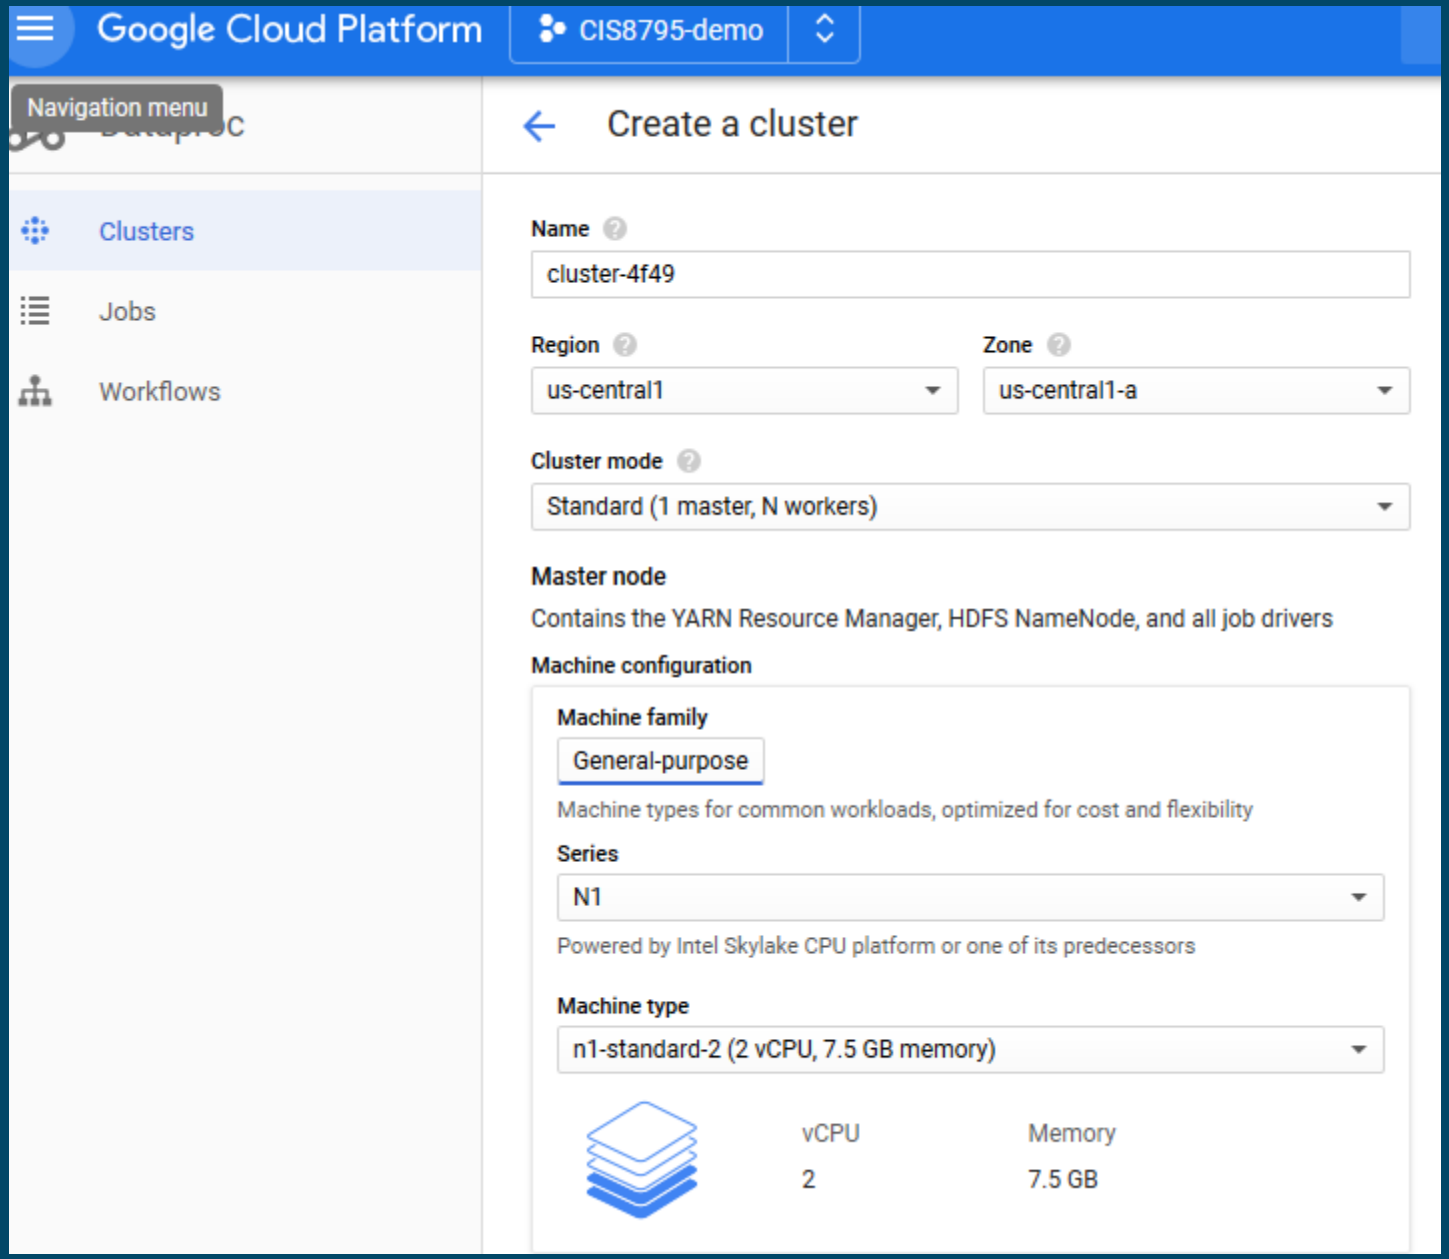
Google Cloud Platform
CIS8795-demo
Navigation menu
Clusters
Jobs
Workflows
Create a cluster
Name
cluster-4f49
Region
us-central1
Zone
us-central1-a
Cluster mode
Standard (1 master, N workers)
Master node
Contains the YARN Resource Manager, HDFS NameNode, and all job drivers
Machine configuration
Machine family
General-purpose
Machine types for common workloads, optimized for cost and flexibility
Series
N1
Powered by Intel Skylake CPU platform or one of its predecessors
Machine type
n1-standard-2 (2 vCPU, 7.5 GB memory)
vCPU
2
Memory
7.5 GB

Google Cloud Platform | CIS8795-demo

Dataproc

- Clusters
- Jobs
- Workflows

Create a cluster

The HDFS replication factor is 2.

Machine configuration

Machine family

General-purpose

Machine types for common workloads, optimized for cost and flexibility

Series

N1

Powered by Intel Skylake CPU platform or one of its predecessors

Machine type

n1-standard-2 (2 vCPU, 7.5 GB memory)

vCPU	Memory
2	7.5 GB

CPU platform and GPU

Primary disk size (minimum 15 GB) 500 GB

Primary disk type Standard persistent disk

Nodes (minimum 2) 3

Local SSDs (0-8) 0 x 375 GB

YARN cores	YARN memory
6	18 GB

Autoscaling policy (Optional)

Enable autoscaling on the cluster.

This project does not currently have any applicable policy to enable autoscaling in this region. Learn how to create autoscaling policy.

Component gateway

Enable access to the web interfaces of default and selected optional components on the cluster. Learn more

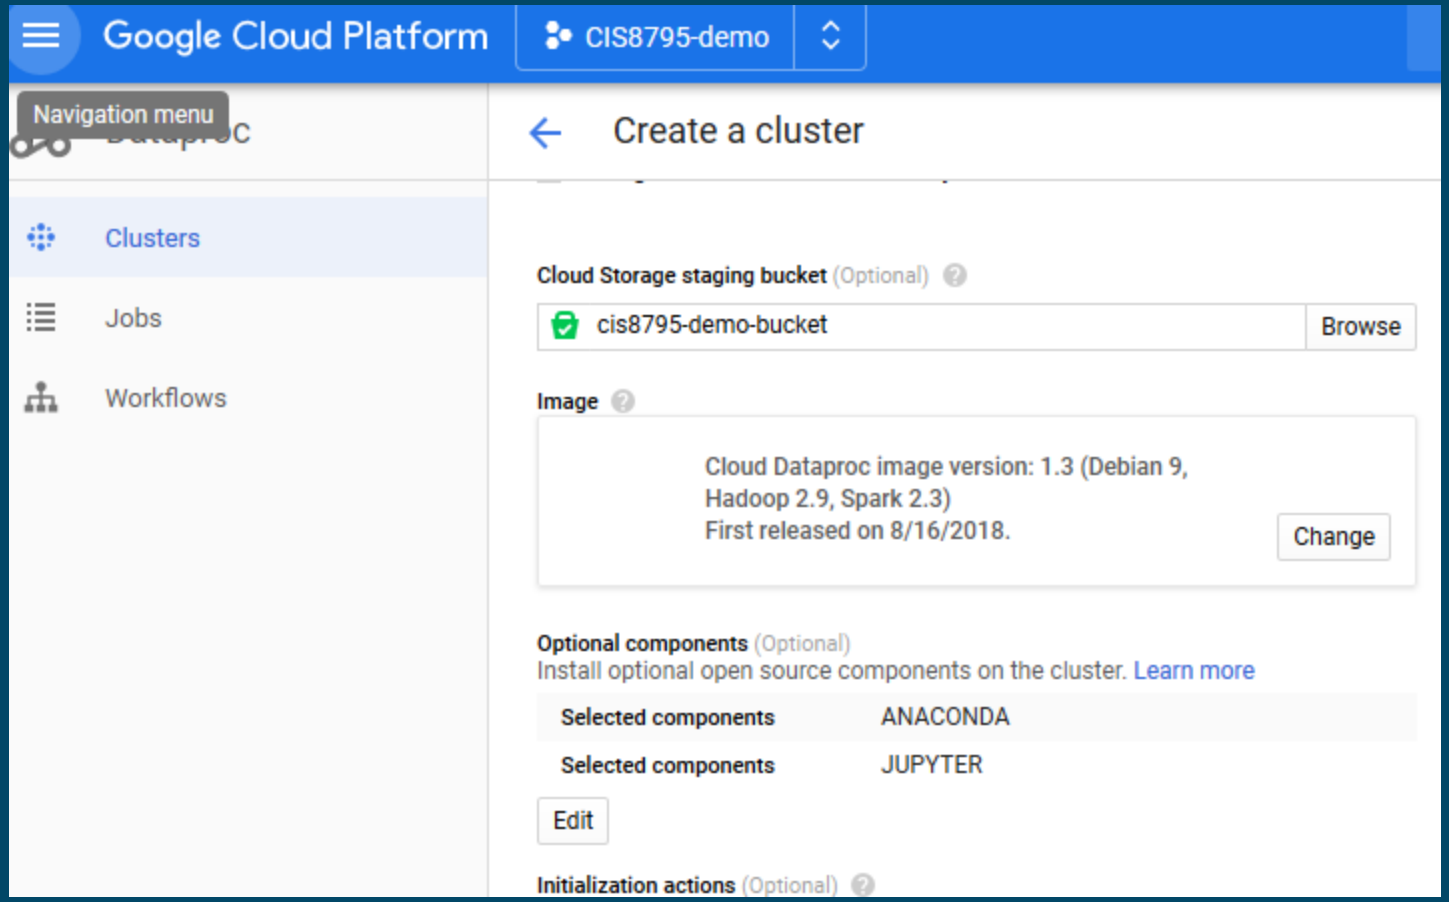
Google Cloud Platform
CIS8795-demo
Navigation menu
Create a cluster
Clusters
Jobs
Workflows
Cloud Storage staging bucket (Optional)
cis8795-demo-bucket
Browse
Image
Cloud Dataproc image version: 1.3 (Debian 9, Hadoop 2.9, Spark 2.3)
First released on 8/16/2018.
Change
Optional components (Optional)
Install optional open source components on the cluster. Learn more
Selected components ANACONDA
Selected components JUPYTER
Edit
Initialization actions (Optional)

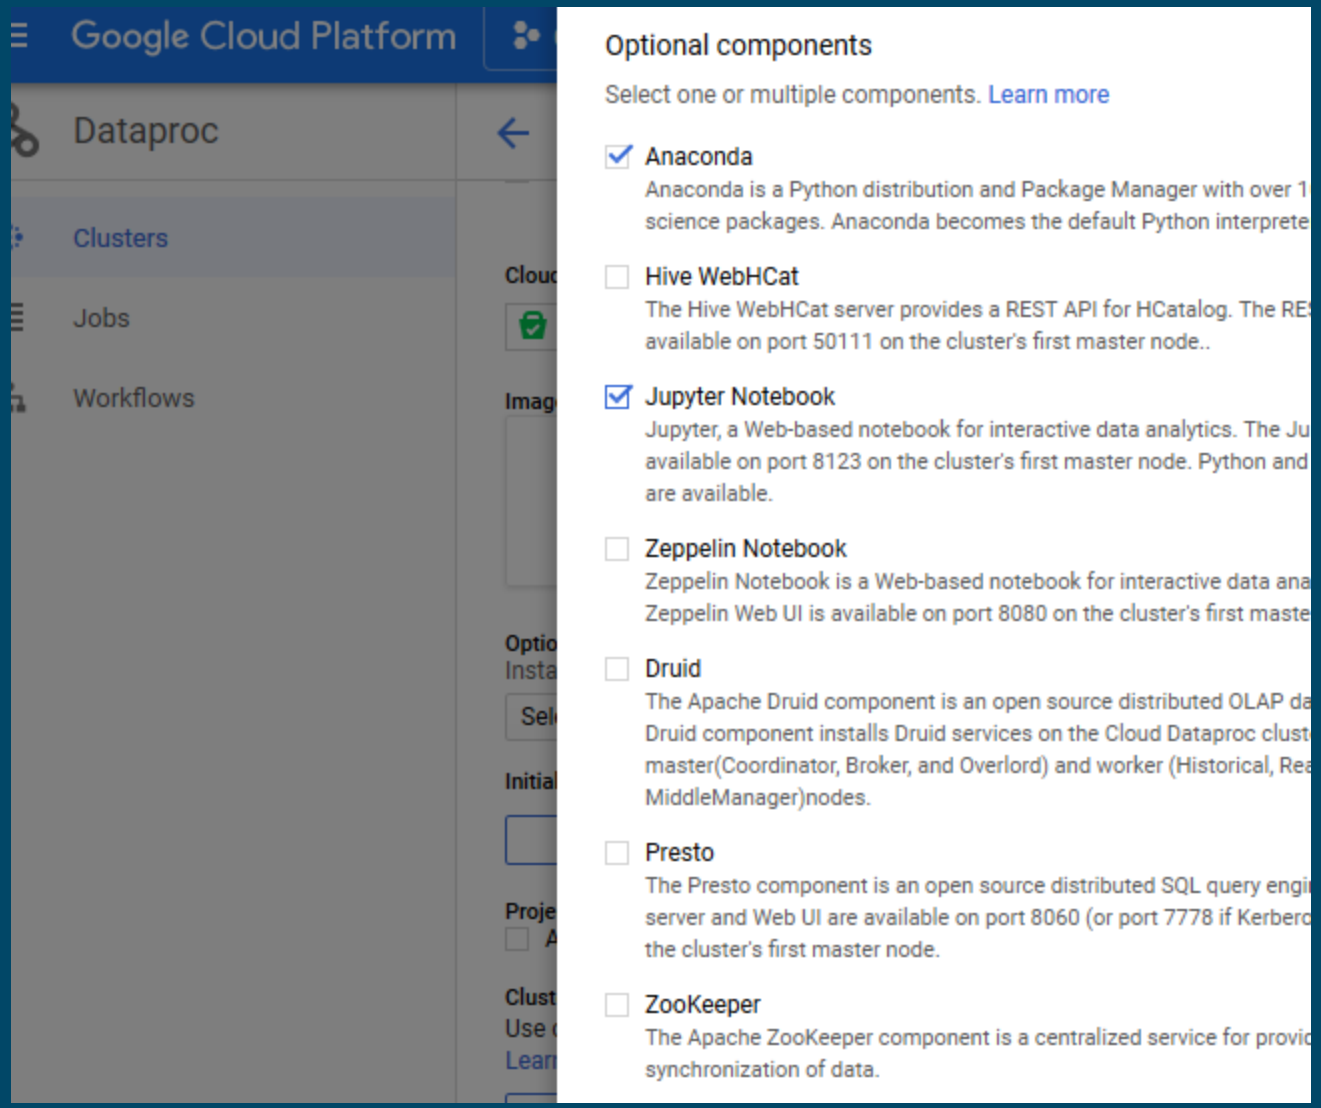
Google Cloud Platform
Dataproc
Clusters
Jobs
Workflows
Optional components
Select one or multiple components. Learn more
Anaconda
Anaconda is a Python distribution and Package Manager with over 1
science packages. Anaconda becomes the default Python interprete
Hive WebHCat
The Hive WebHCat server provides a REST API for HCatalog. The RE
available on port 50111 on the cluster's first master node..
Jupyter Notebook
Jupyter, a Web-based notebook for interactive data analytics. The Ju
available on port 8123 on the cluster's first master node. Python and
are available.
Zeppelin Notebook
Zeppelin Notebook is a Web-based notebook for interactive data ana
Zeppelin Web UI is available on port 8080 on the cluster's first maste
Druid
The Apache Druid component is an open source distributed OLAP da
Druid component installs Druid services on the Cloud Dataproc clust
master(Coordinator, Broker, and Overlord) and worker (Historical, Rea
MiddleManager)nodes.
Presto
The Presto component is an open source distributed SQL query engi
server and Web UI are available on port 8060 (or port 7778 if Kerberc
the cluster's first master node.
ZooKeeper
The Apache ZooKeeper component is a centralized service for provic
synchronization of data.

After about 15 minutes...

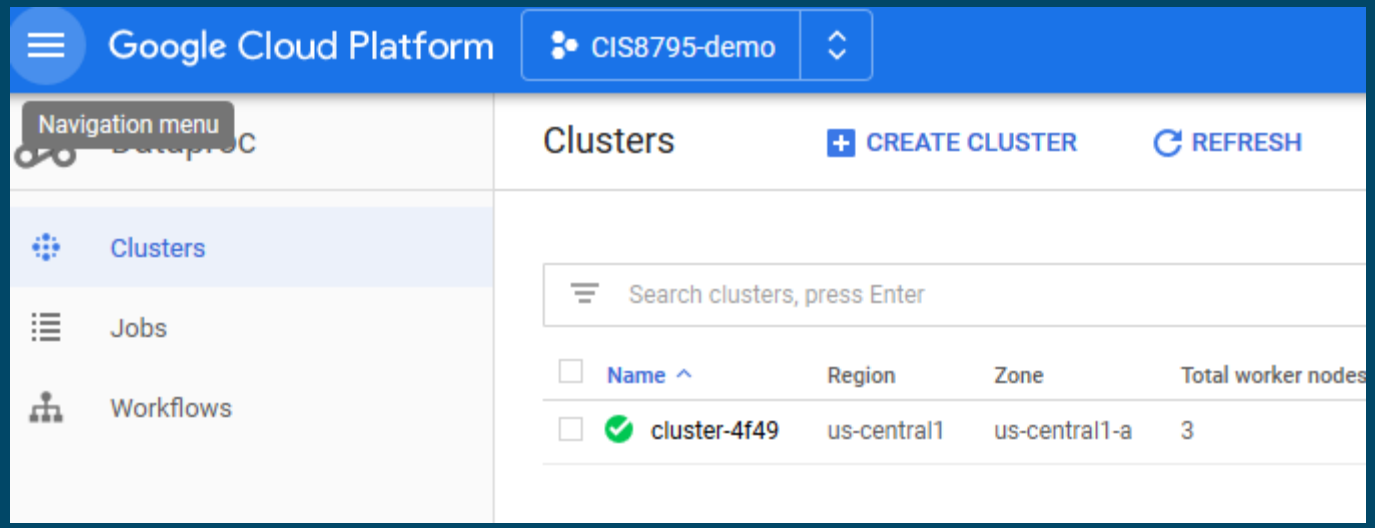

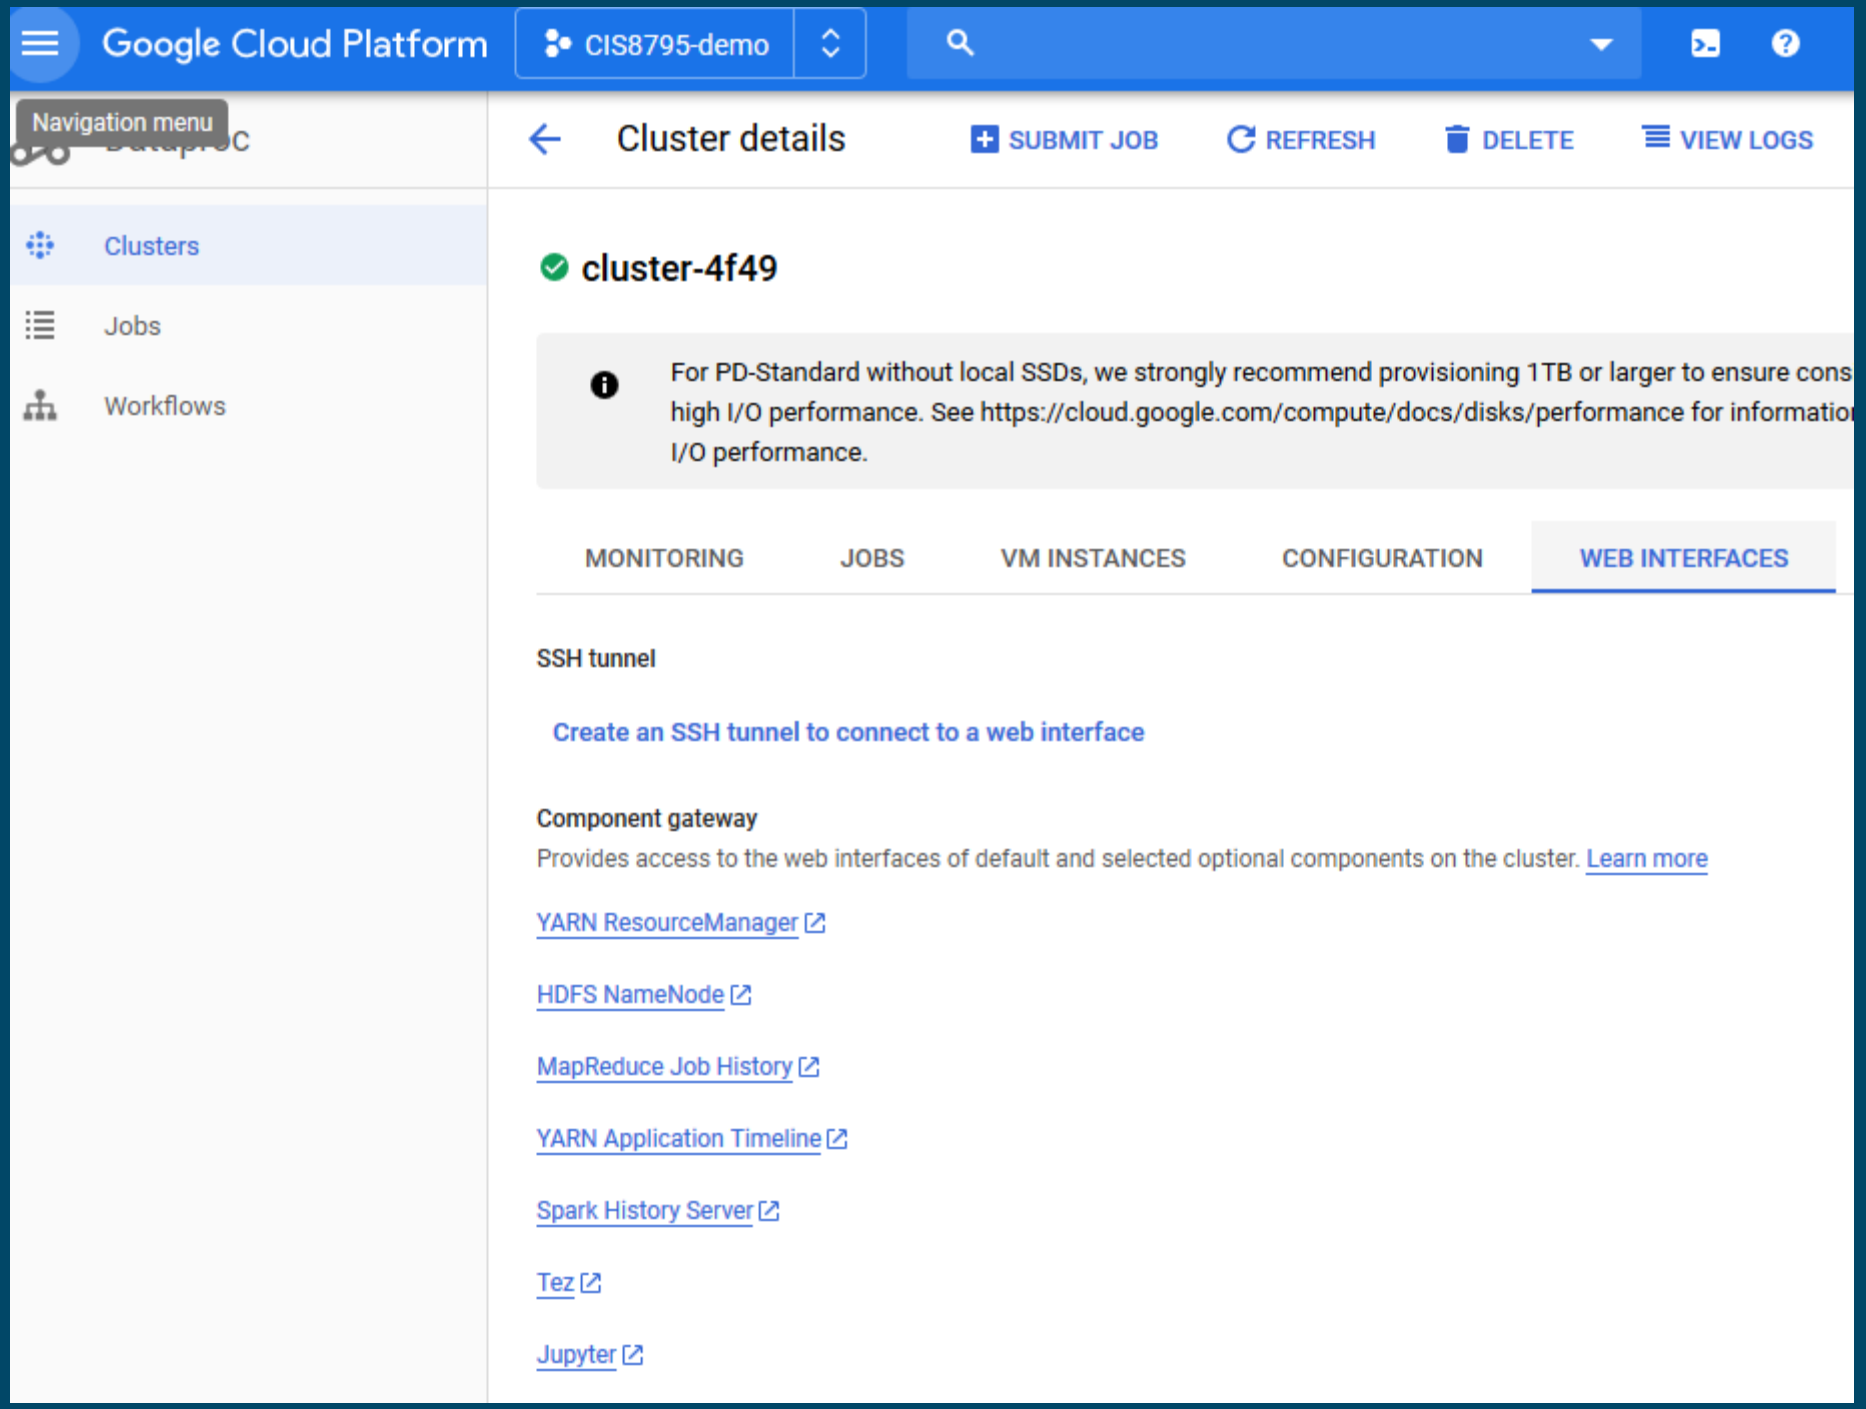
Google Cloud Platform
CIS8795-demo
Navigation menu
Clusters
Jobs
Workflows
Cluster details
SUBMIT JOB
REFRESH
DELETE
VIEW LOGS
cluster-4f49
For PD-Standard without local SSDs, we strongly recommend provisioning 1TB or larger to ensure cons
high I/O performance. See https://cloud.google.com/compute/docs/disks/performance for informatio
I/O performance.
MONITORING
JOBS
VM INSTANCES
CONFIGURATION
WEB INTERFACES
SSH tunnel
Create an SSH tunnel to connect to a web interface
Component gateway
Provides access to the web interfaces of default and selected optional components on the cluster. Learn more
YARN ResourceManager
HDFS NameNode
MapReduce Job History
YARN Application Timeline
Spark History Server
Tez
Jupyter

Opening the Jupyter Notebook

- Wait for 10 minutes or so
- Click on the cluster
- Navigate to “WEB INTERFACES”
- Select “Jupyter”
- Click on the link and navigate to the home page
- Select the .ipynb file
- Note: source files should be accessed in the below way

  Format : df = spark.read.csv('gs://<bucket-name>/<csv-file>', header=True, inferSchema=True)

  Example : df = spark.read.csv('gs://**newbucket-bdidemo**/cruise_ship_info.csv', header=True, inferSchema=True)

Important to remember

- Dataproc is managed infrastructure for spark clusters
  - Allows user to configure nodes
    - Number & size
    - Security & network options
  - Supports PySpark notebooks