

Docker illustrated

<https://training.play-with-docker.com/>

Docker (and it's Swarm)

- Supports Infrastructure as service (IaaS)
 - Declaratively specify apps in containers and their service connections over their own network
 - Isolation, increased security, simplified maintenance, supporting Agile software development
- Elastic run time
 - Increase/decrease nodes as needed for demand or maintenance
 - More automation with Kubernetes than Swarm
- Such container technologies behind Amazon, Databricks, etc.
 - Yarn / HDFS a specialized variant for data processing
 - It's possible to run Spark in Docker swarm

Online Docker

- <https://training.play-with-docker.com/ops-sl-hello/>
 - Training labs
 - Simpler interface
- <https://labs.play-with-docker.com/>
 - Better interface

Hello World

- First run
 - Must pull image from store.docker.com

```
$ docker container run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
1b930d010525: Pull complete
Digest: sha256:92695bc579f31df7a63da6922075d0666e565ceccad16b59c3374d2cf4e8e50e
Status: Downloaded newer image for hello-world:latest

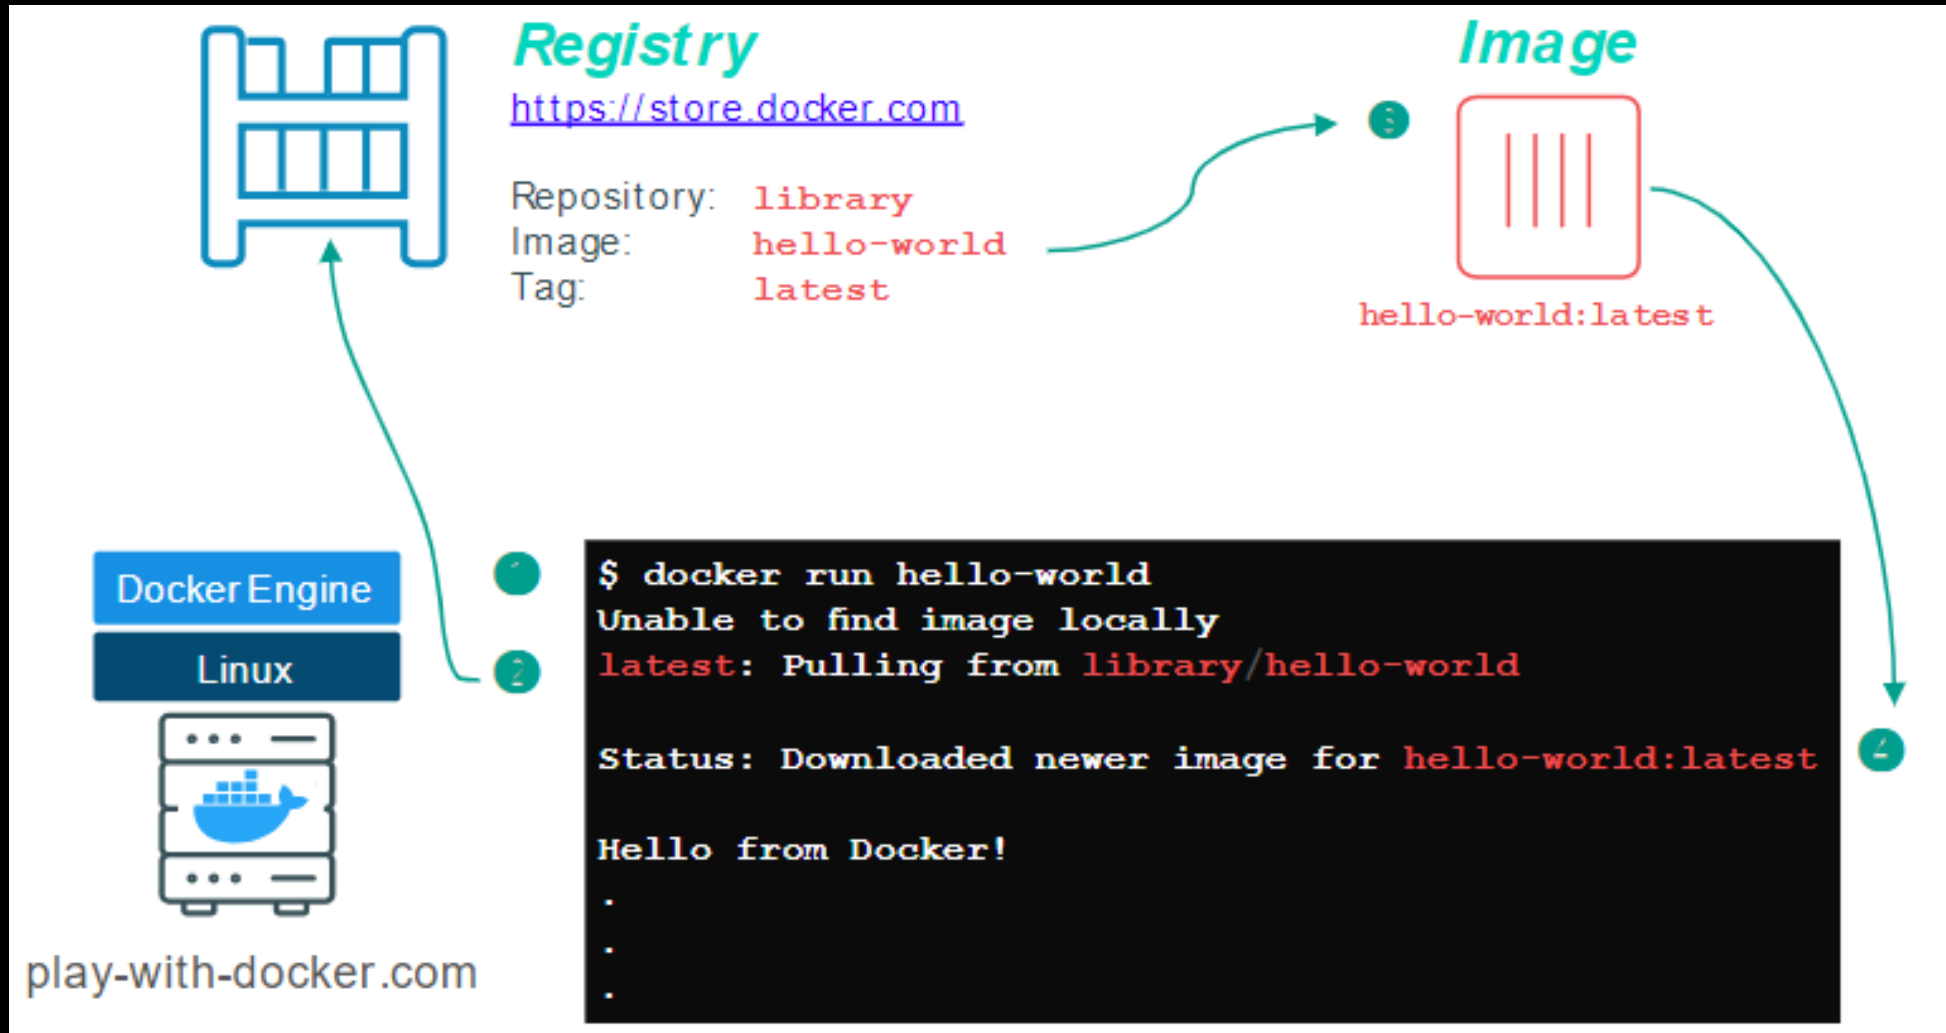
Hello from Docker!
```

- Second run
 - Image is local
 - Run it

```
$ docker container run hello-world

Hello from Docker!
This message shows that your installation appears to be working correctly.
```

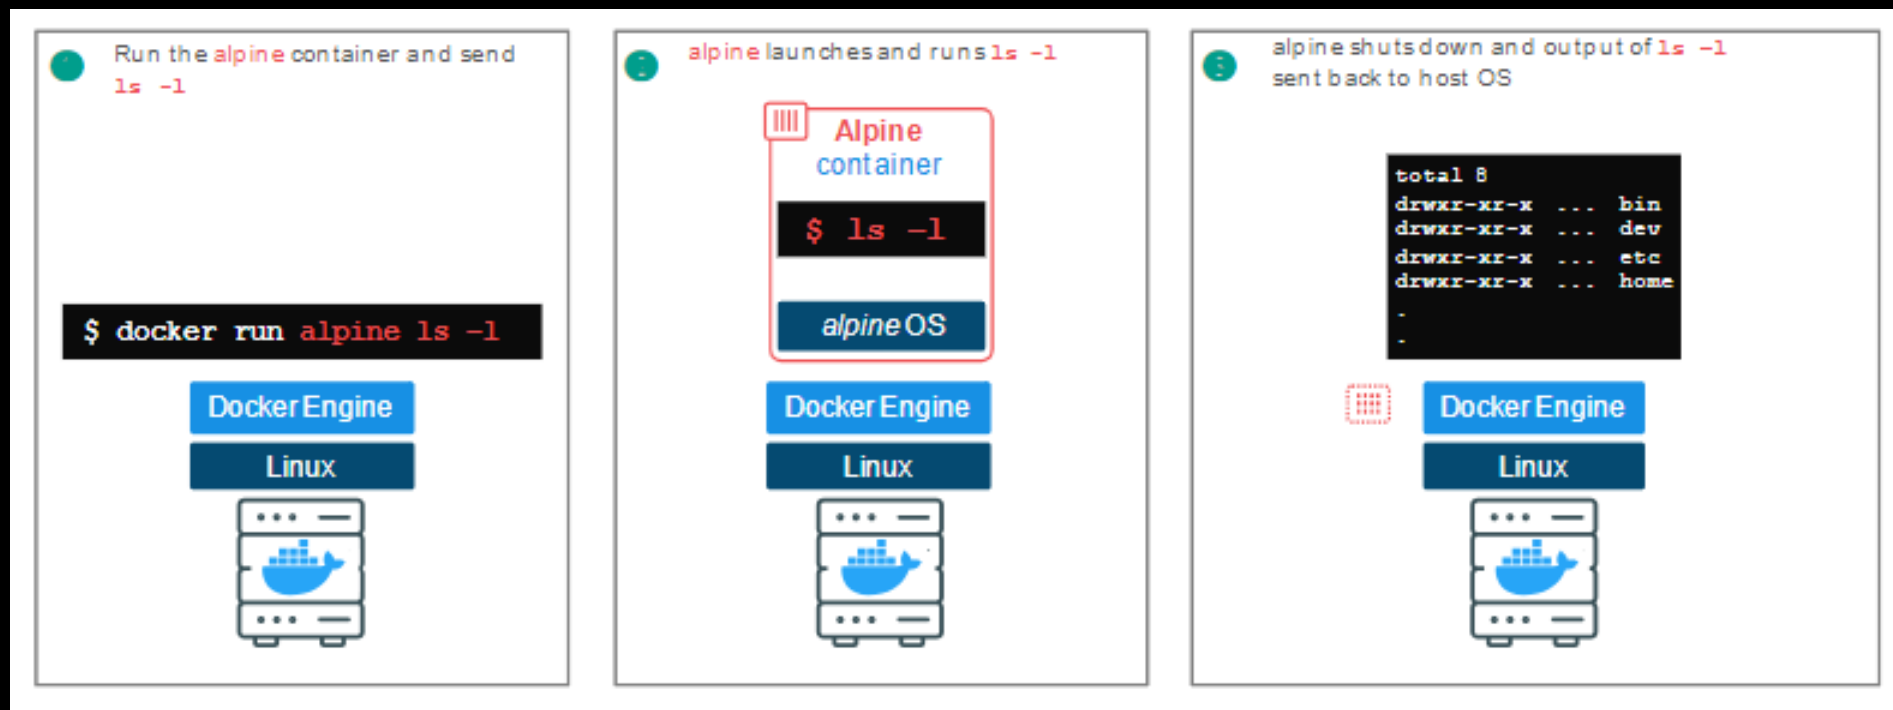
Hello World illustrated



run alpine ls -l

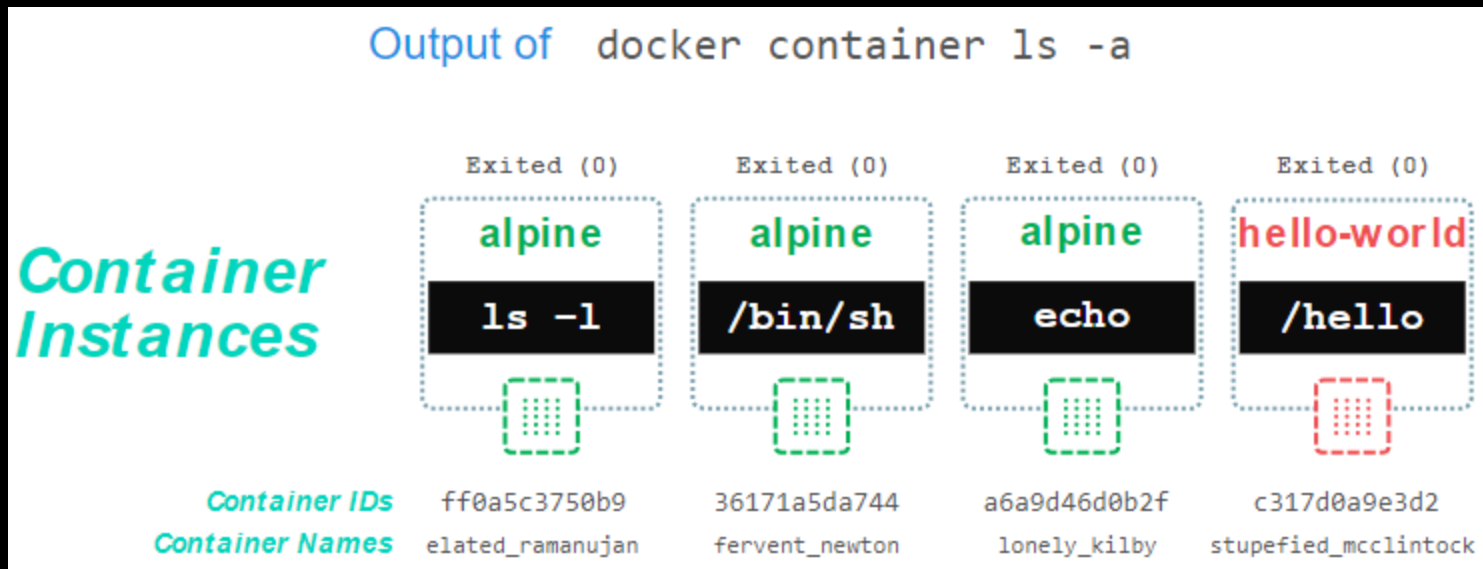
- Start container
- Send “ls -l” to container

```
$ docker container run alpine ls -l
total 8
drwxr-xr-x    2 root    root          4096 Apr  8 20:30 bin
drwxr-xr-x    5 root    root          340 Apr 16 14:49 dev
drwxr-xr-x    1 root    root           66 Apr 16 14:49 etc
drwxr-xr-x    2 root    root           6 Apr  8 20:30 home
drwxr-xr-x    5 root    root         185 Apr  8 20:30 lib
drwxr-xr-x    5 root    root          44 Apr  8 20:30 media
```



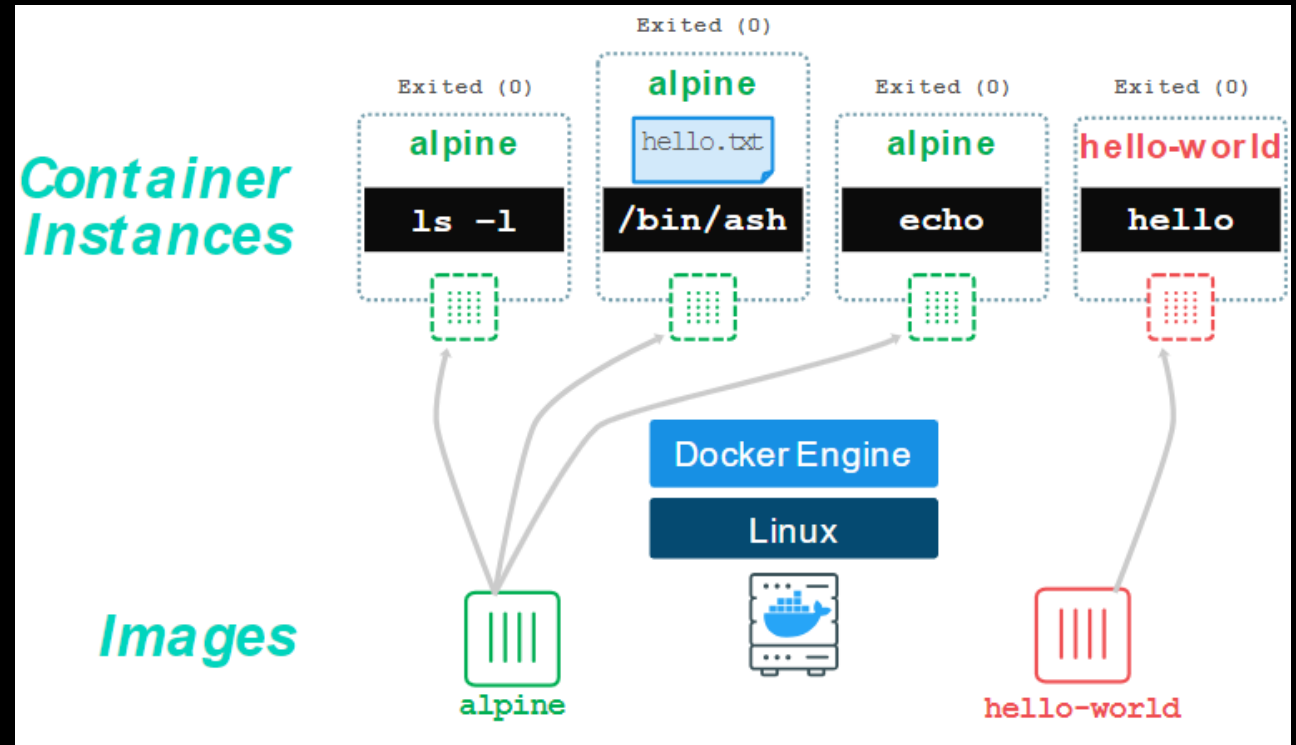
Container history

- **Running containers:** `docker ps`
- **All containers:** `docker ps -a`



Containers are isolated from each other

- Run creates a new container
- Typically Linux
 - Linux and Microsoft containers
- Custom software loaded



Docker run vs start

1. Run: create a new container of an image, and execute the container. You can create N clones of the same image. The command is: `docker run IMAGE_ID` **and not** `docker run CONTAINER_ID`

```
PS C:\Users\Daniele> docker run --help
Usage:  docker run [OPTIONS] IMAGE [COMMAND] [ARG...]
Run a command in a new container
```

2. Start: Launch a container previously stopped. For example, if you had stopped a database with the command `docker stop CONTAINER_ID`, you can relaunch the same container with the command `docker start CONTAINER_ID`, and the data and settings will be the same.

```
PS C:\Users\Daniele> docker start --help
Usage:  docker start [OPTIONS] CONTAINER [CONTAINER...]
Start one or more stopped containers
```

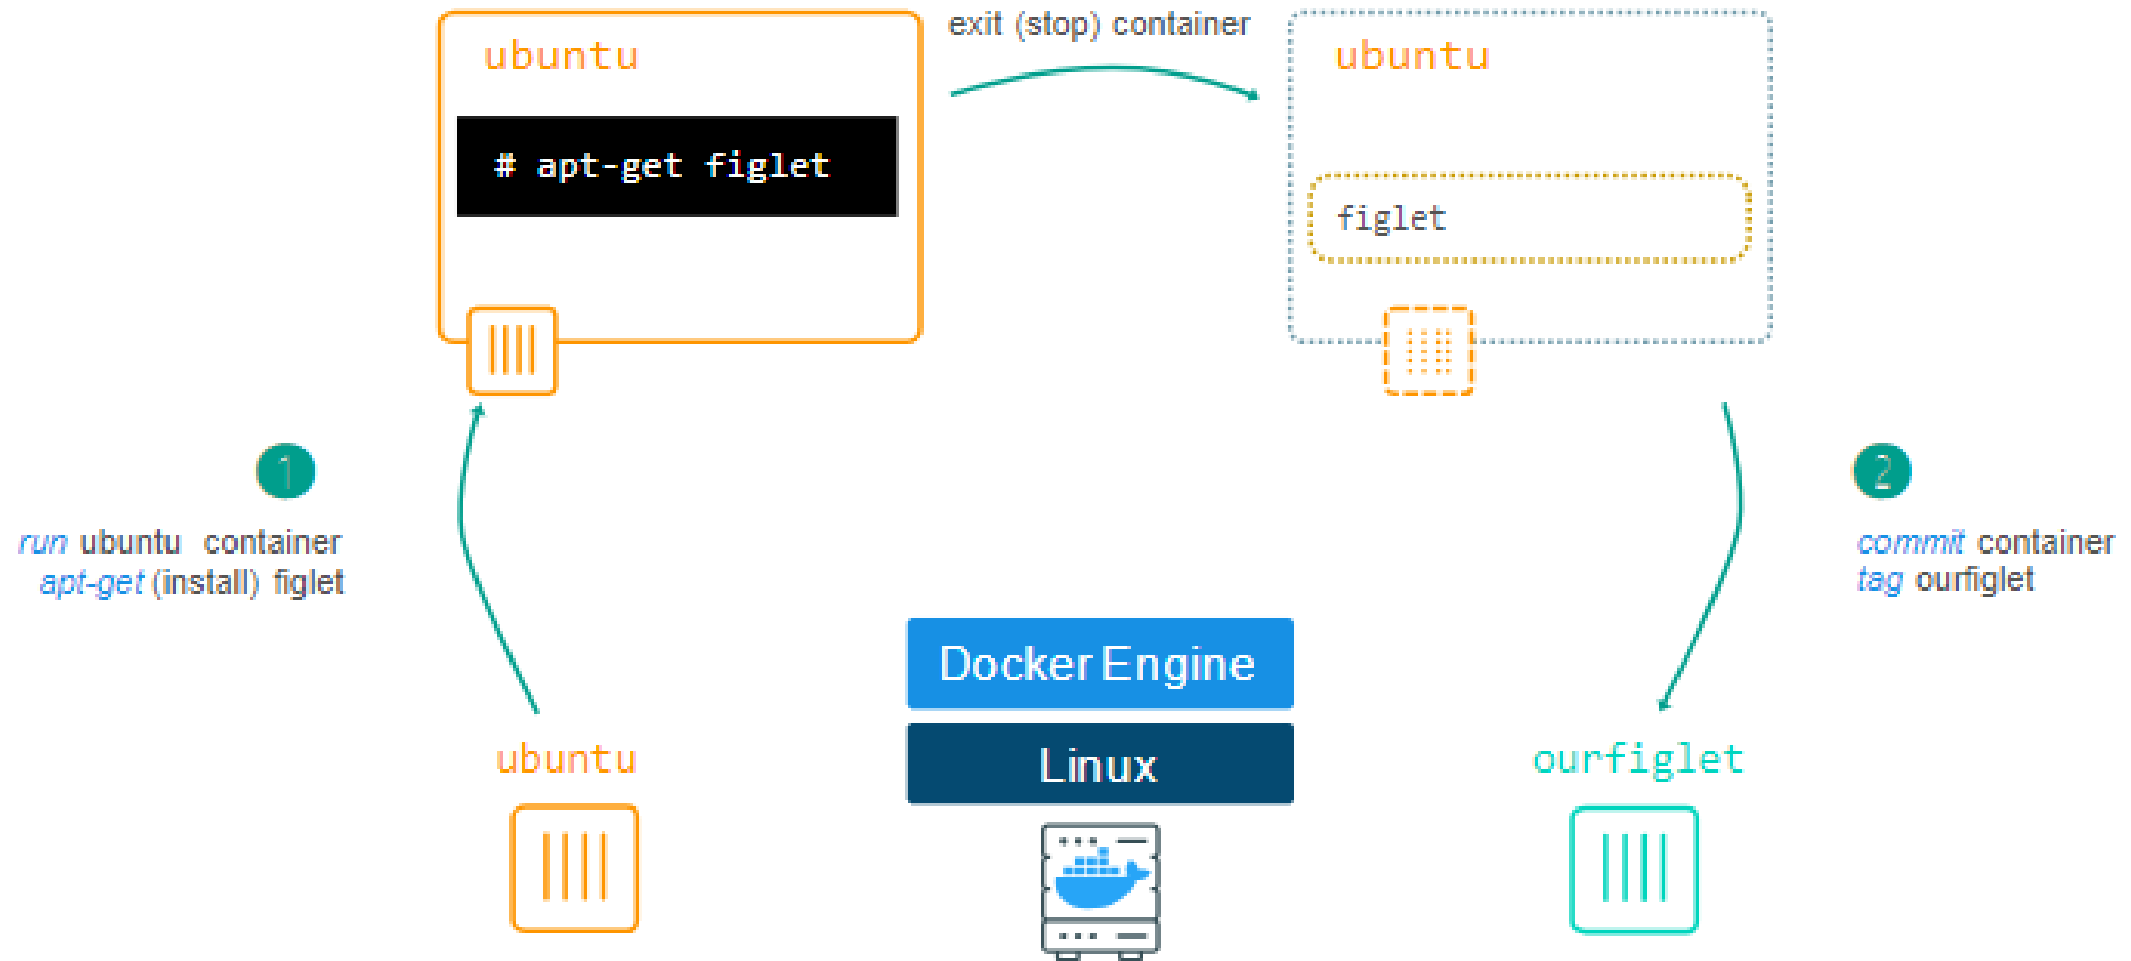
Docker image

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
ubuntu	latest	94e814e2efa8	5 weeks ago	88.9MB

- A **Docker image** is a file, comprised of multiple layers, used to execute code in a **Docker** container.
- An **image** is essentially built from the instructions for a complete and executable version of an application, which relies on the host OS kernel.

Create an image

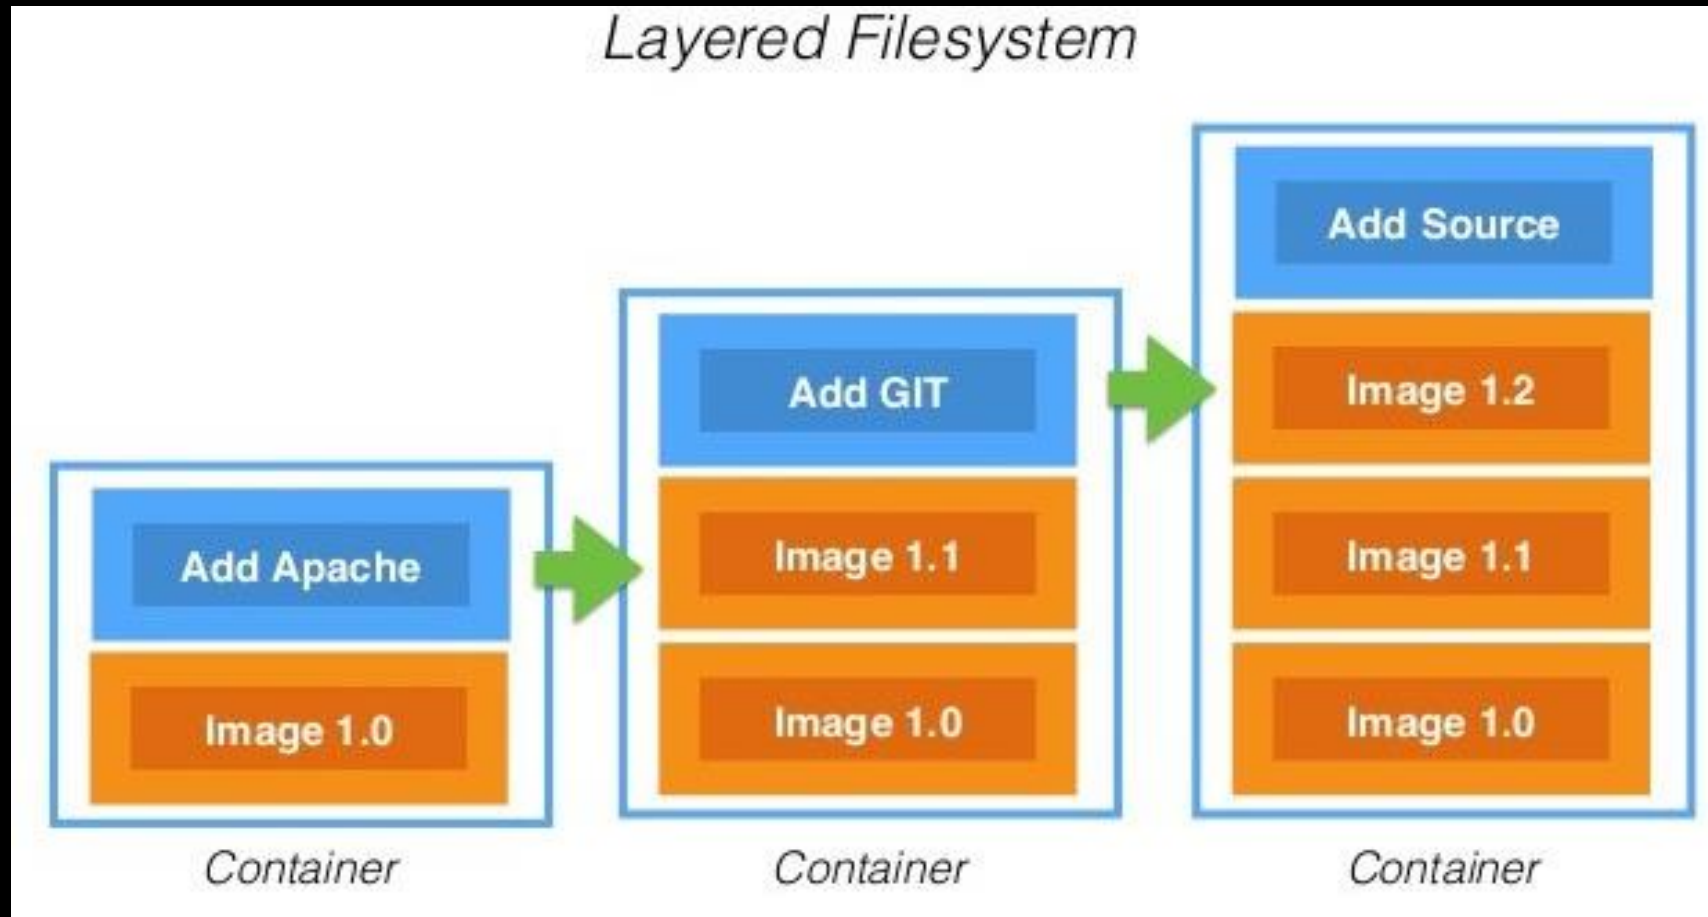


Running our new image in a container

```
$ docker image tag fb ourfiglet
[node1] (local) root@192.168.0.13 ~
$ docker container run ourfiglet figlet hello

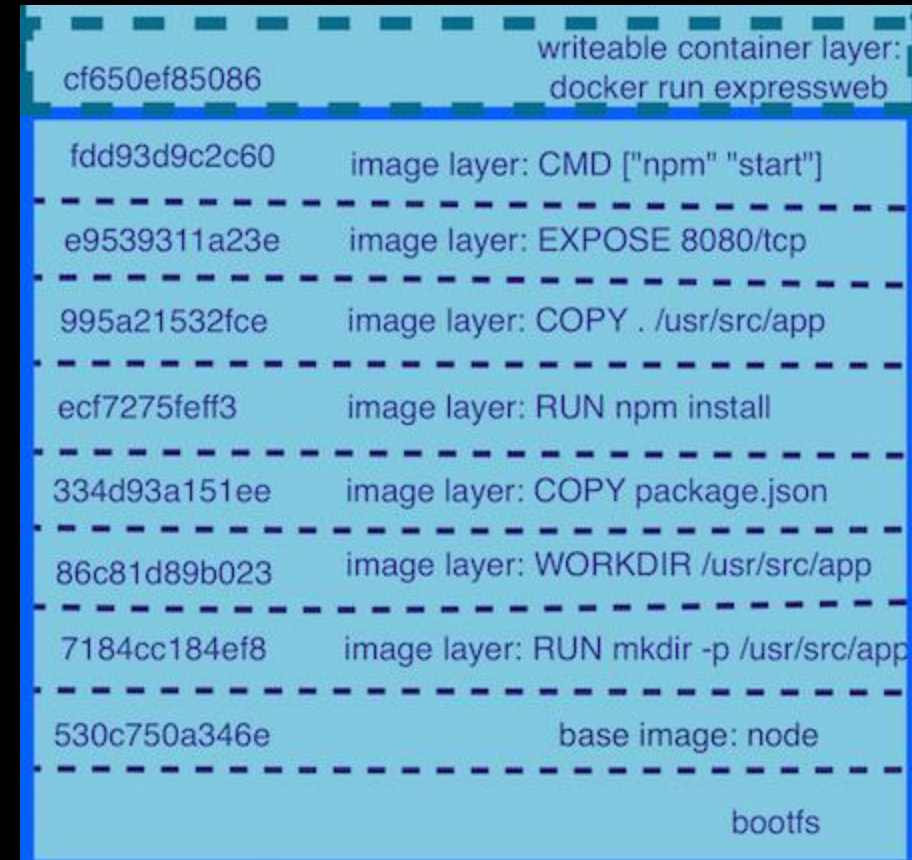
 _
| | _ _ _ | | | _
| ' \ / _ \ | | / _ \
| | | _ / | | ( ) |
| | | \ _ | | | \ _ /
```

Docker images defined with layers

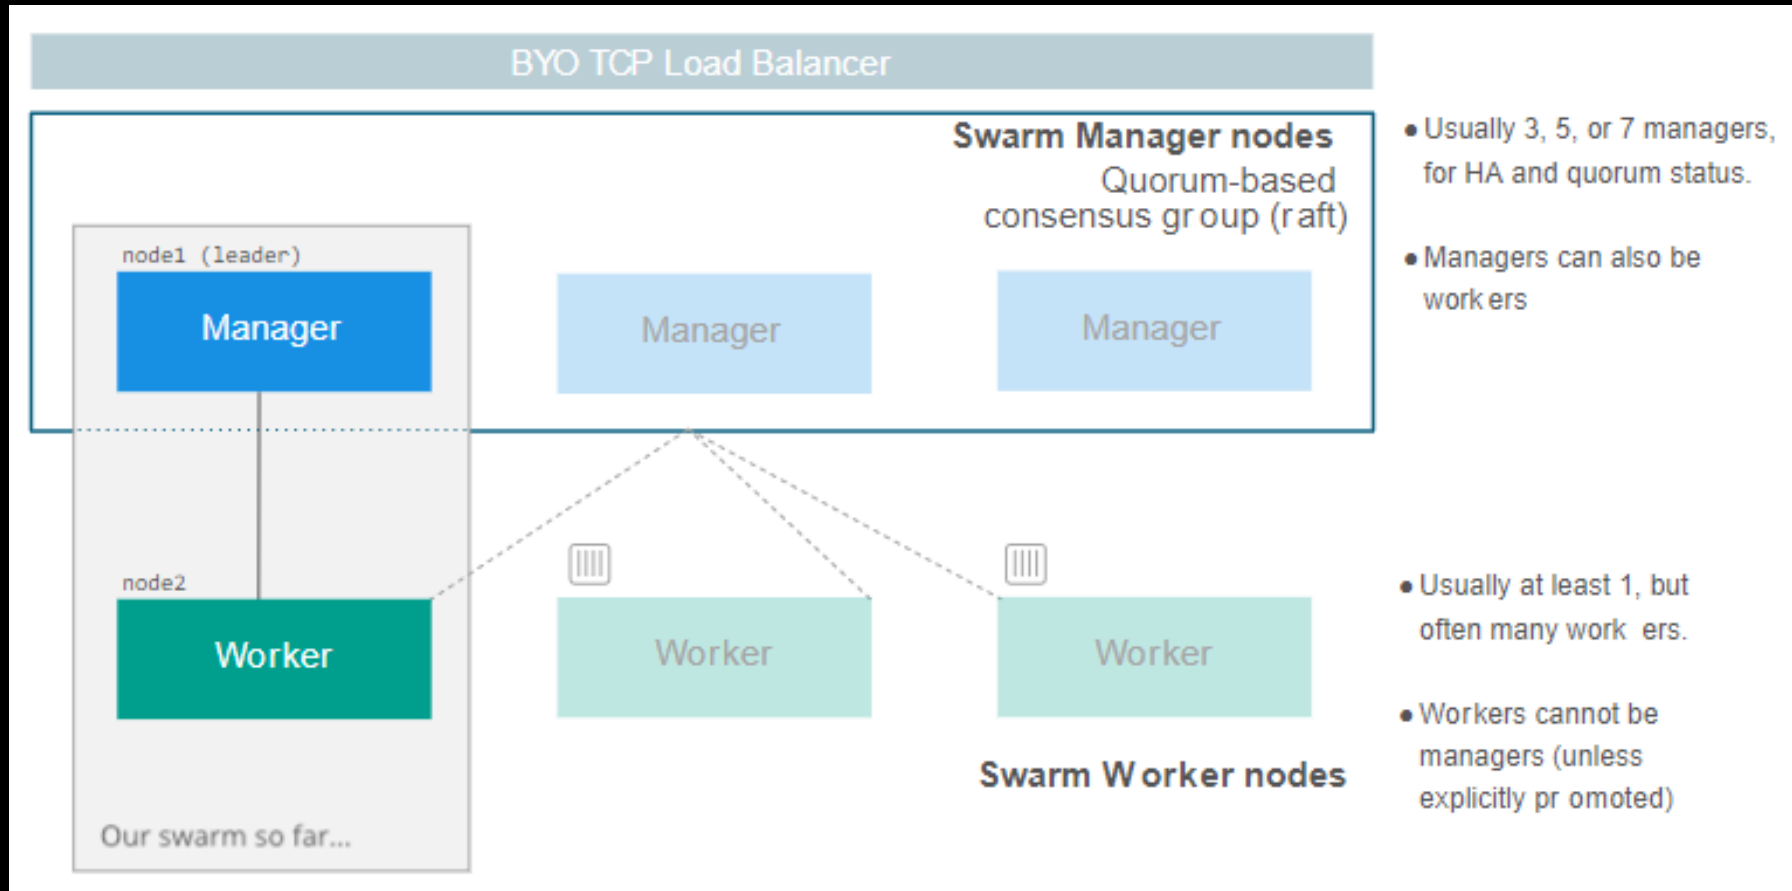


Docker image layers

- Layers (also called intermediate images) are generated when the commands in the Dockerfile are executed during the Docker image build.
- Updating or customizing an image is done by adding layers
 - Can be very fast for simple updates



Docker Swarm (networked containers)



Example of running a Swarm

- `git clone https://github.com/docker/example-voting-app`
 - Downloads sources
 - Contains docker-compose file, which has instructions on composing containers (pull and deploy)

```
git clone https://github.com/docker/example-voting-app
cd example-voting-app
```

- Deploy the stack
 - A stack is a group of services that are deployed together
 - Each individual service is made up of one or more containers, called tasks
 - The tasks & services together make up a stack
 - Defined in a YAML file

```
docker stack deploy --compose-file=docker-stack.yml voting_stack
```

```
docker stack ls
```


Voting stack YAML file

services:

vote:

build: ./vote

command: `python app.py`

volumes:

- ./vote:/app

ports:

- "5000:80"

networks:

- front-tier
- back-tier

result:

build: ./result

command: `nodemon server.js`

volumes:

- ./result:/app

ports:

- "5001:80"
- "5858:5858"

networks:

- front-tier
- back-tier

worker:

build:

context: `./worker`

depends_on:

- "redis"

networks:

- back-tier

redis:

image: `redis:alpine`

container_name: redis

ports: ["6379"]

networks:

- back-tier

db:

image: `postgres:9.4`

container_name: db

volumes:

- "db-data:/var/lib/postgresql/data"

networks:

- back-tier

volumes:


db-data:

networks:

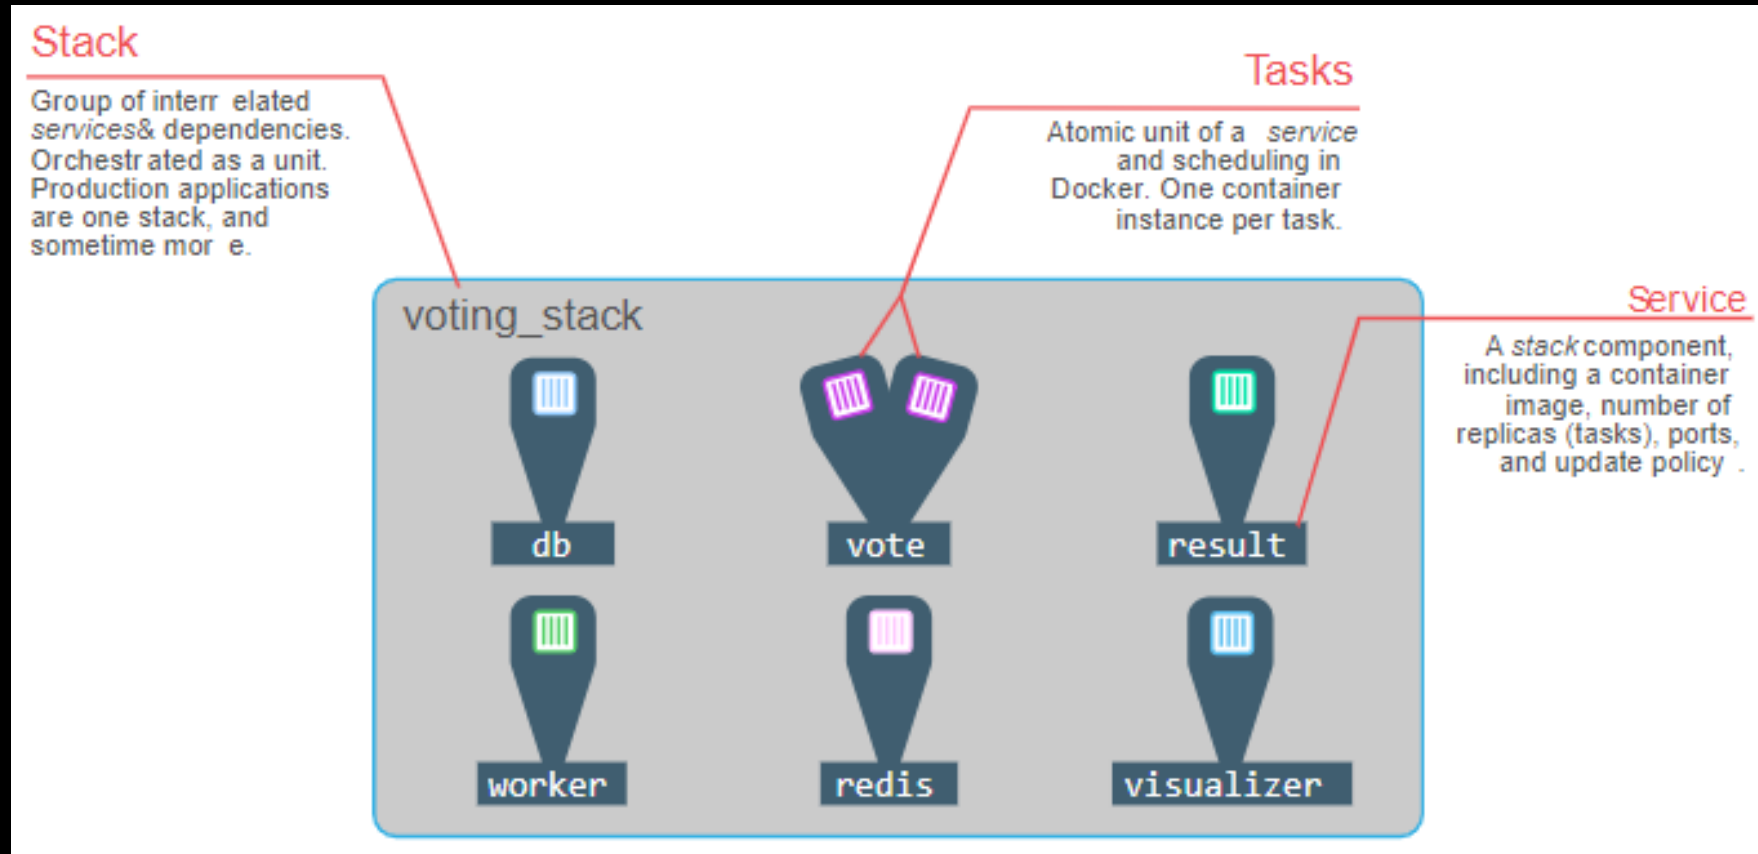
front-tier:

back-tier:

App Container architecture



Swarm: Stacks, Services, Tasks



Docker stack services

- View the **services** of a running **stack**

```
docker stack services voting_stack
```

```
$ docker stack services voting_stack
```

ID	NAME	MODE	REPLICAS	IMAGE
	PORTS			
h00br1j1eewu	voting_stack_vote	replicated	2/2	dockersamples/examplevotingap
p_vote:before	*:5000->80/tcp			
nwu8tihc7t1l	voting_stack_redis	replicated	1/1	redis:alpine
r9ms1sg92fuq	voting_stack_visualizer	replicated	1/1	dockersamples/visualizer:stab
le	*:8080->8080/tcp			
sqmkg3xgt7xp	voting_stack_db	replicated	1/1	postgres:9.4
v8ufgmz3f8rv	voting_stack_result	replicated	1/1	dockersamples/examplevotingap
p_result:before	*:5001->80/tcp			
z6az706cz16f	voting_stack_worker	replicated	1/1	dockersamples/examplevotingap
p_worker:latest				

- If there are 0 replicas just wait a few seconds and enter the command again

Docker tasks from running services

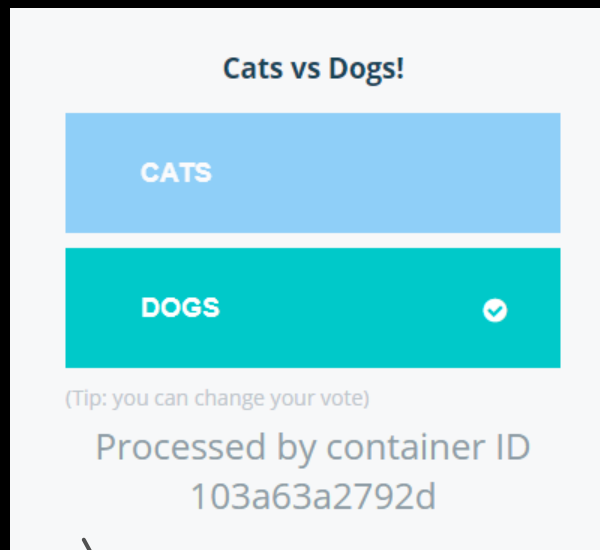
- View the **tasks** (replicas) of a running **service**
 - service ps only shows the **running** tasks of the service
 - Here, just the voting stack

```
docker service ps voting_stack_vote
```

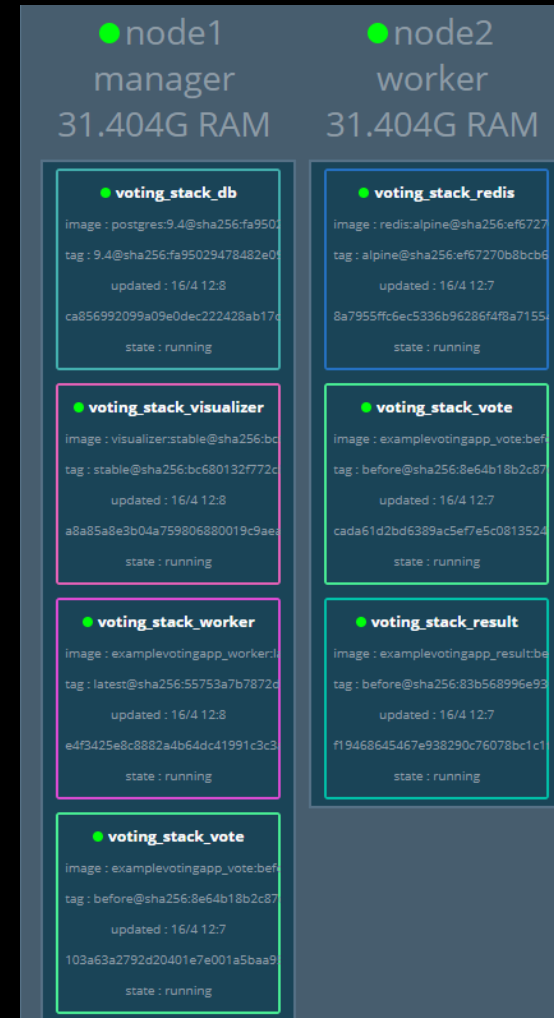
```
$ docker service ps voting_stack_vote
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STA
ec0er9ap3f4a	voting_stack_vote.1	dockersamples/examplevotingapp_vote:before	node2	Running	Running 7 m
5wiq4fosmyeb	voting_stack_vote.2	dockersamples/examplevotingapp_vote:before	node1	Running	Running 7 m

Swarm visualizer: nodes and their containers



Select built-in SWARM VISUALIZER, front-end web UI, and the result page on play with docker




Scale the application

- Assume that you need more computing for the voting service
- To increase the replicas of a `docker service scale voting_stack_vote=5`

```
$ docker stack services voting_stack
```

ID	NAME	MODE	REPLICAS
h00br1j1eewu	voting_stack_vote	replicated	5/5
0->80/tcp			
nwu8tihc7t11	voting_stack_redis	replicated	1/1
r9ms1sg92fuq	voting_stack_visualizer	replicated	1/1
0->8080/tcp			
sqmkg3xgt7xp	voting_stack_db	replicated	1/1
v8ufgmz3f8rv	voting_stack_result	replicated	1/1
1->80/tcp			
z6az706cz16f	voting_stack_worker	replicated	1/1



node2
worker
31.404G RAM

voting_stack_db

voting_stack_vote

voting_stack_redis

voting_stack_visualizer

voting_stack_vote

voting_stack_worker

voting_stack_vote

node2
worker
31.404G RAM

voting_stack_vote

voting_stack_vote

voting_stack_vote

Reduce scale & Drain a node

- Scale down to 2 replicas

```
$ docker service scale voting_stack_vote=2
voting_stack_vote scaled to 2
overall progress: 2 out of 2 tasks
1/2: running
2/2: running
```

- Bring a node down for maintenance

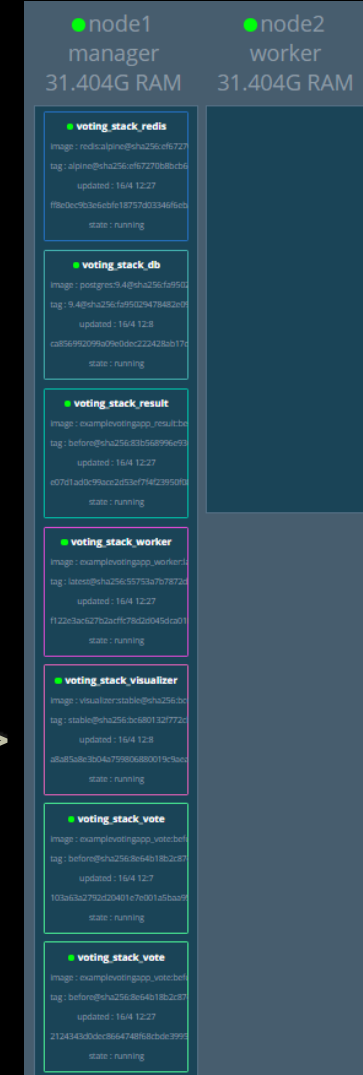
- docker node update --availability drain node2

- Bring it back up

- docker node update --availability active <NODE-ID>

- Ensure services migrate to updated node

```
$ docker service update --force voting_stack_vote
voting_stack_vote
overall progress: 2 out of 2 tasks
1/2: running [=====>]
2/2: running [=====>]
verify: Service converged
[node1] (local) root@192.168.0.12 ~/example-voting-app
$ docker service update --force voting_stack_redis
voting_stack_redis
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
[node1] (local) root@192.168.0.12 ~/example-voting-app
$ docker service update --force voting_stack_result
```



Common Docker swarm commands

- `docker service --help`
 - `logs`
 - `ls`
 - `ps`
 - `Scale`
- `docker stack --help`
 - `deploy`
 - `ls`
 - `ps`
 - `rm`
 - `services`

Do it yourself: Swarm lab

- On play with docker
 - <https://training.play-with-docker.com/ops-sl-swarm-intro/>

Docker Hub Image Demos

- <https://labs.play-with-docker.com/>
- Docker images
 - <https://hub.docker.com/r/jupyter/pyspark-notebook>
 - Too big for PWD, but on your own notebook
 - `docker run -p 8888:8888 jupyter/pyspark-notebook`
 - https://hub.docker.com/_/wordpress/
 - <https://hub.docker.com/r/bign8/games>

```
import pyspark
if not 'sc' in globals():
    sc = pyspark.SparkContext()

rdd = sc.parallelize(range(100))
rdd.count()
```

Docker (and it's Swarm)

- Supports Infrastructure as service (IaaS)
 - Declaratively specify apps in containers and their service connections over their own network
 - Isolation, increased security, simplified maintaince, supporting Agile software development
- Elastic run time
 - Increase/decrease nodes as needed for demand or maintenance
 - More automation with Kubernetes than Swarm
- Such container technologies behind Amazon, Databricks, etc.
 - Yarn / HDFS a specialized variant for data processing
 - It's possible to run Spark in Docker swarm